

Acquisition

Umberto Castellani
Robotics, vision and control

3D modelling from reality pipeline



Overall aim



Real Object



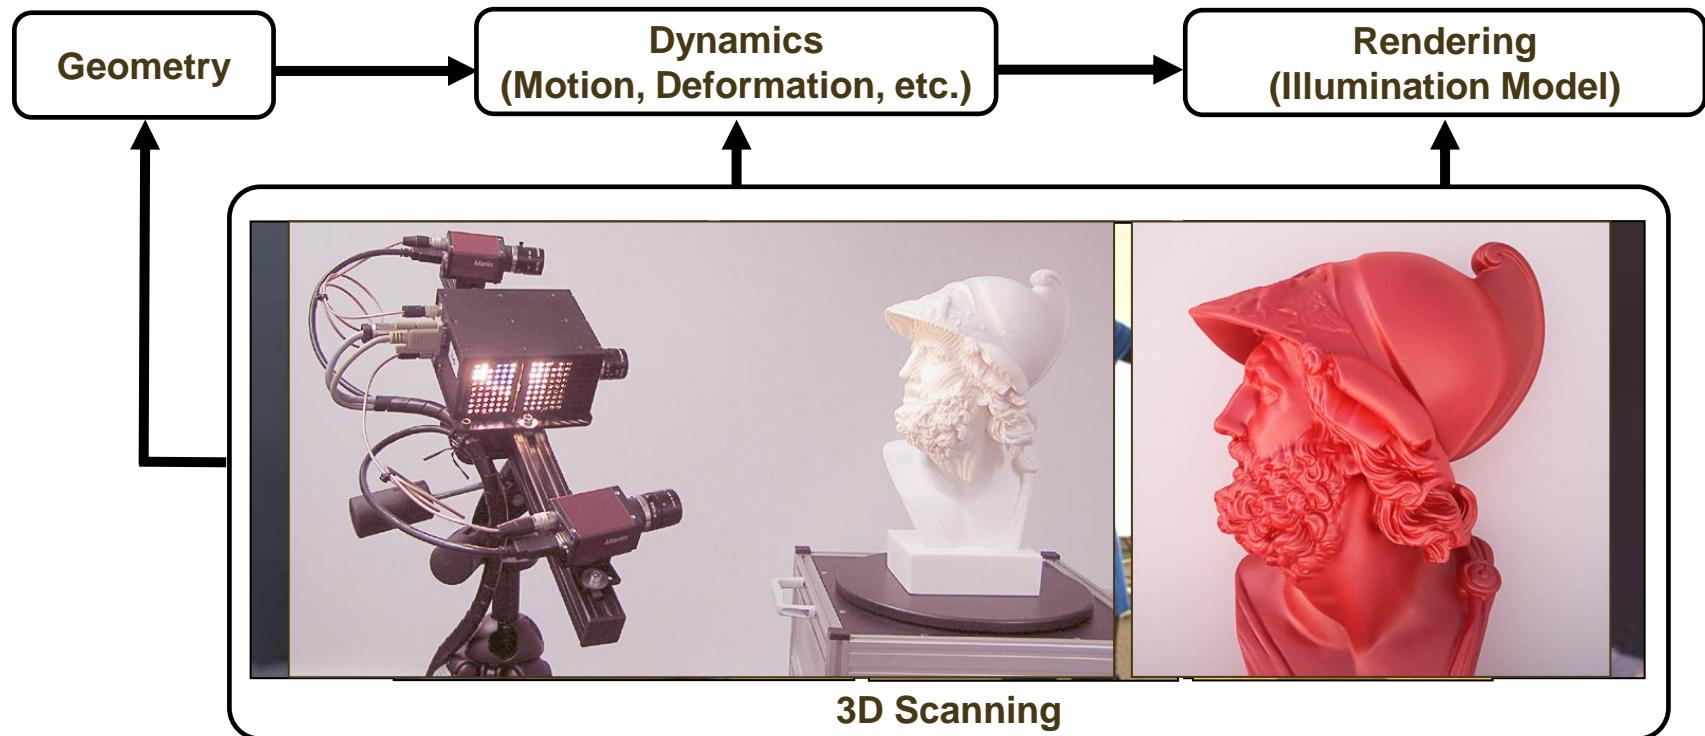
3D Scanning



Virtual Model

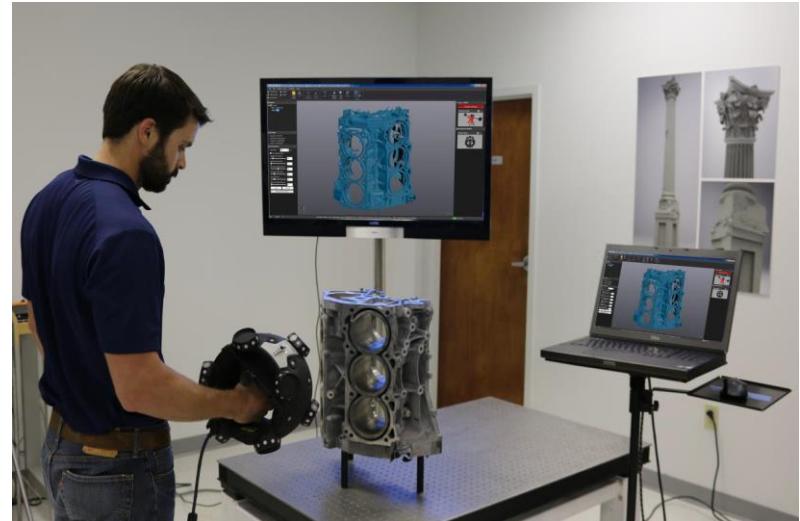
<https://youtu.be/r1e088hGee4>

Overall aim

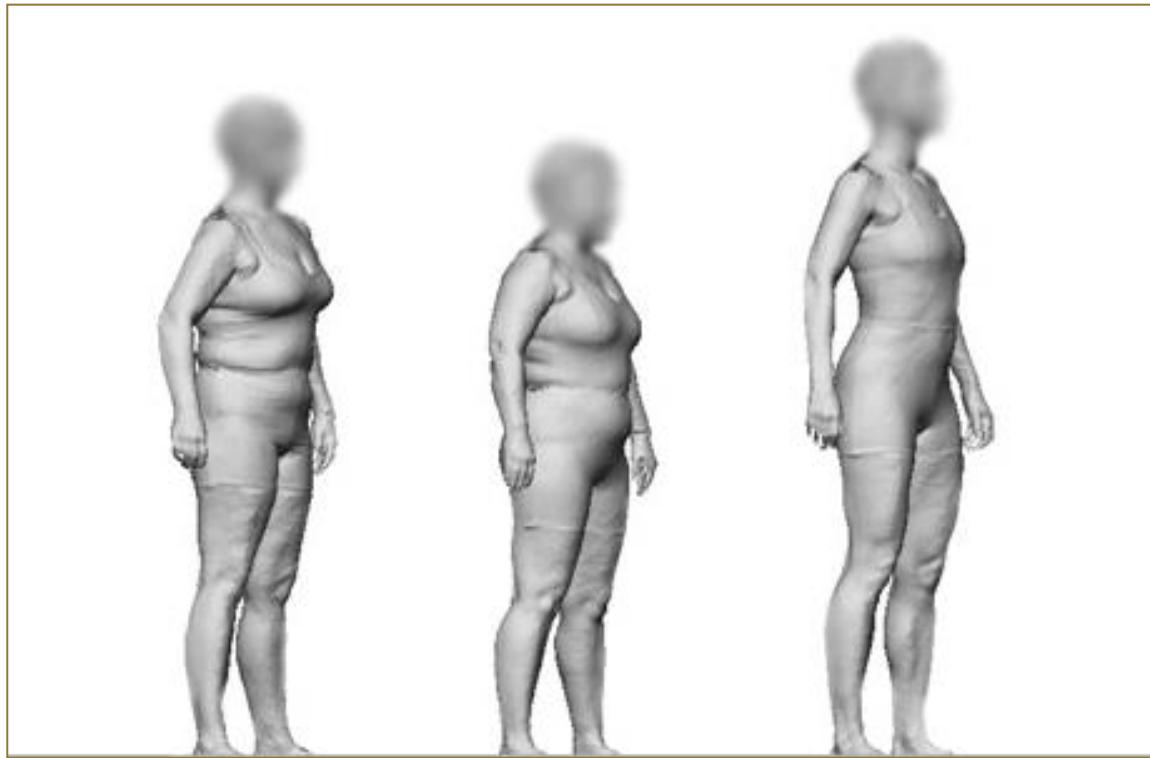


Overall aim

- Create a **virtual representation** of the real word through the **observation** of the real word
 - Geometric properties,
 - Dynamic properties,
 - Photometric properties,



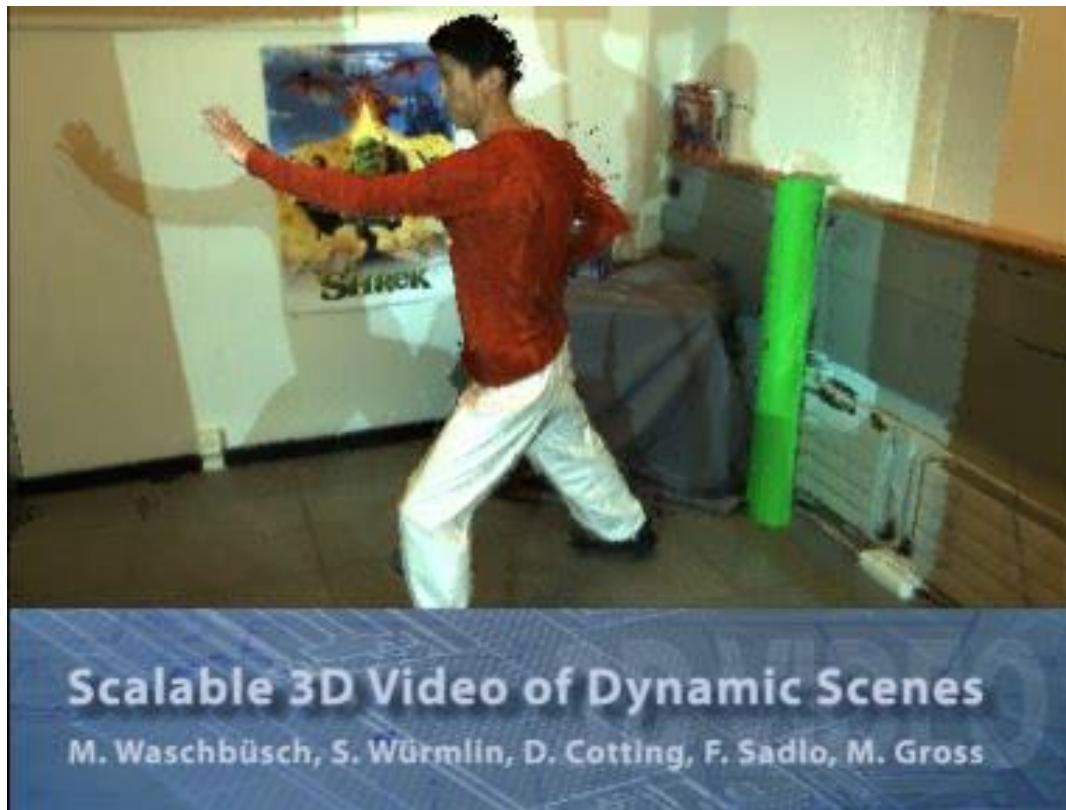
Geometric properties



Body scanner

Fit3D Body Scanner

Dynamic properties



Realtime capture: RGBD Scanning

Photometric properties



Light Dome

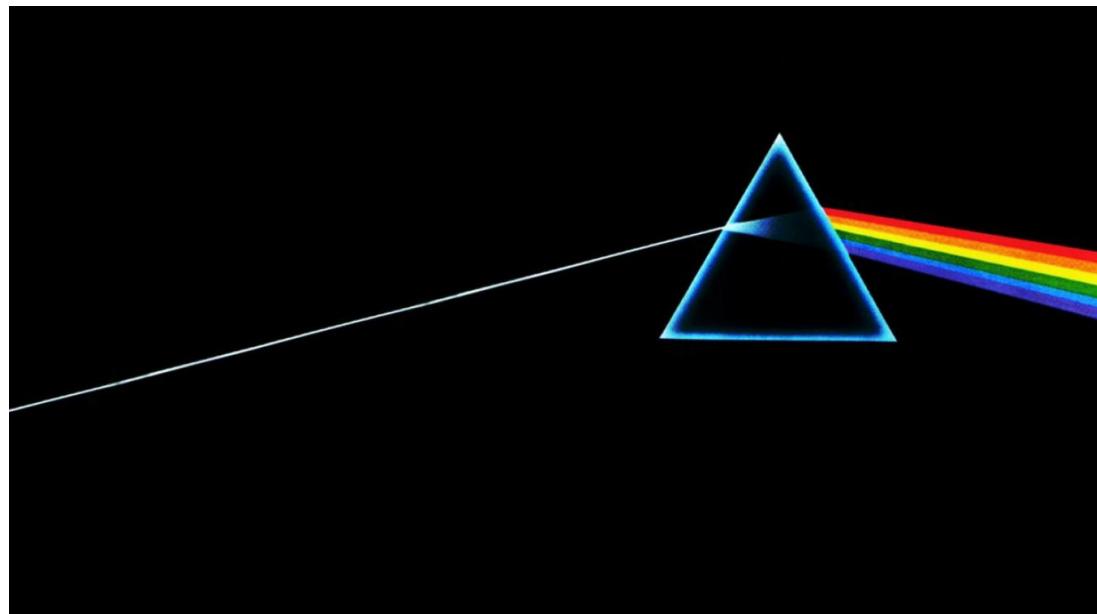
Overall aim

- Some names:

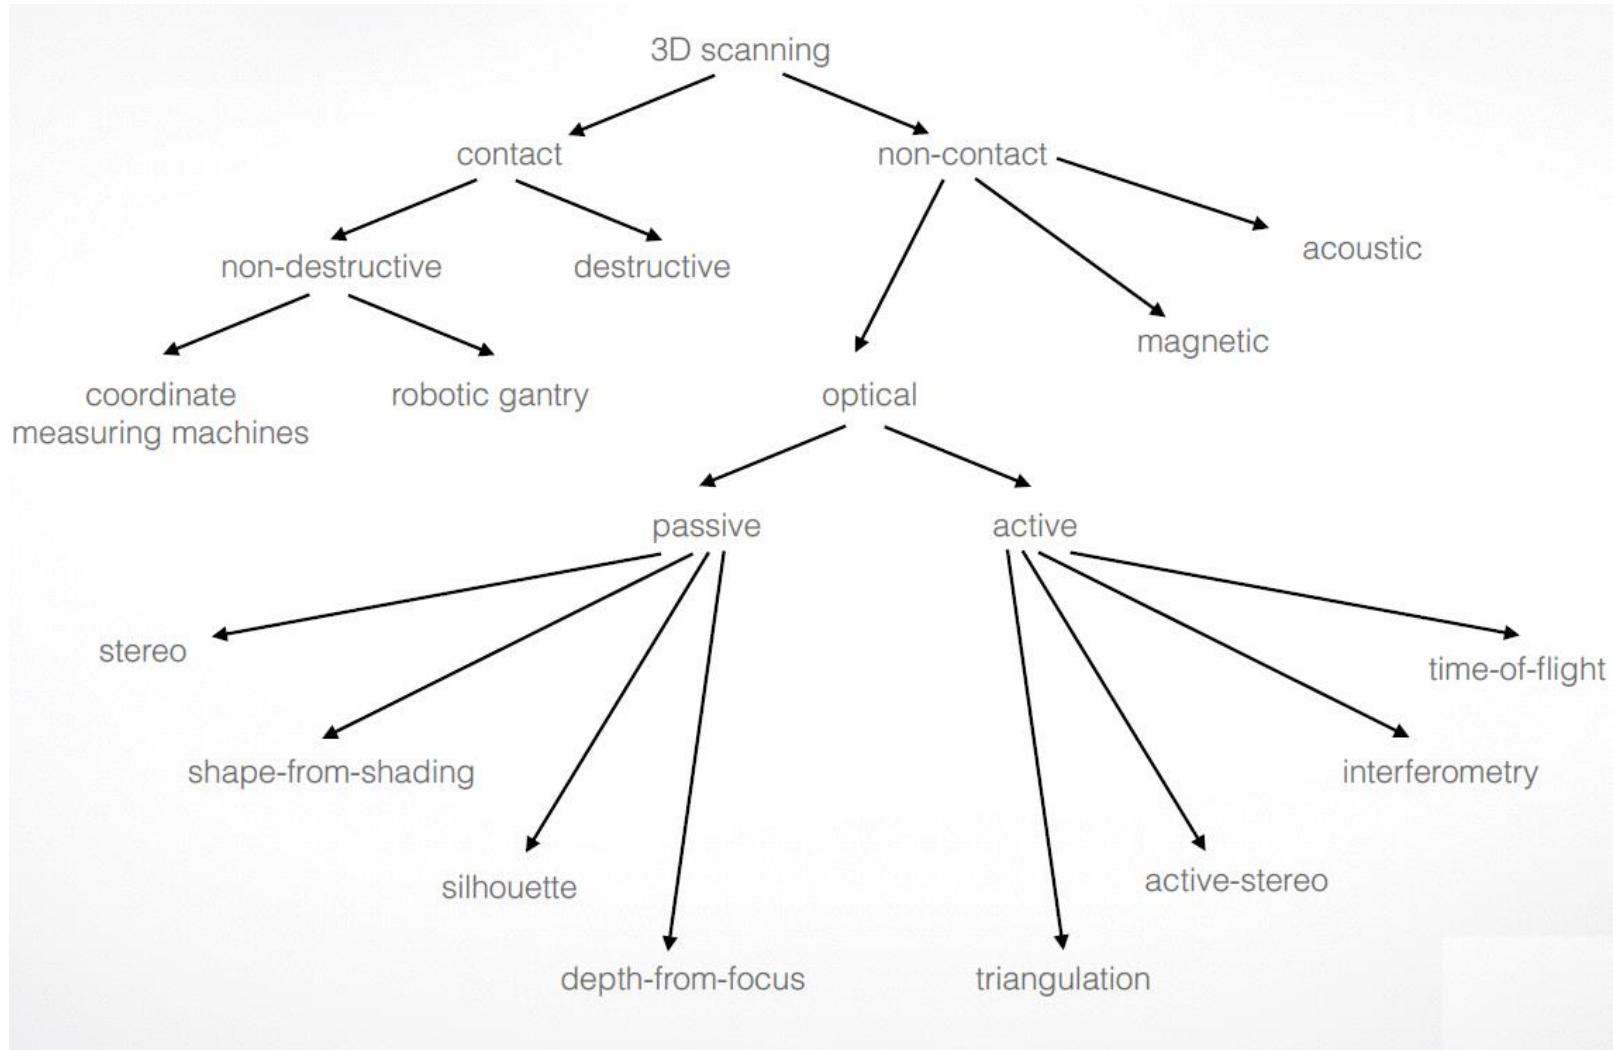
Reality capture, 3D scanning, 3D acquisition, Digital content creation, Range finder, Virtualized reality, RGBD, dynamic scene capture, reverse engineering, photogrammetry, computational photography, 3DTV,...

Principles

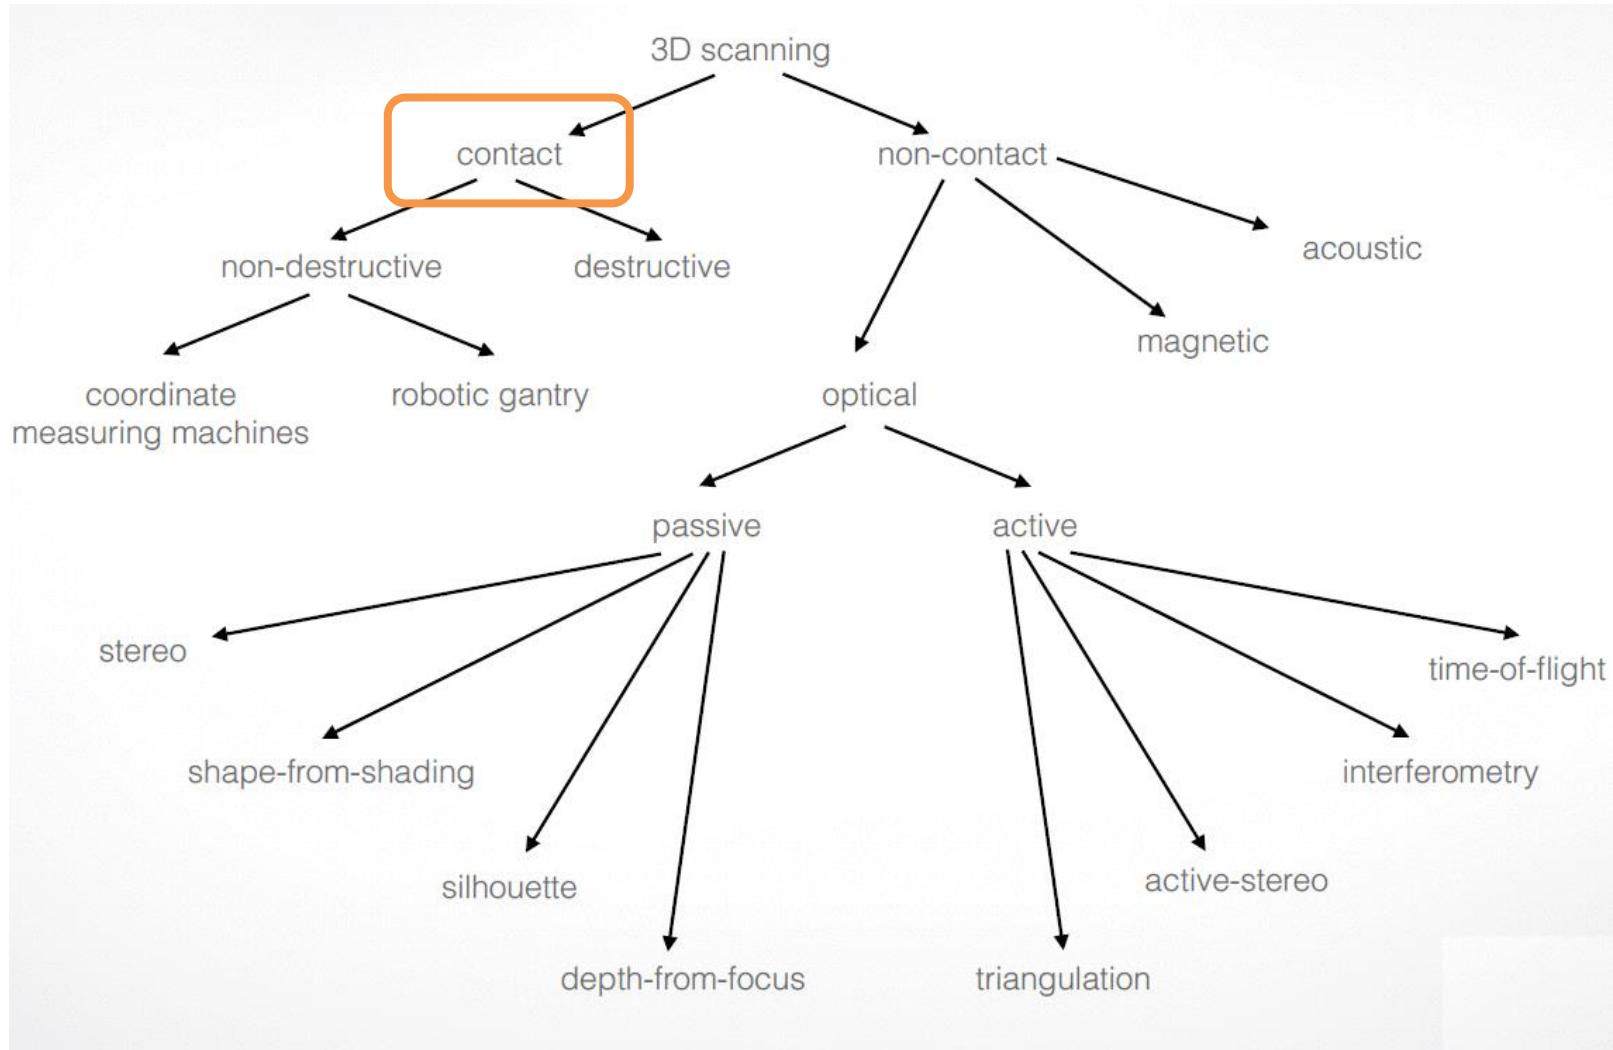
- So many acquisition techniques..., physics helps!
 - Each technique is based on a well founded physical principle



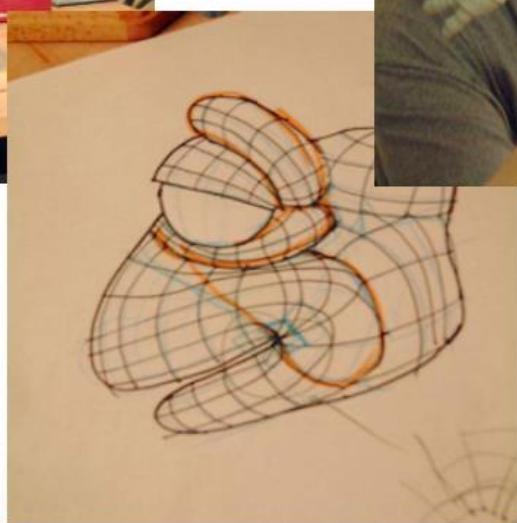
A taxonomy



A taxonomy



Contact Scanner



[immersion microscribe, magnetic dreams]: <http://www.hao-li.com/cs599-ss2015/slides/Lecture04.2.pdf>

Contact Scanner

Probe object by physical touch!

- used in manufacturing control,
- highly accurate,
- reflectance independent (transparency!)
- slow scanning, sparse set of samples,
- for rigid and non-fragile objects



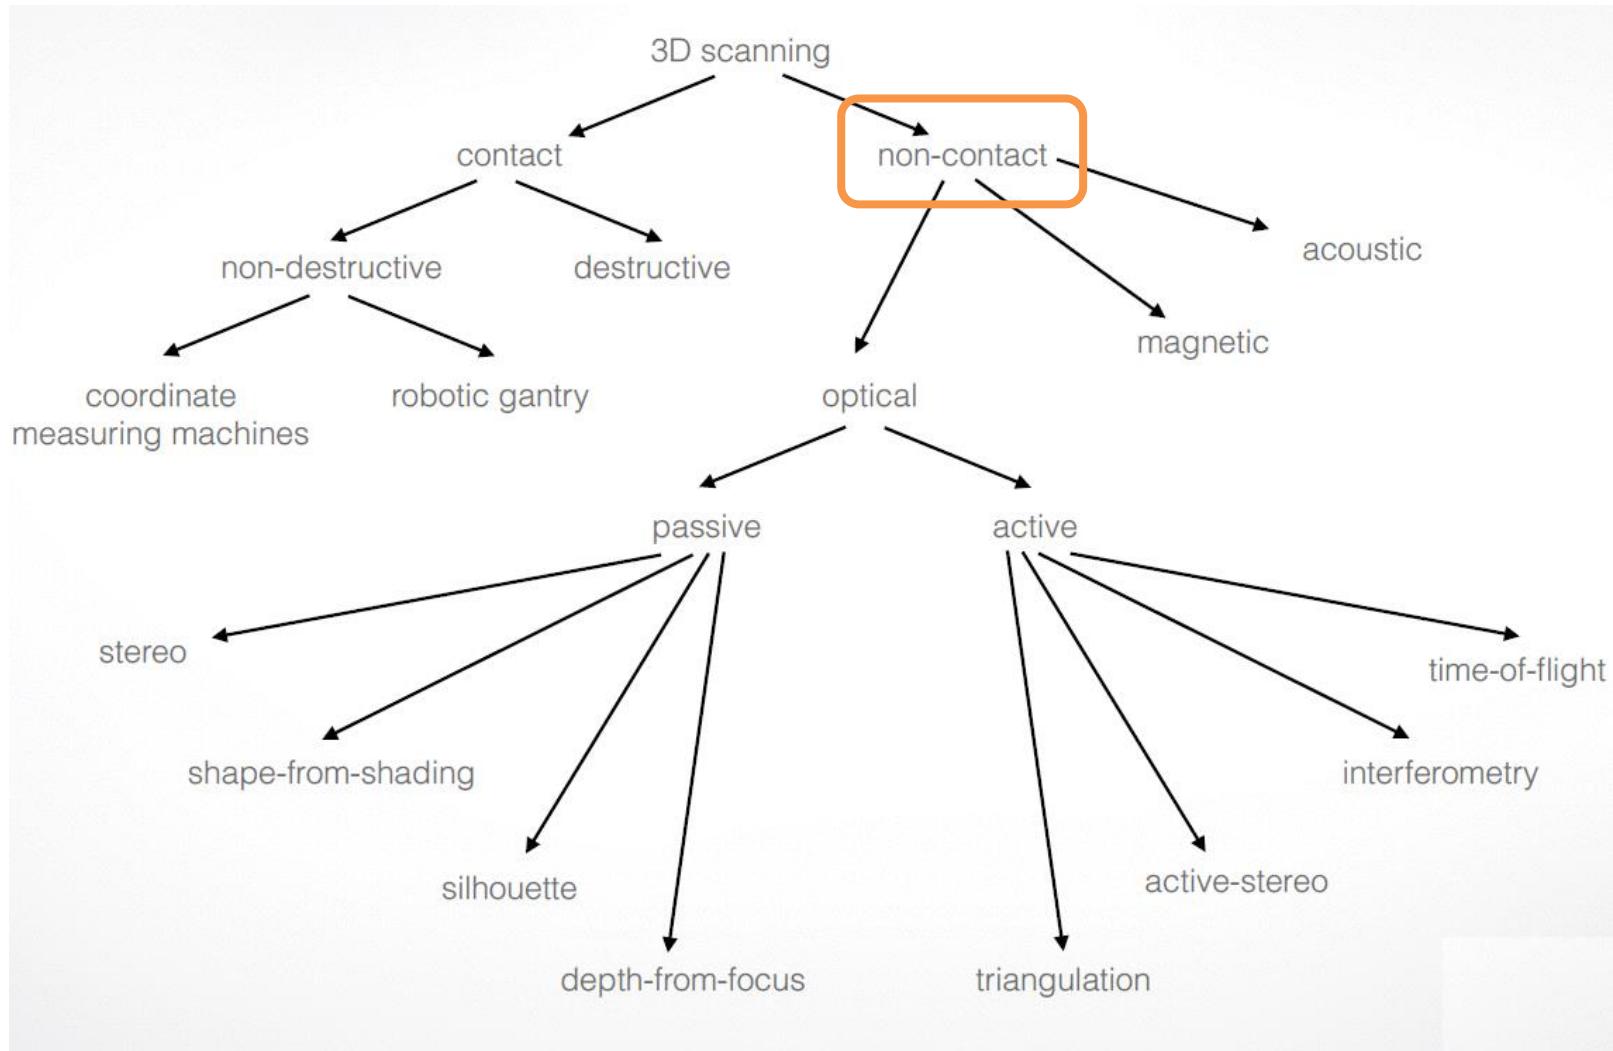
Contact Scanner

Probe object by physical touch!

- hand-held scanners!
- less accurate
- slow scanning, sparse set of samples, it requires experts



A taxonomy



Non-contact Scanner

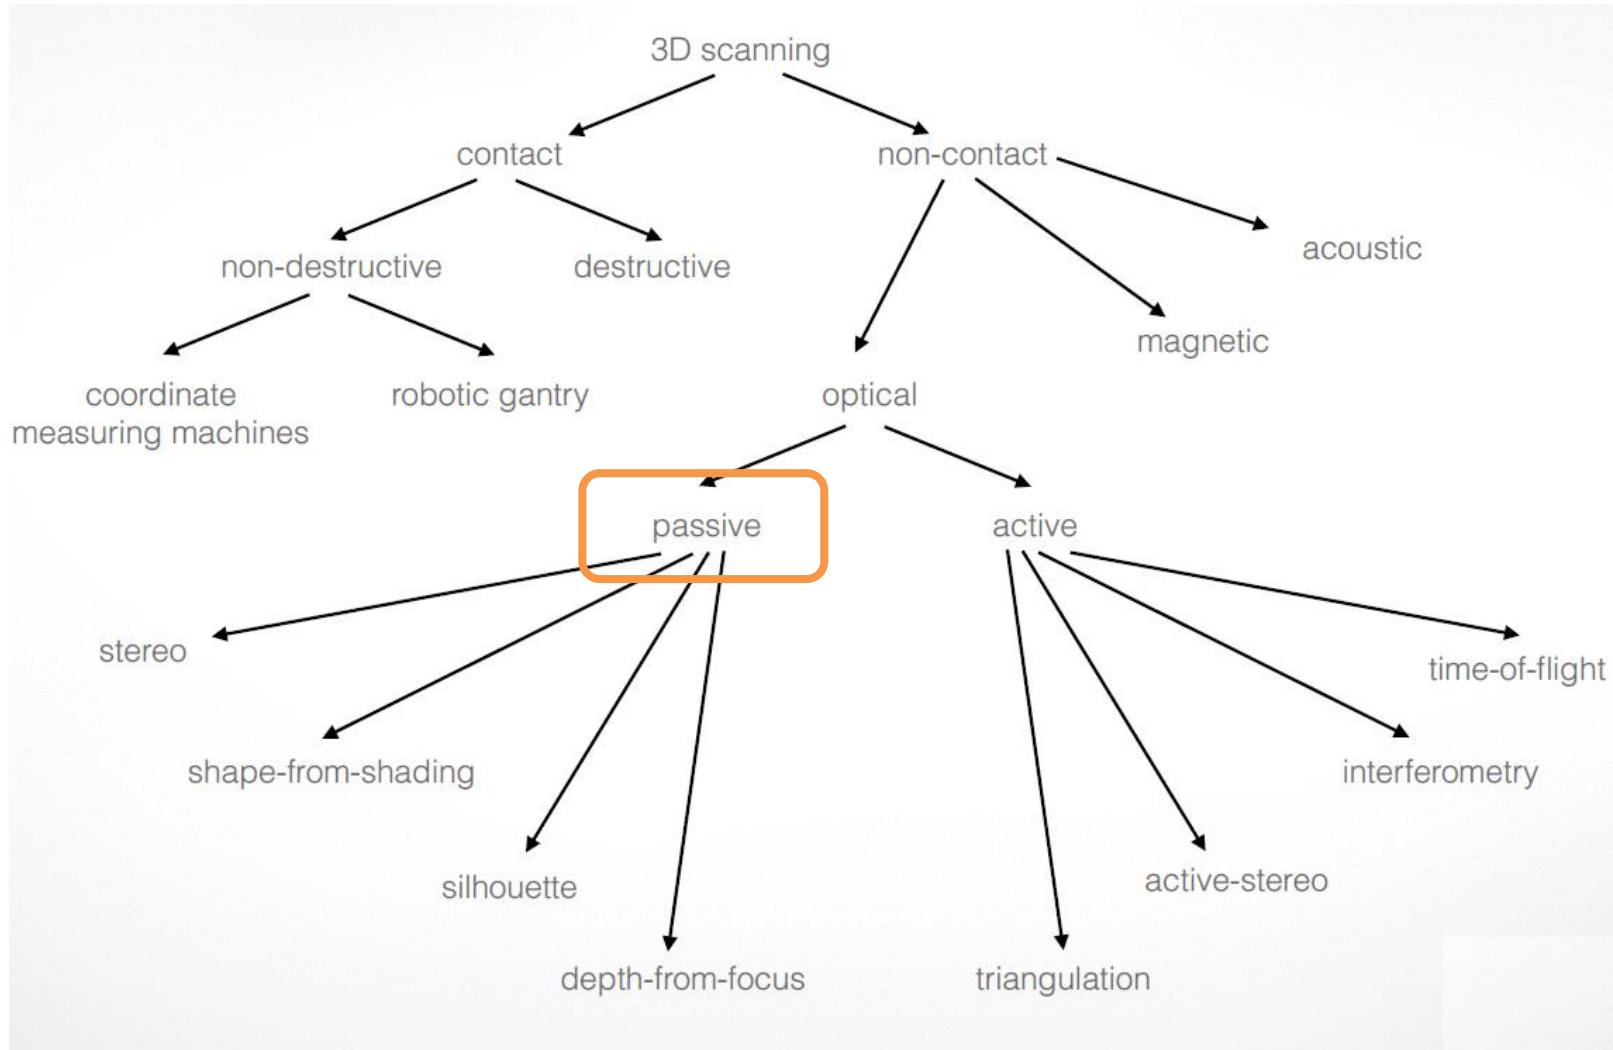
Advantages

- longer and safer distance capture
- potentially faster acquisition
- more automated

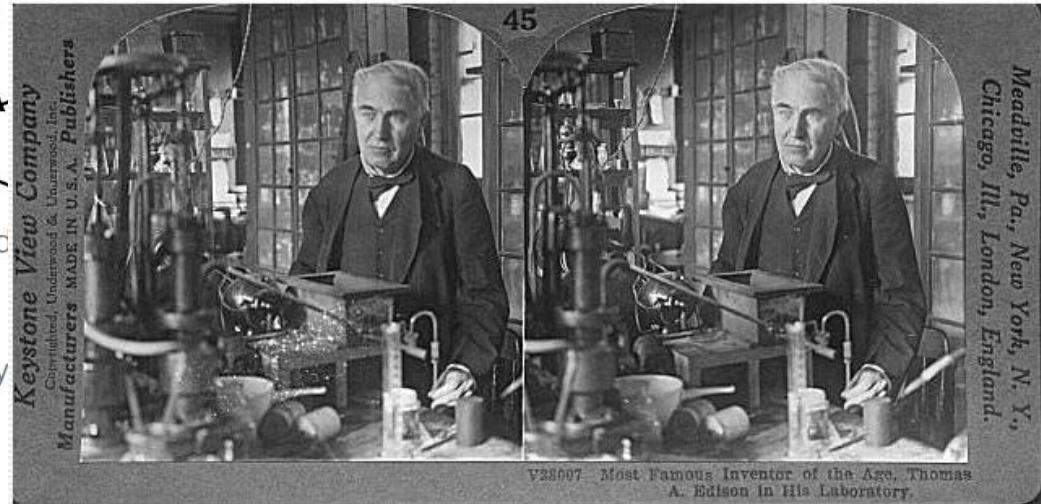
Optical Approaches

- most relevant and used (no special hardware requirements)
- highly flexible
- most accurate
- **passive** and **active** approaches

A taxonomy



A taxonomy



stereo

shape-from-shading

silhouette

depth-from-focus

active-stereo

interferometry

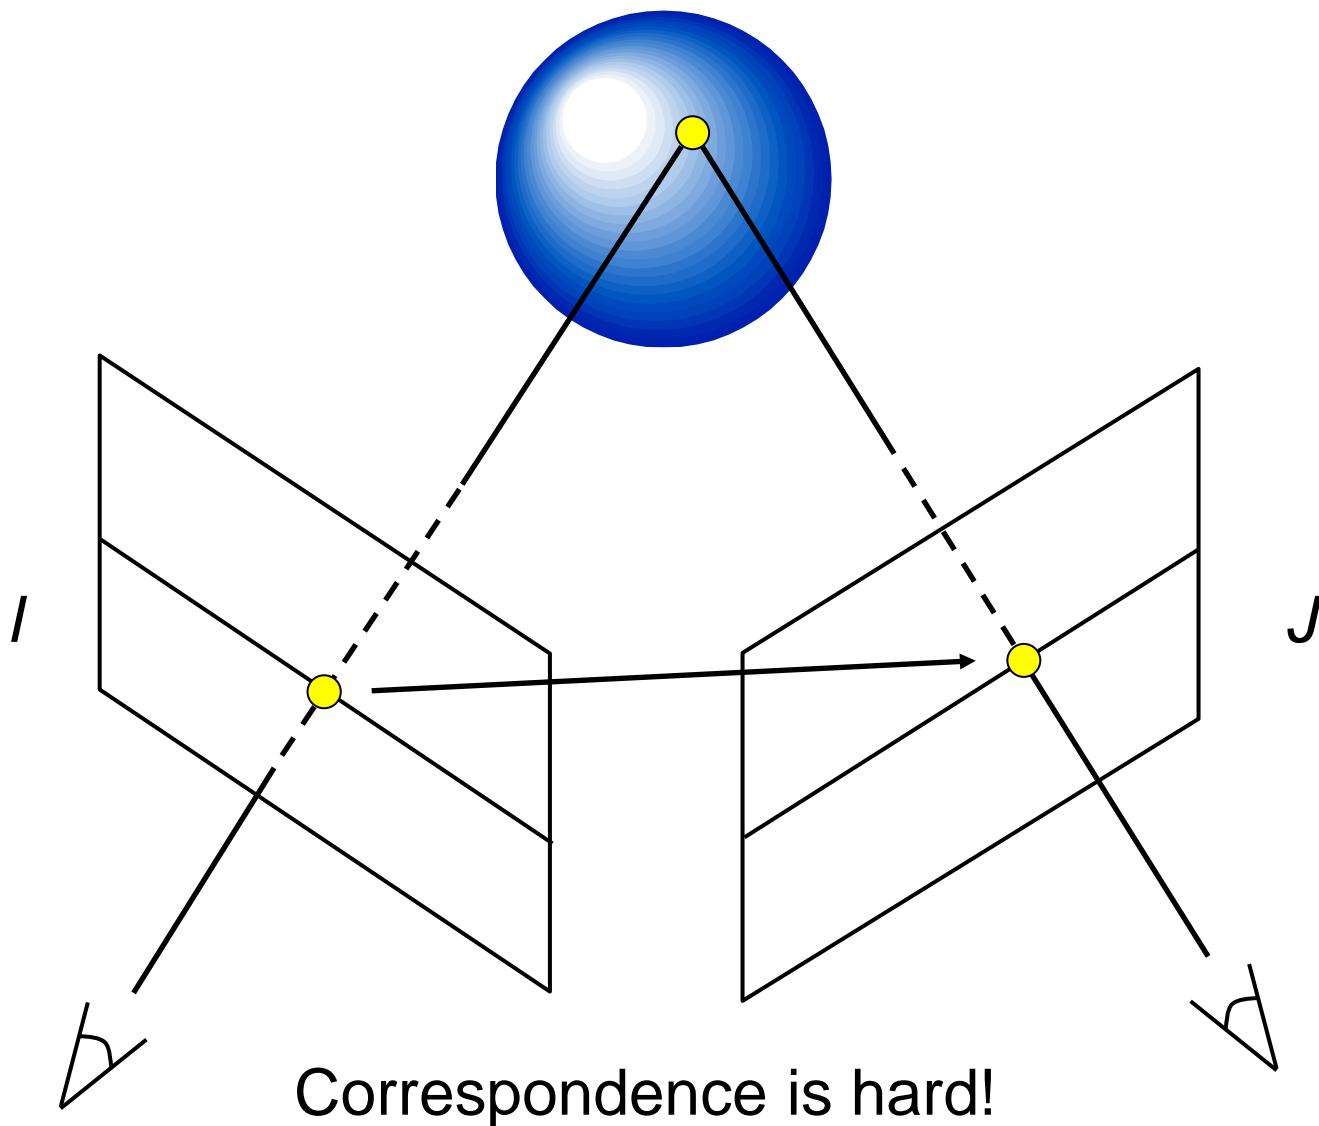
time-of-flight

triangulation

Passive stereo

- Stereo (two views)
- Multiple views
- Structure and motion

Stereo Triangulation



Passive stereo

image 1



image 2

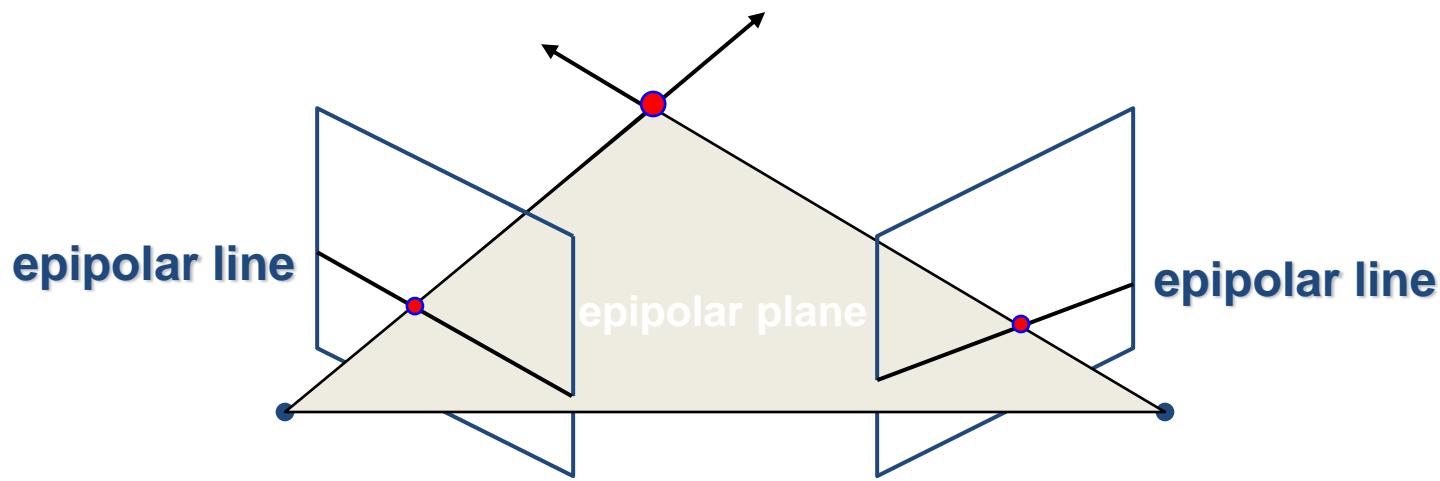


Dense depth (or disparity) map



Stereo correspondence

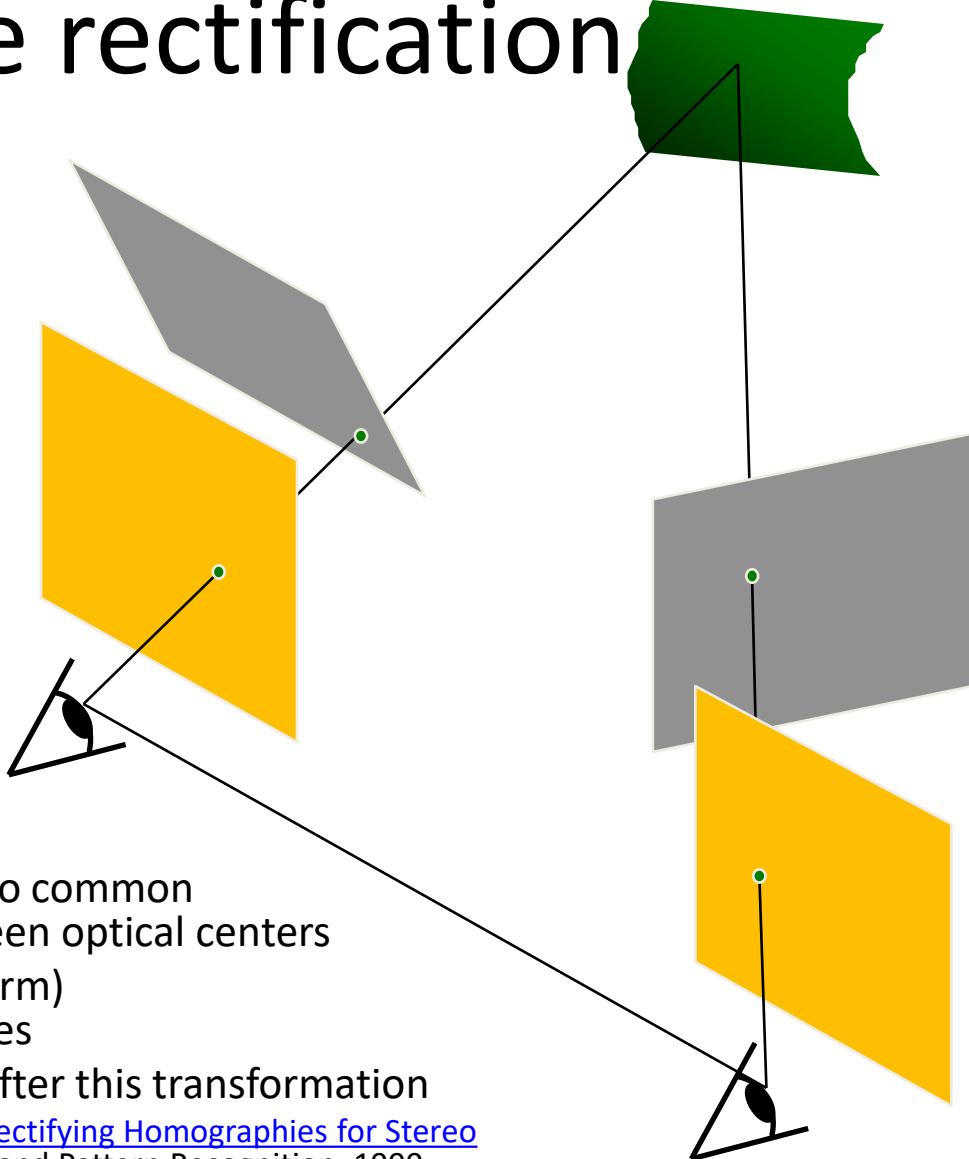
Pairs of points that correspond to same scene point



Epipolar Constraint

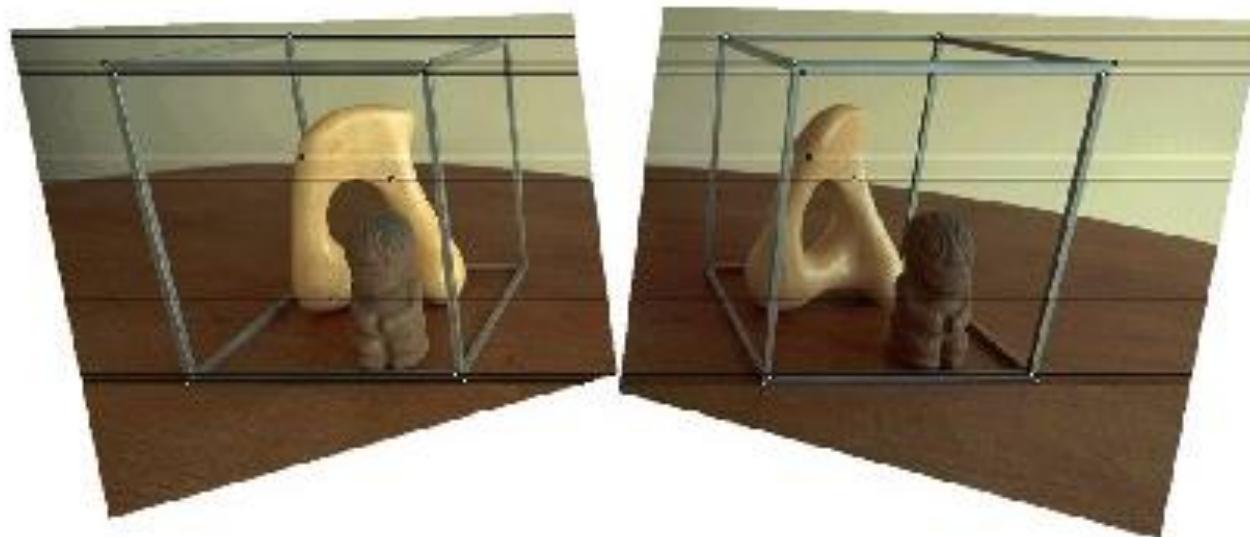
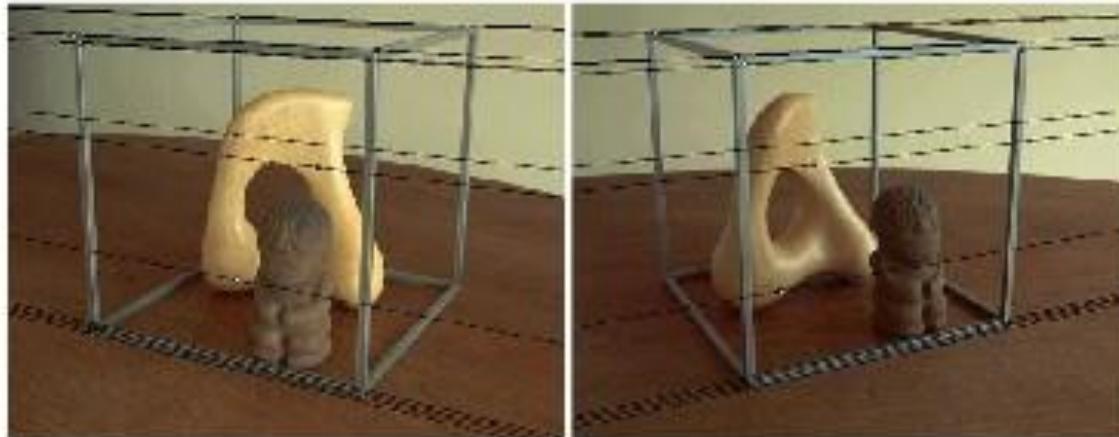
- Reduces correspondence problem to 1D search along *conjugate epipolar lines*

Image rectification

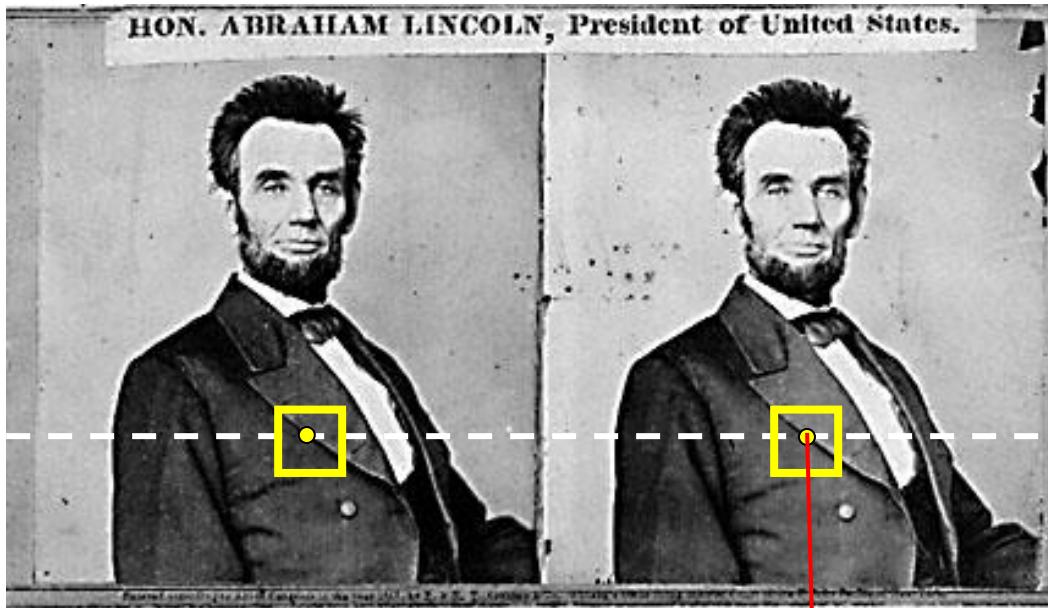


- Image Reprojection
 - reproject image planes onto common plane parallel to line between optical centers
 - a homography (3×3 transform) applied to both input images
 - pixel motion is horizontal after this transformation
 - C. Loop and Z. Zhang. [Computing Rectifying Homographies for Stereo Vision](#). IEEE Conf. Computer Vision and Pattern Recognition, 1999.

Image rectification



Your basic stereo algorithm

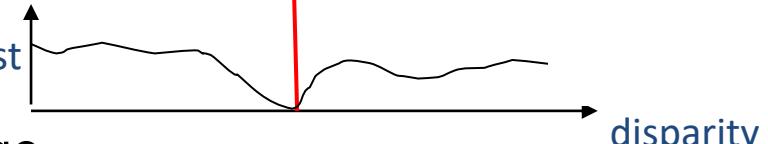


For each epipolar line

Matching cost

For each pixel in the left image

- compare with every pixel on same epipolar line in right image
- pick pixel with minimum match cost



Improvement: match **windows**

Stereo camera



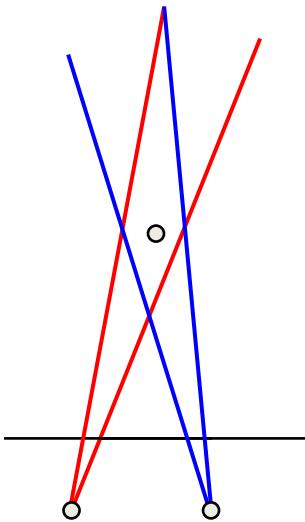
<https://www.stereolabs.com/>

Multiview passive stereo

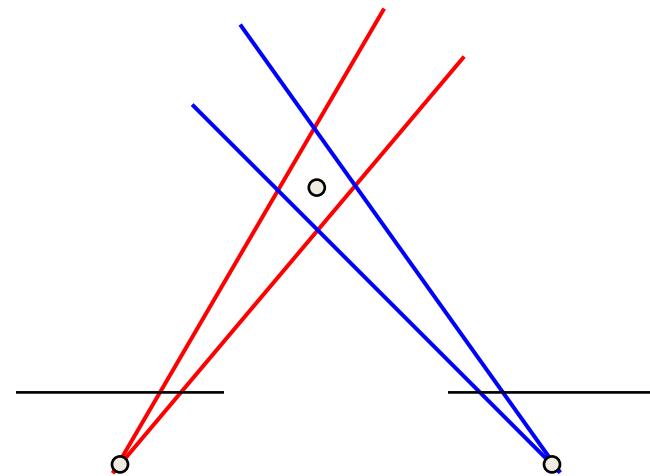


Larger baseline (i.e., distance between cameras)

The role of the baseline



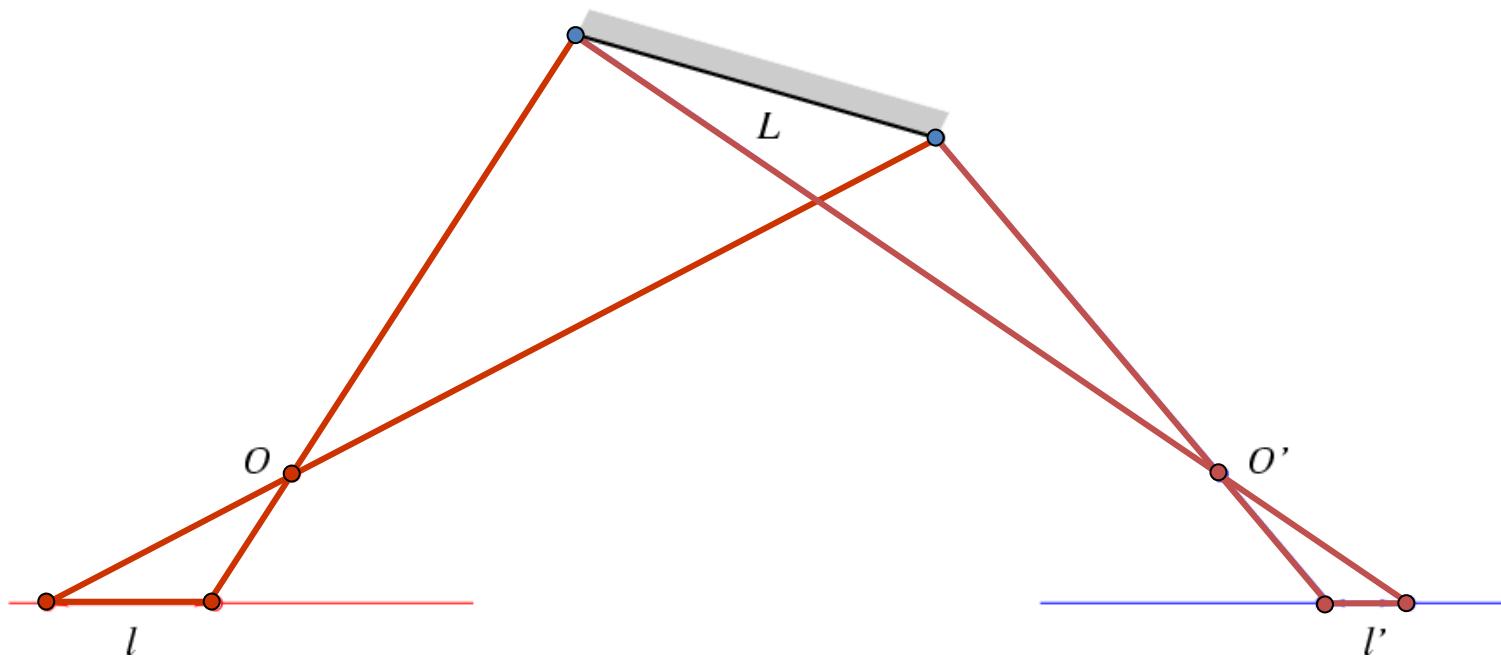
Small Baseline



Large Baseline

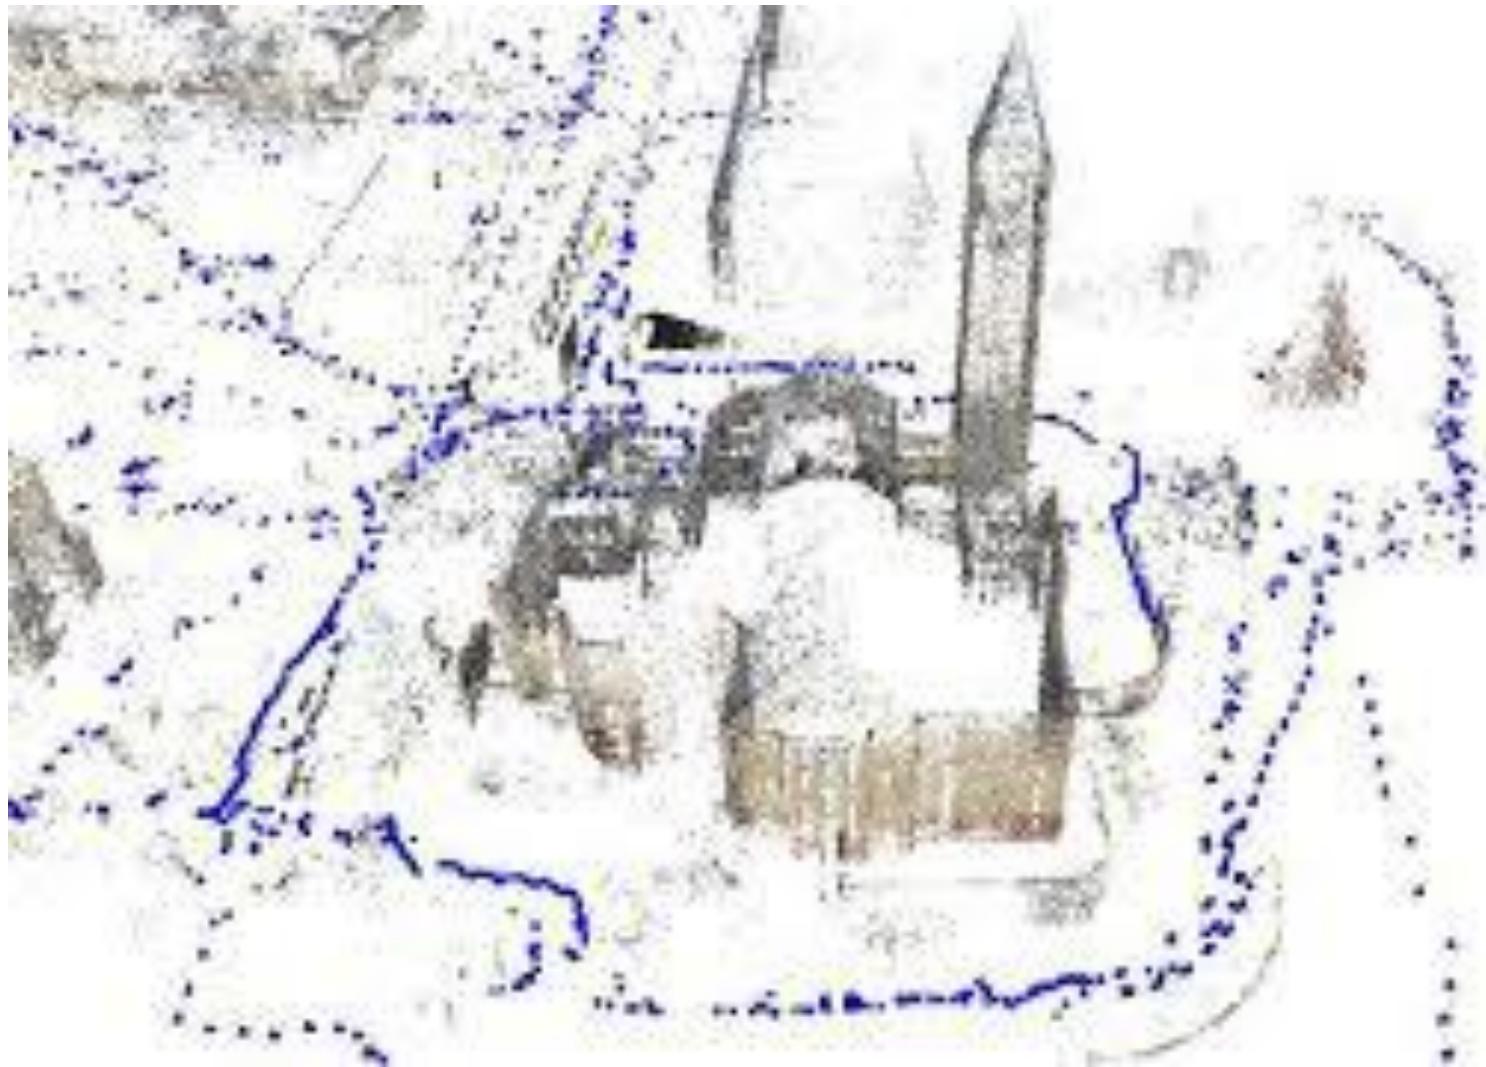
- Small baseline: large depth error
- Large baseline: difficult search problem

Problem for wide baselines: Foreshortening

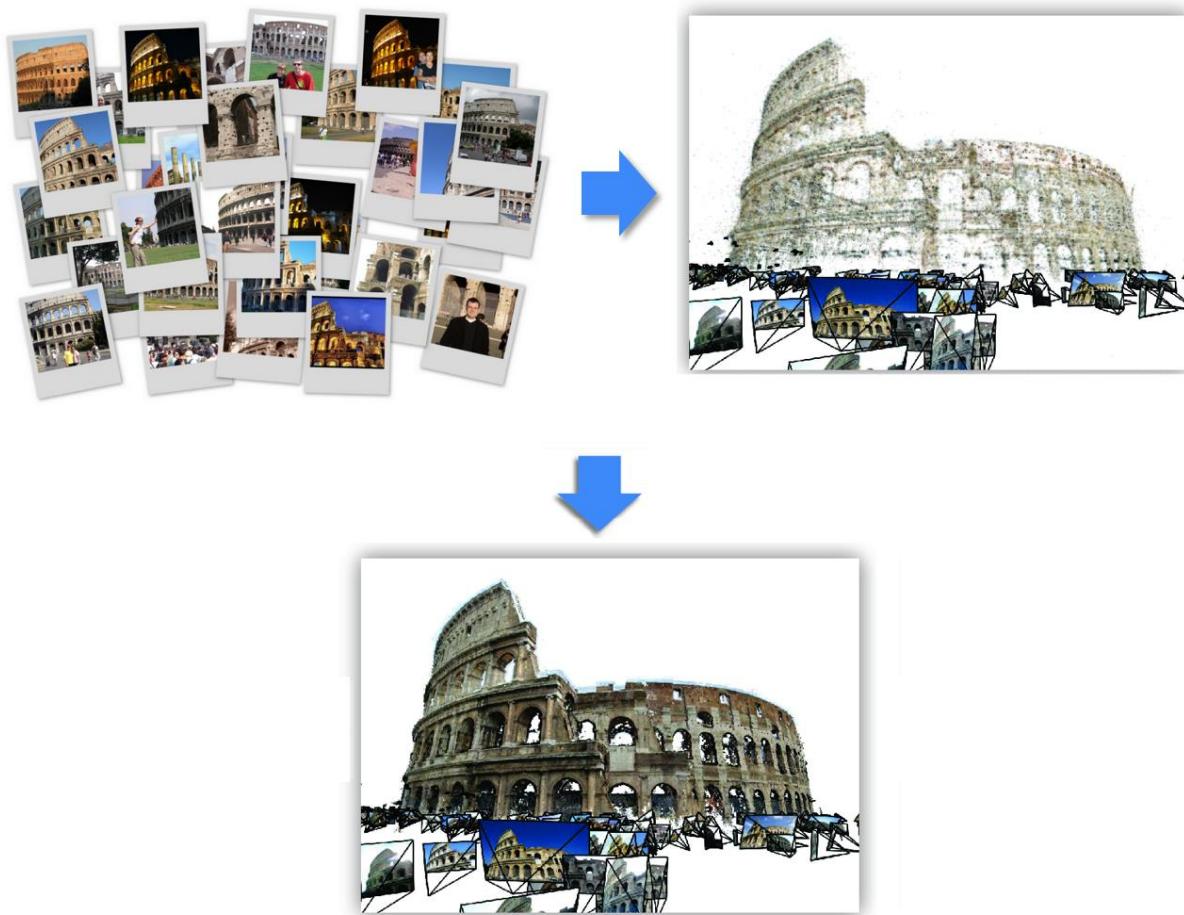


- Matching with fixed-size windows will fail!
- Possible solution: adaptively vary window size

Structure and motion (S&M)

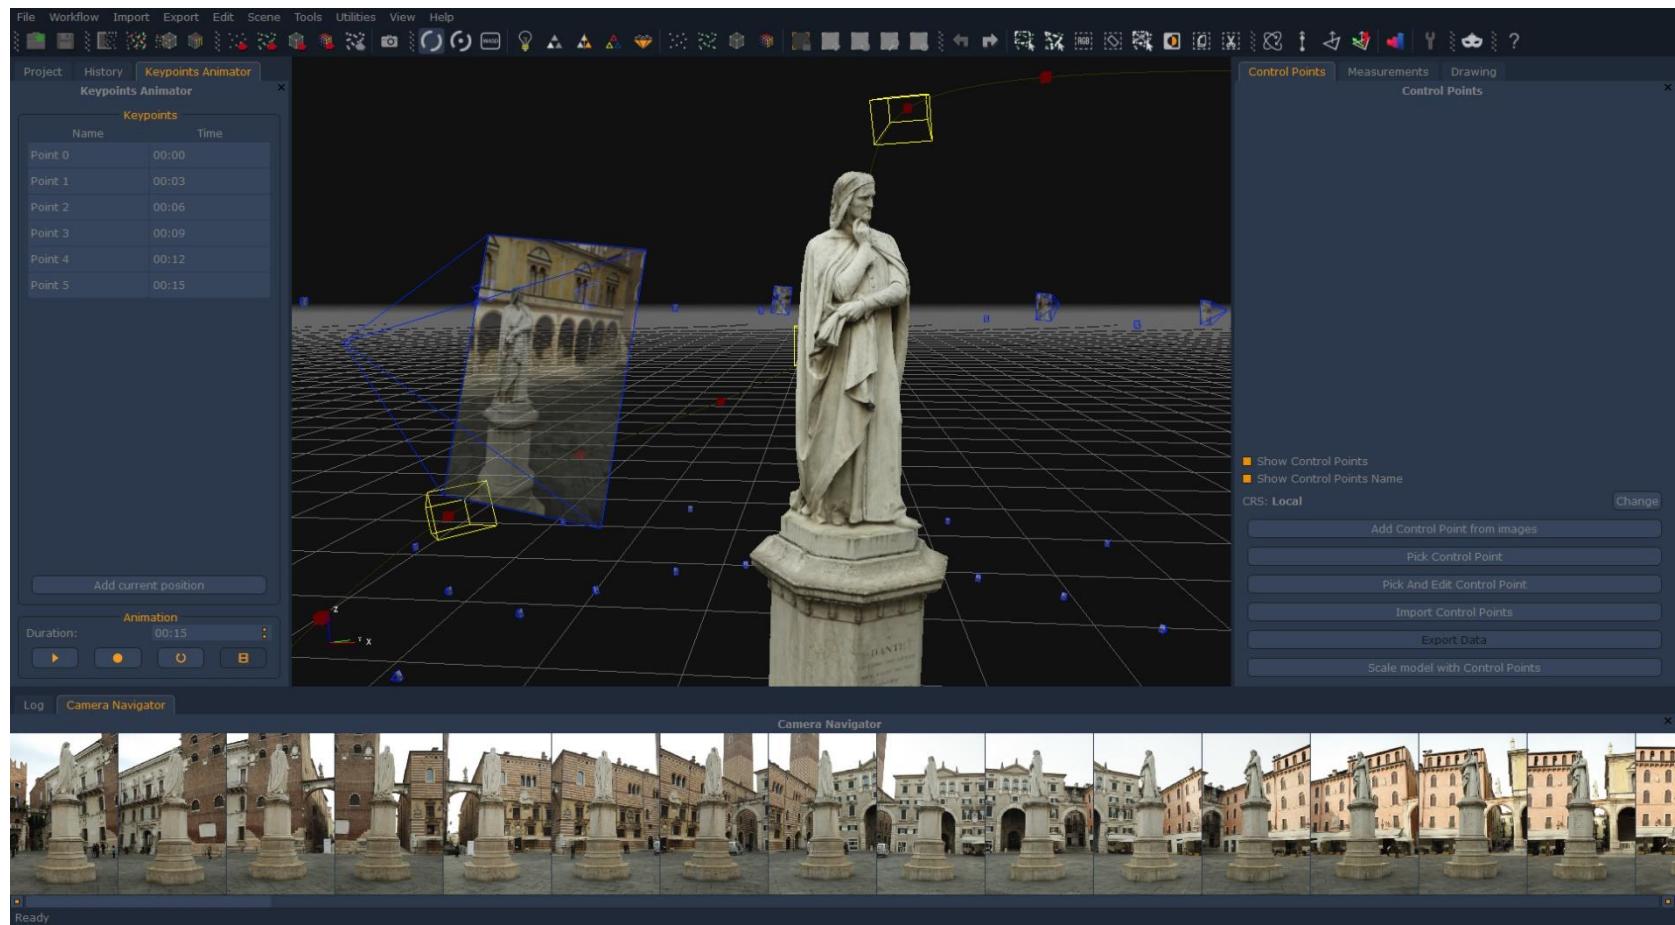


Large scale S&M



Building Rome in a day: <https://grail.cs.washington.edu/rome/>

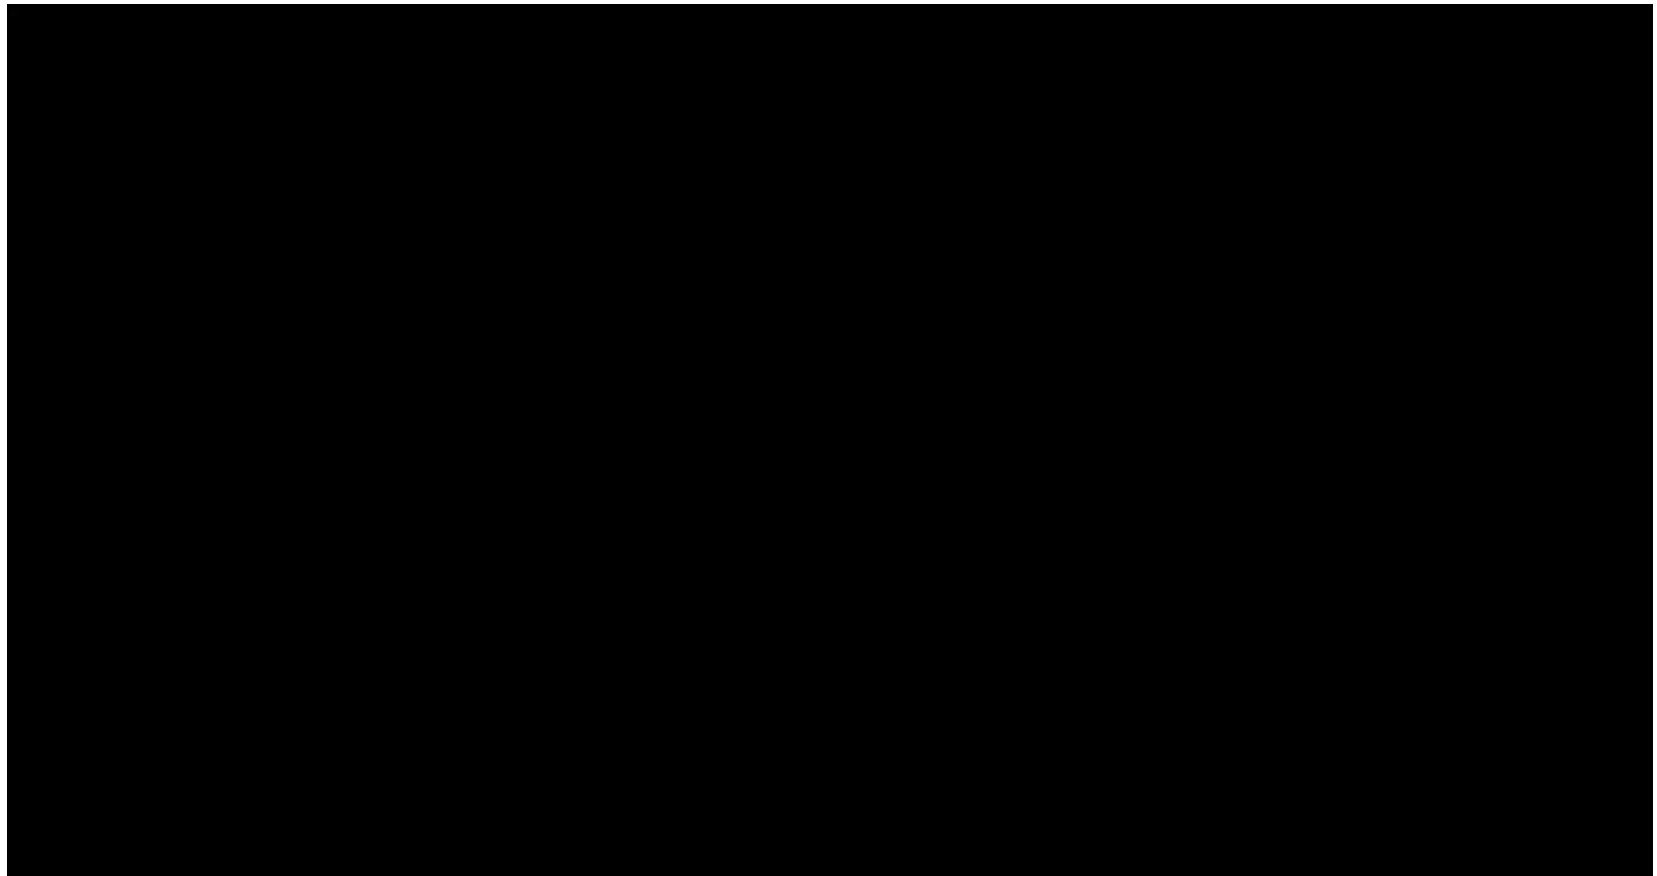
Structure and motion



3DF Zephyr Pro



Structure and motion

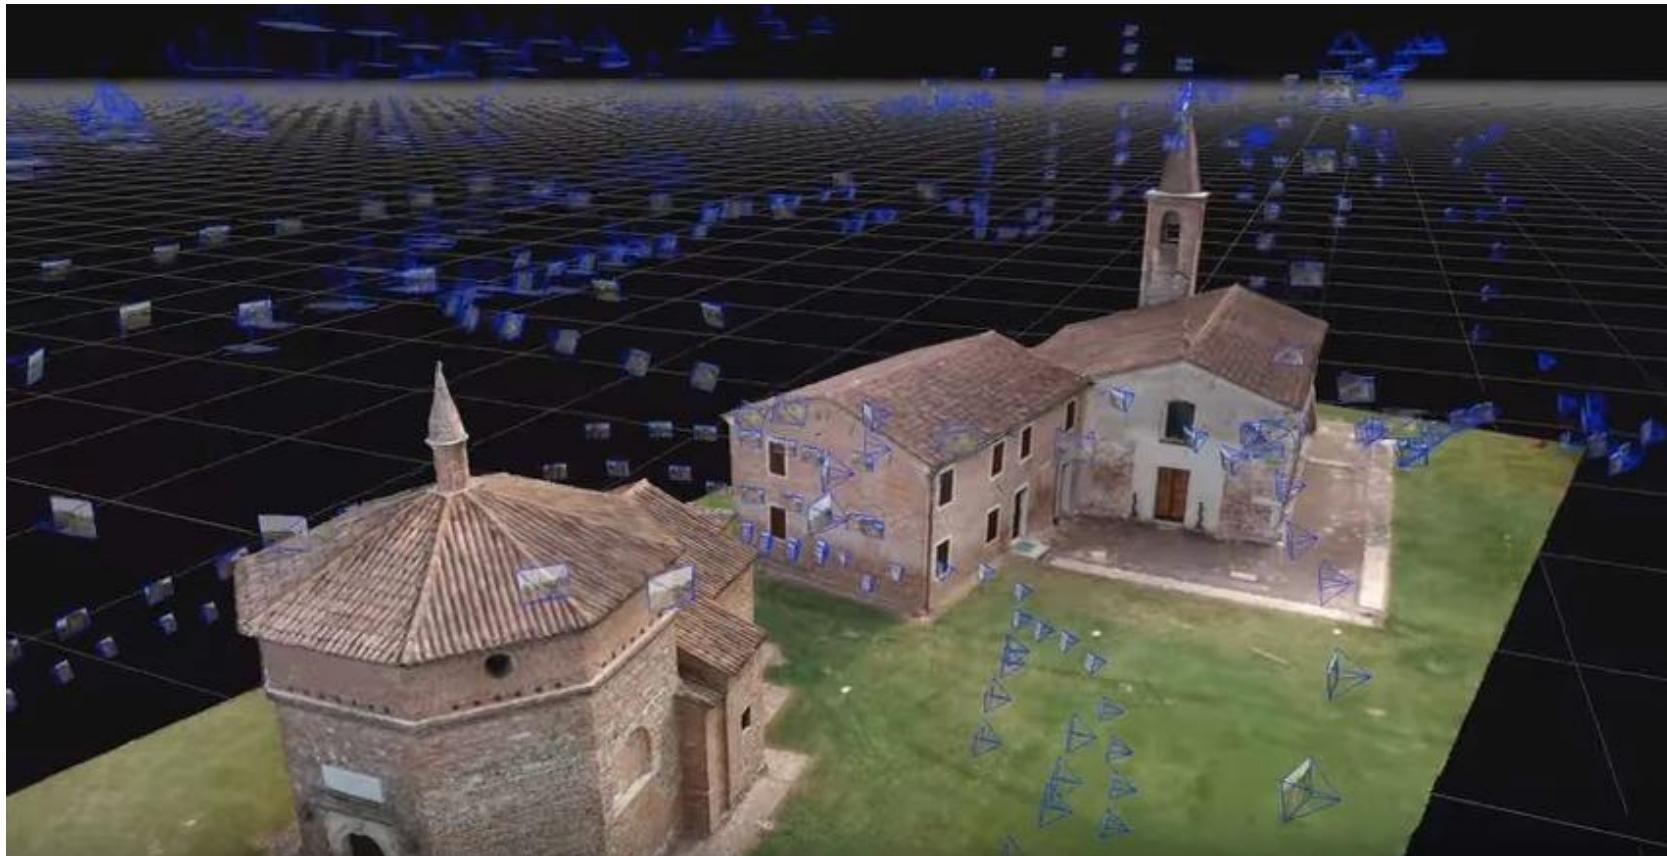


<https://www.3dflow.net/3df-zephyr-photogrammetry-software/>

3DF Zephyr Pro



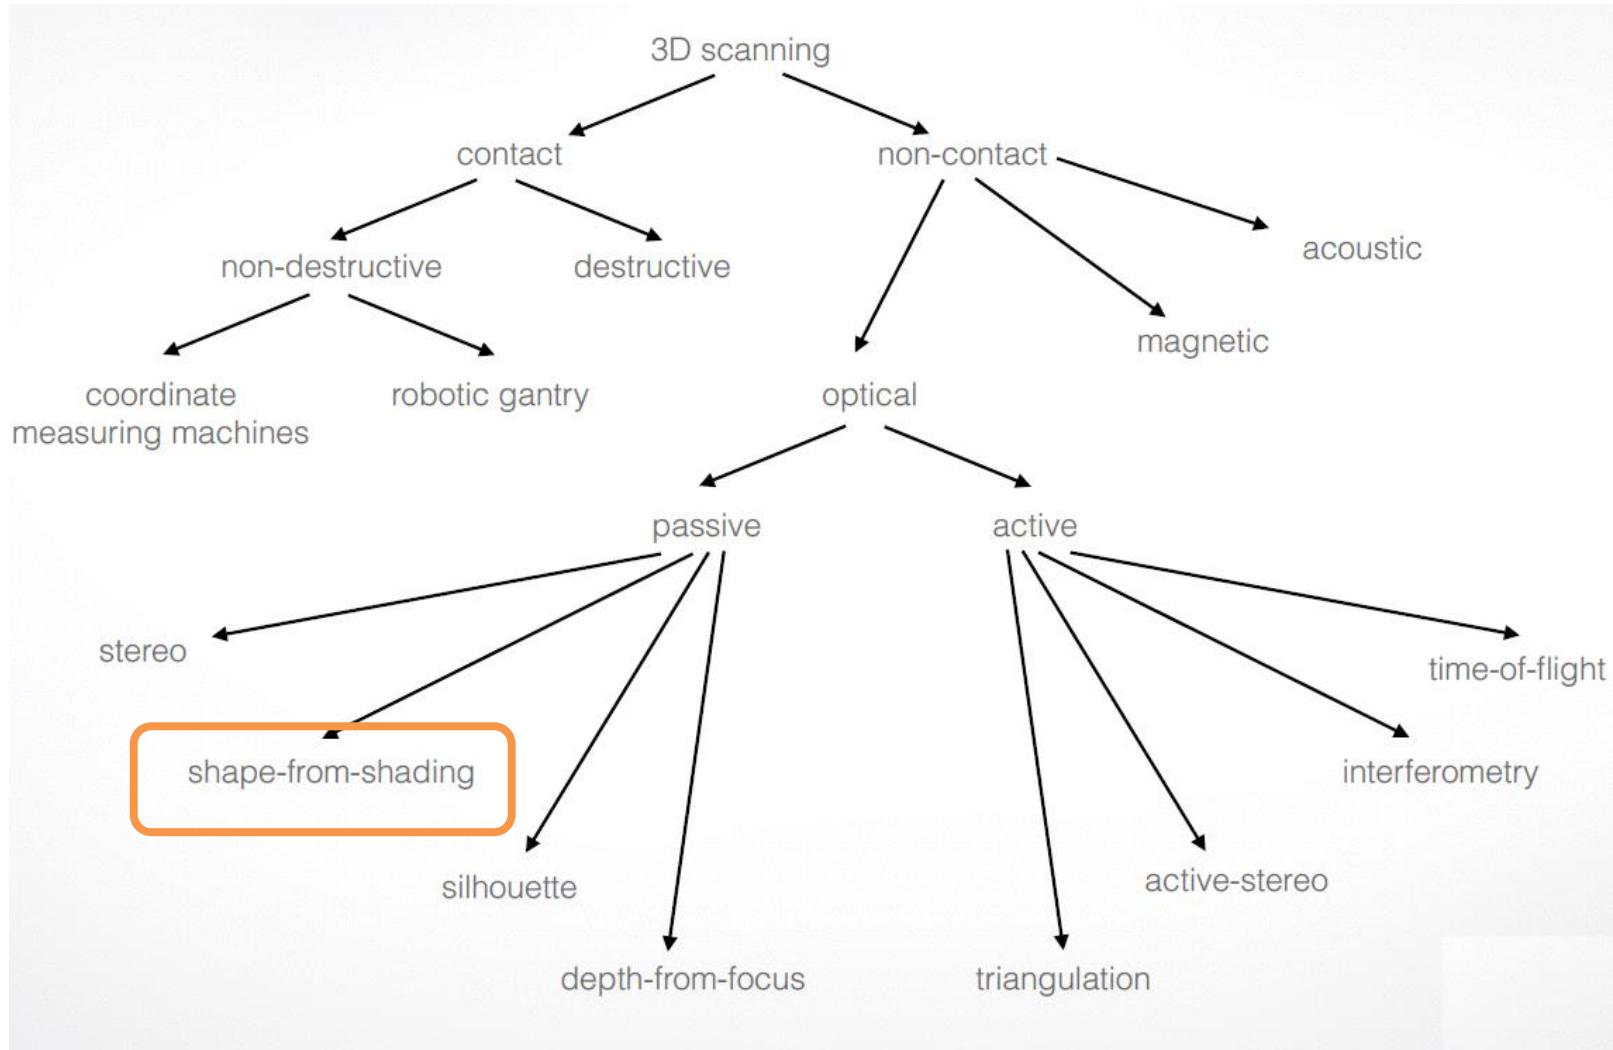
From aerial images



3DF Zephyr Pro



A taxonomy



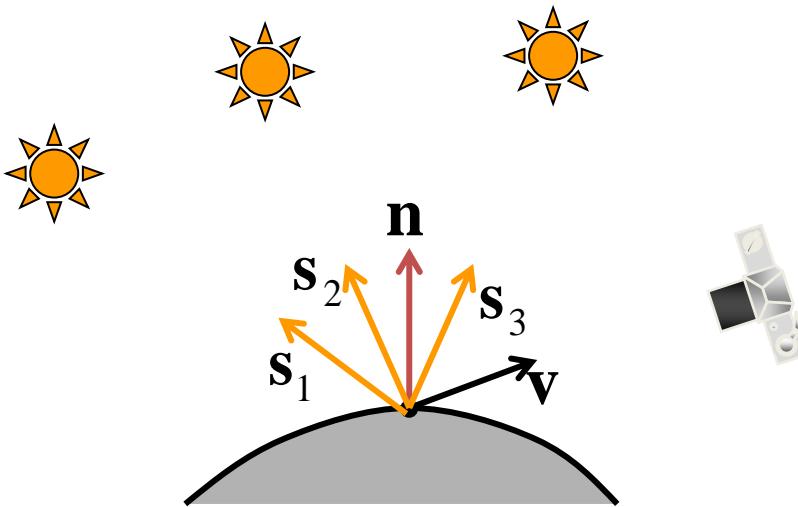
Shape from shading

- Image with a controlled illumination set-up
 - Not a proper ‘passive’ system...,
- Only surface normals are estimated,
- Better performance when more images with different light sources are used



Photometric stereo!

Photometric Stereo



Lambertian case:

$$I = \frac{\rho}{\pi} k c \cos \theta_i = \rho \mathbf{n} \cdot \mathbf{s} \quad \left(\frac{kc}{\pi} = 1 \right)$$

Image irradiance:

$$\begin{aligned} I_1 &= \rho \mathbf{n} \cdot \mathbf{s}_1 \\ I_2 &= \rho \mathbf{n} \cdot \mathbf{s}_2 \\ I_3 &= \rho \mathbf{n} \cdot \mathbf{s}_3 \end{aligned}$$

- We can write this in matrix form:

$$\begin{bmatrix} I_1 \\ I_2 \\ I_3 \end{bmatrix} = \rho \begin{bmatrix} \mathbf{s}_1^T \\ \mathbf{s}_2^T \\ \mathbf{s}_3^T \end{bmatrix} \mathbf{n}$$

Photometric Stereo

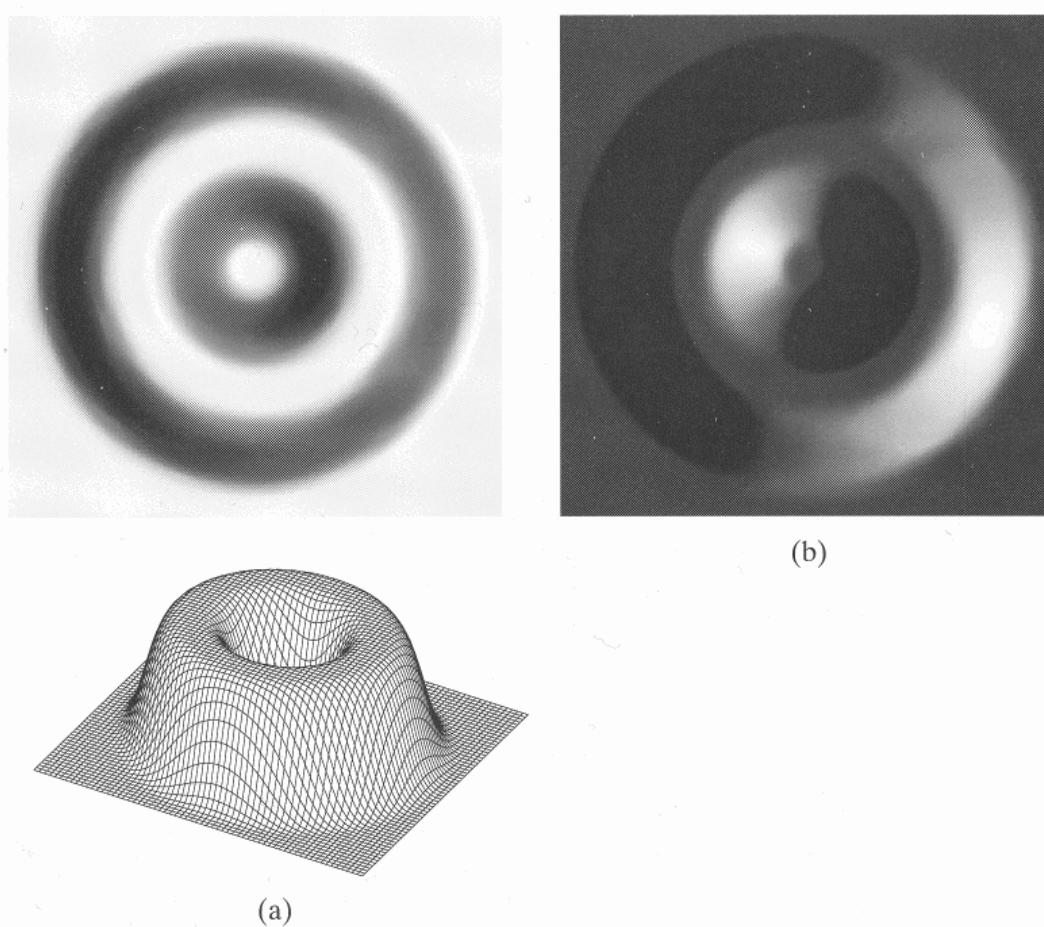
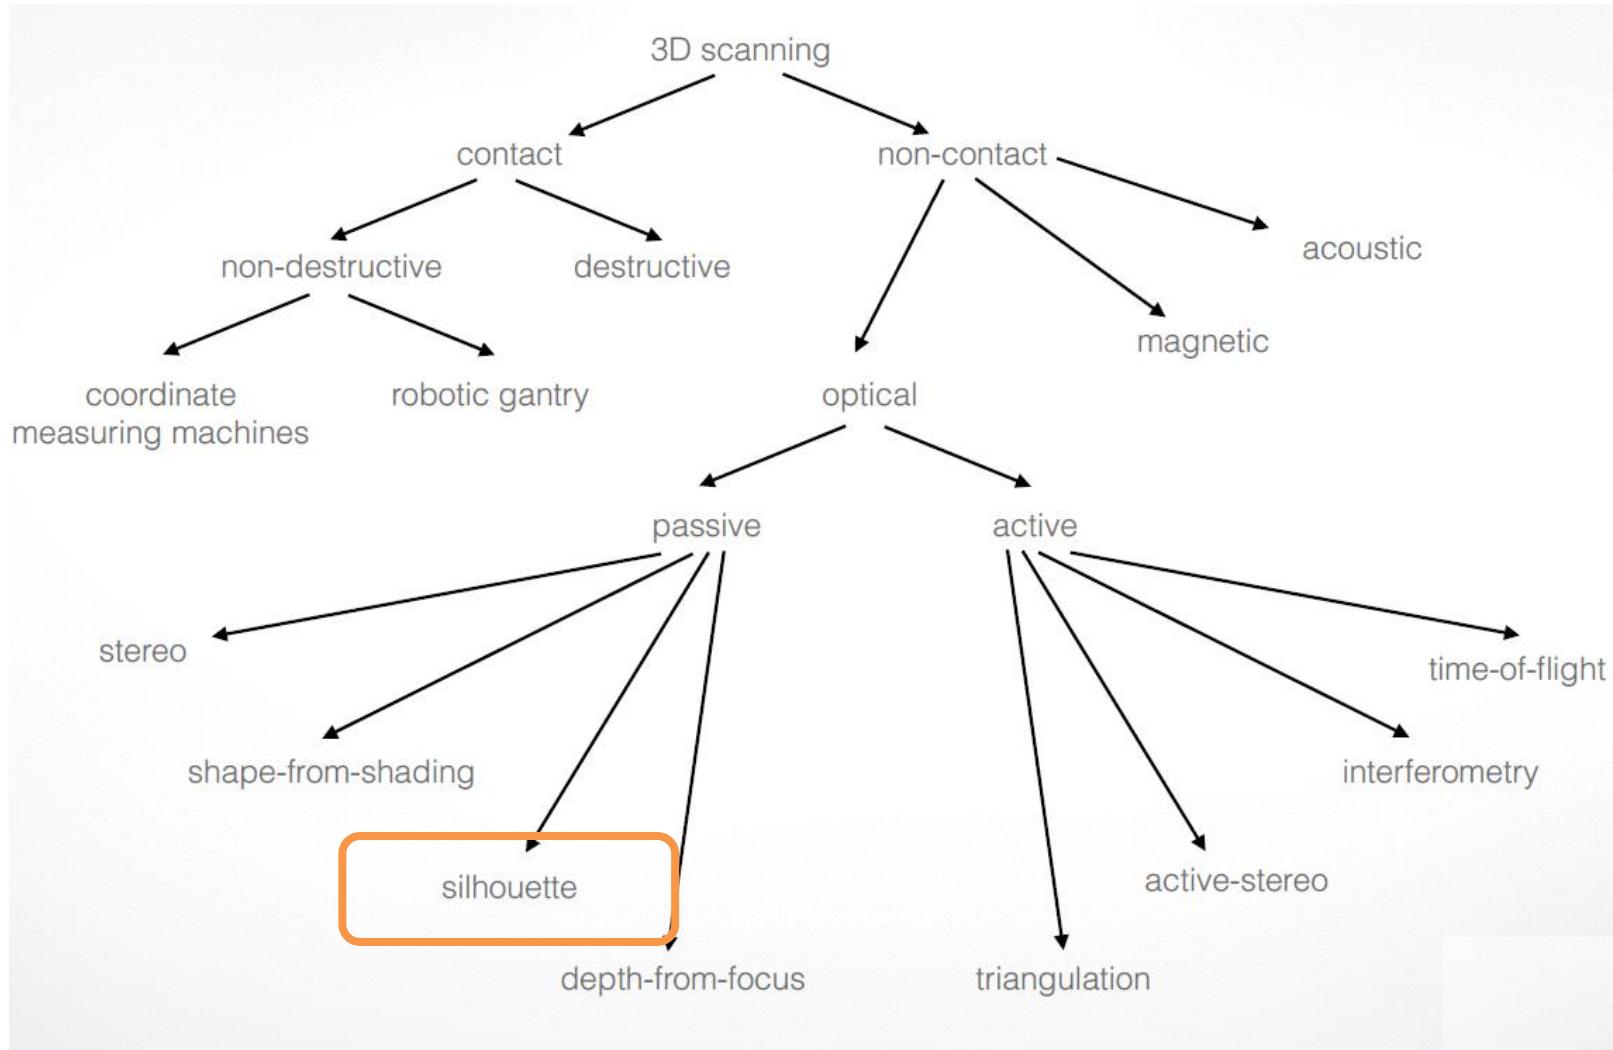
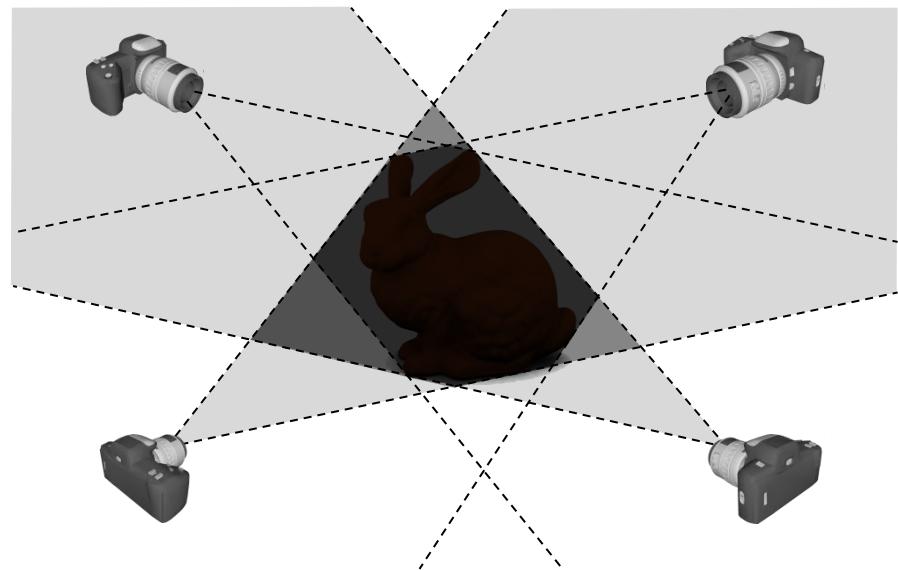


Figure 9.2 Two images of the same Lambertian surface seen from above but illuminated from different directions and 3-D rendering of the surface. Practically all the points in the top left image receive direct illumination ($\mathbf{i} = [0.20, 0, 0.98]^\top$); some regions of the top right image are in the dark due to self-shadowing effects ($\mathbf{i} = [0.94, 0.31, 0.16]^\top$).

A taxonomy

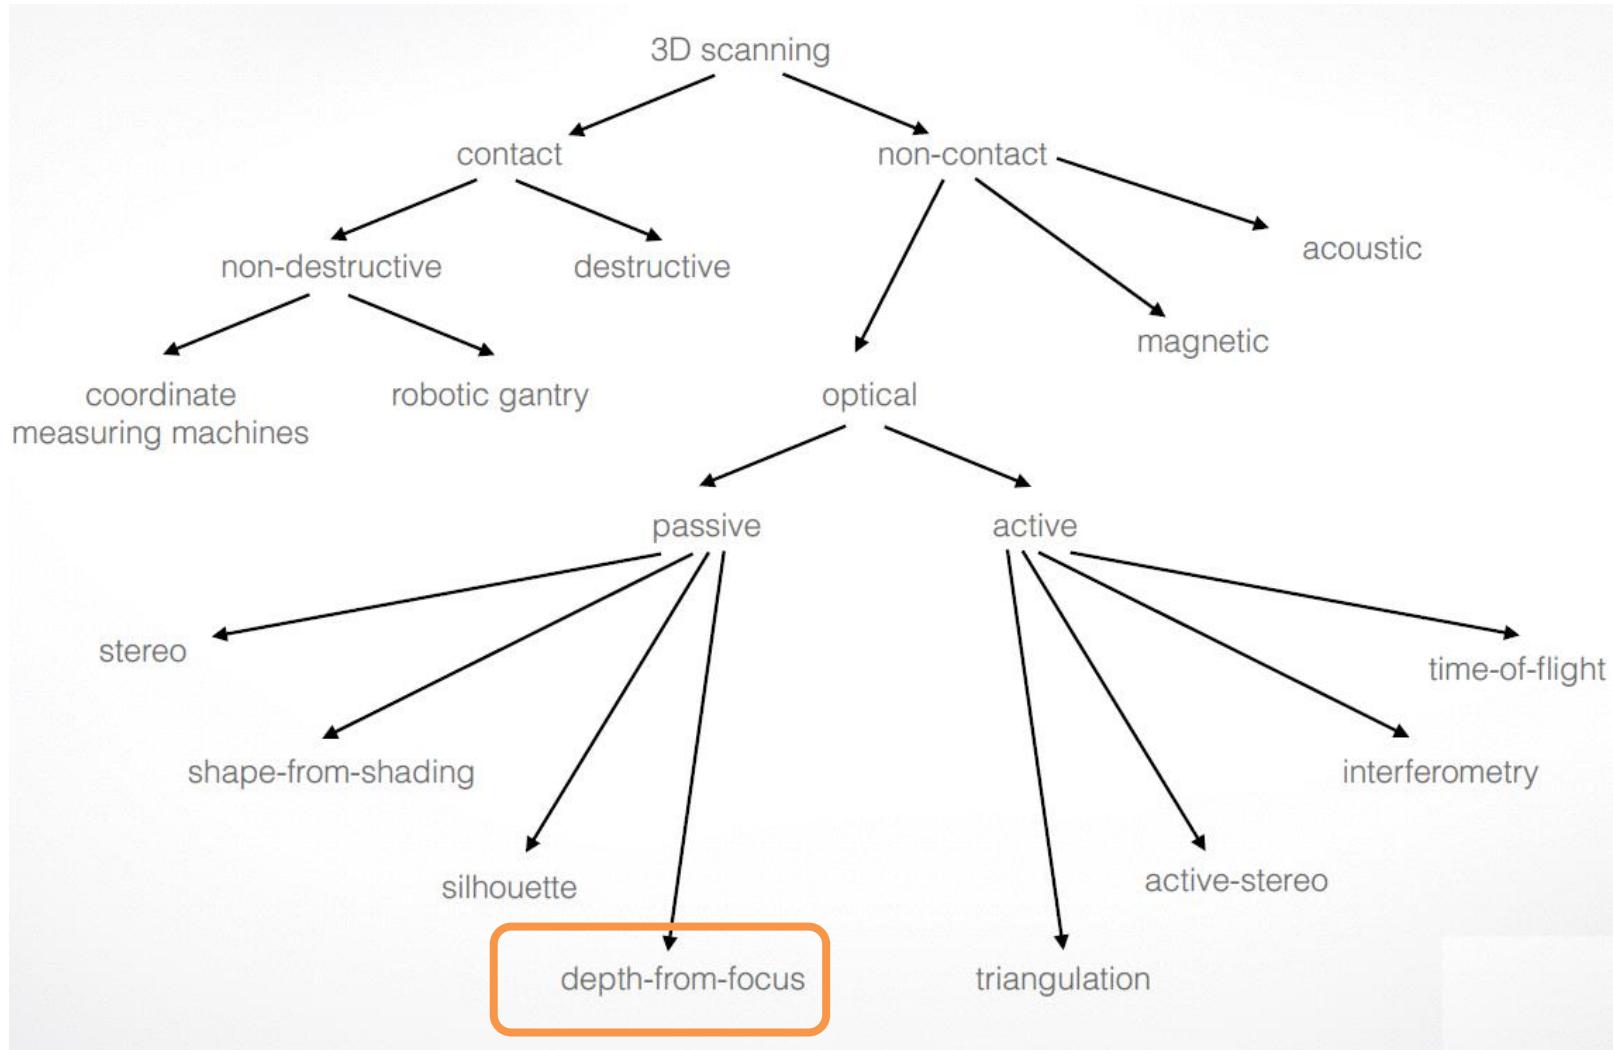


Shape from Silhouette

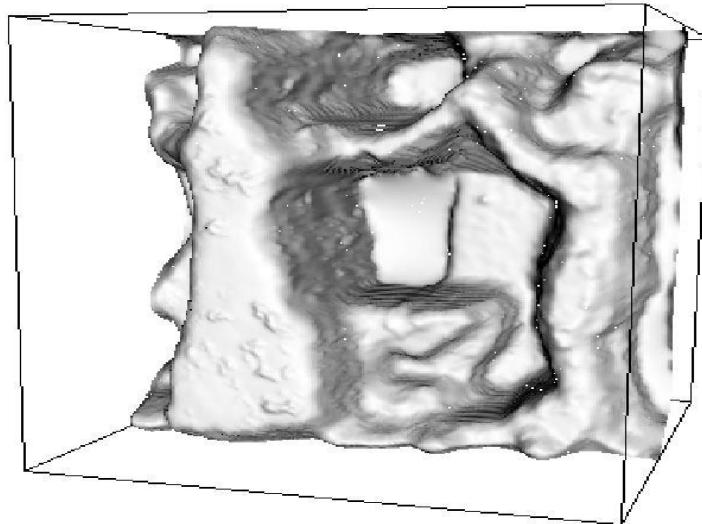
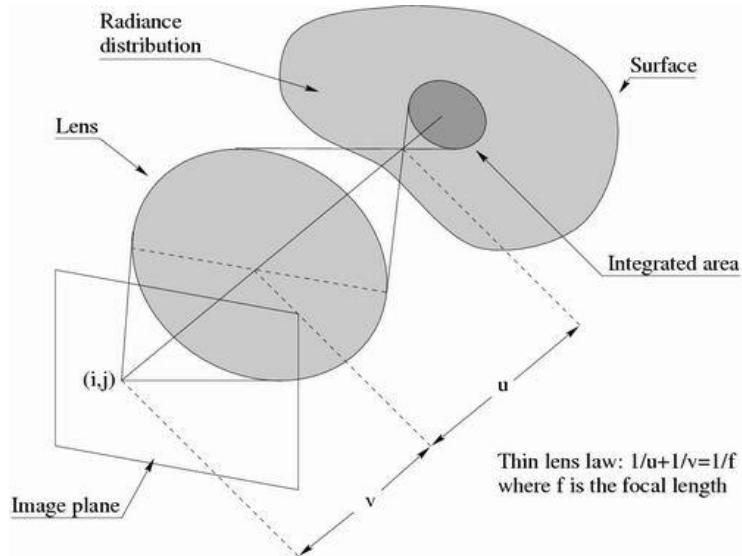


J. Starck and A. Hilton. *Surface Capture for Performance-Based Animation*. *IEEE Computer Graphics and Applications*, 2007

A taxonomy

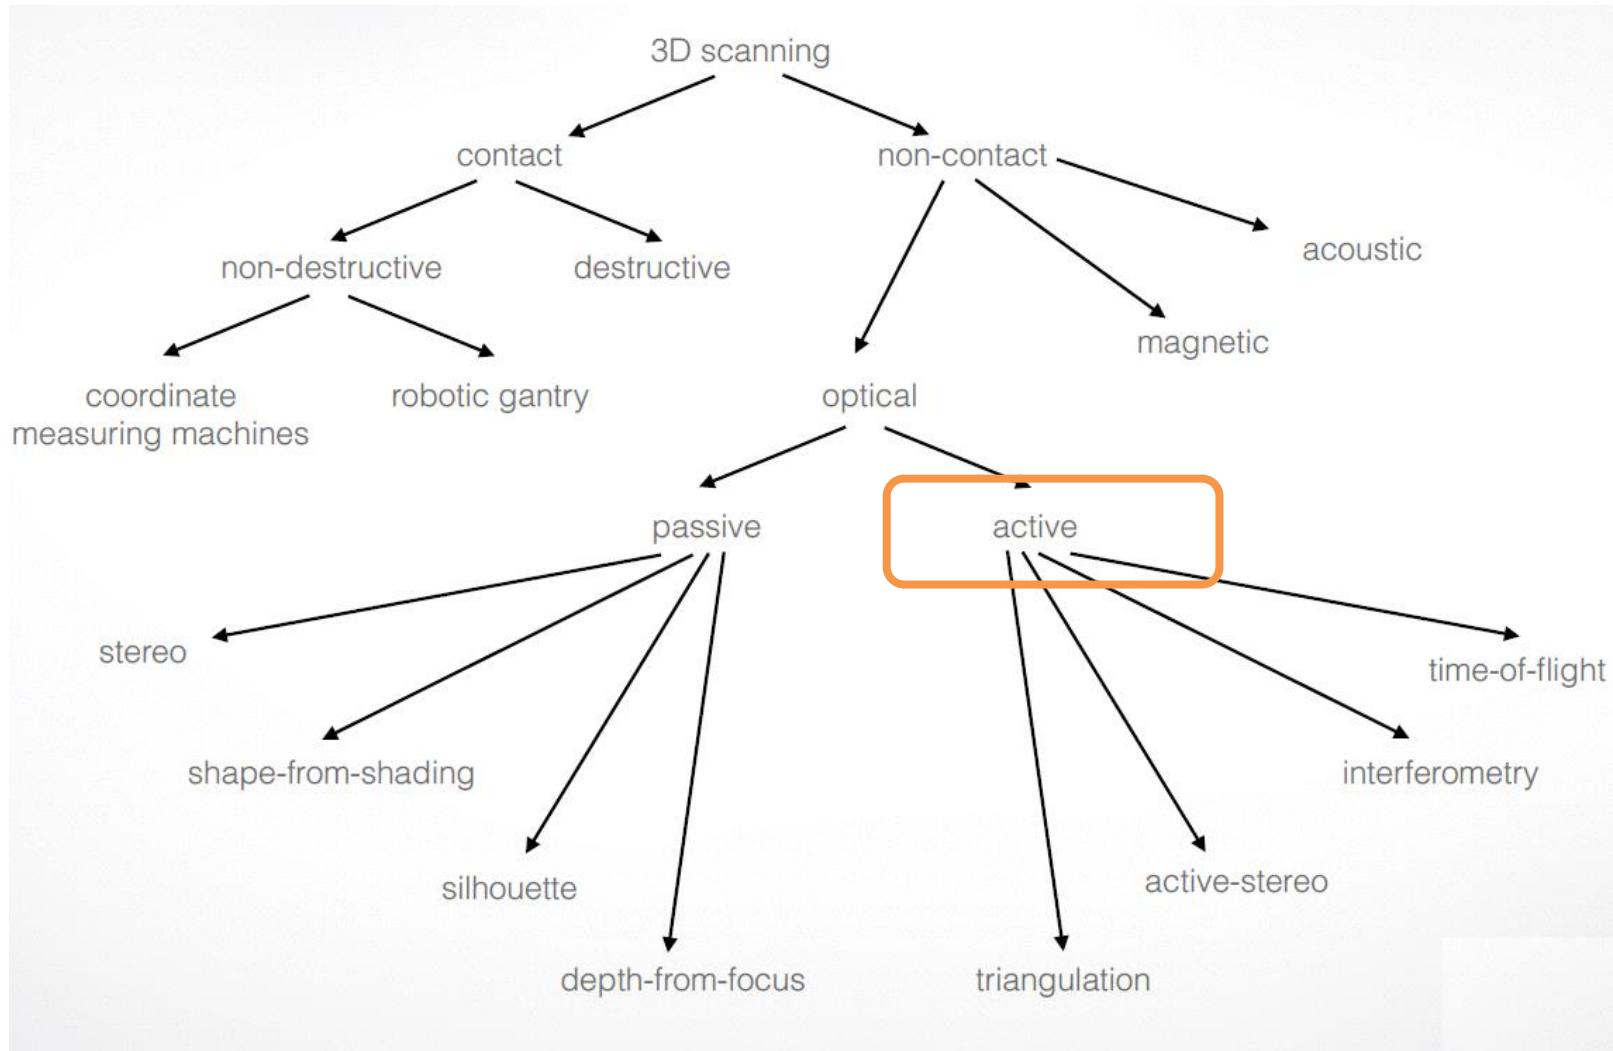


Depth from focus



M. Watanabe and S. Nayar. [Rational filters for passive depth from defocus](#). *Intl. J. of Comp. Vision*, 27(3):203-225, 1998

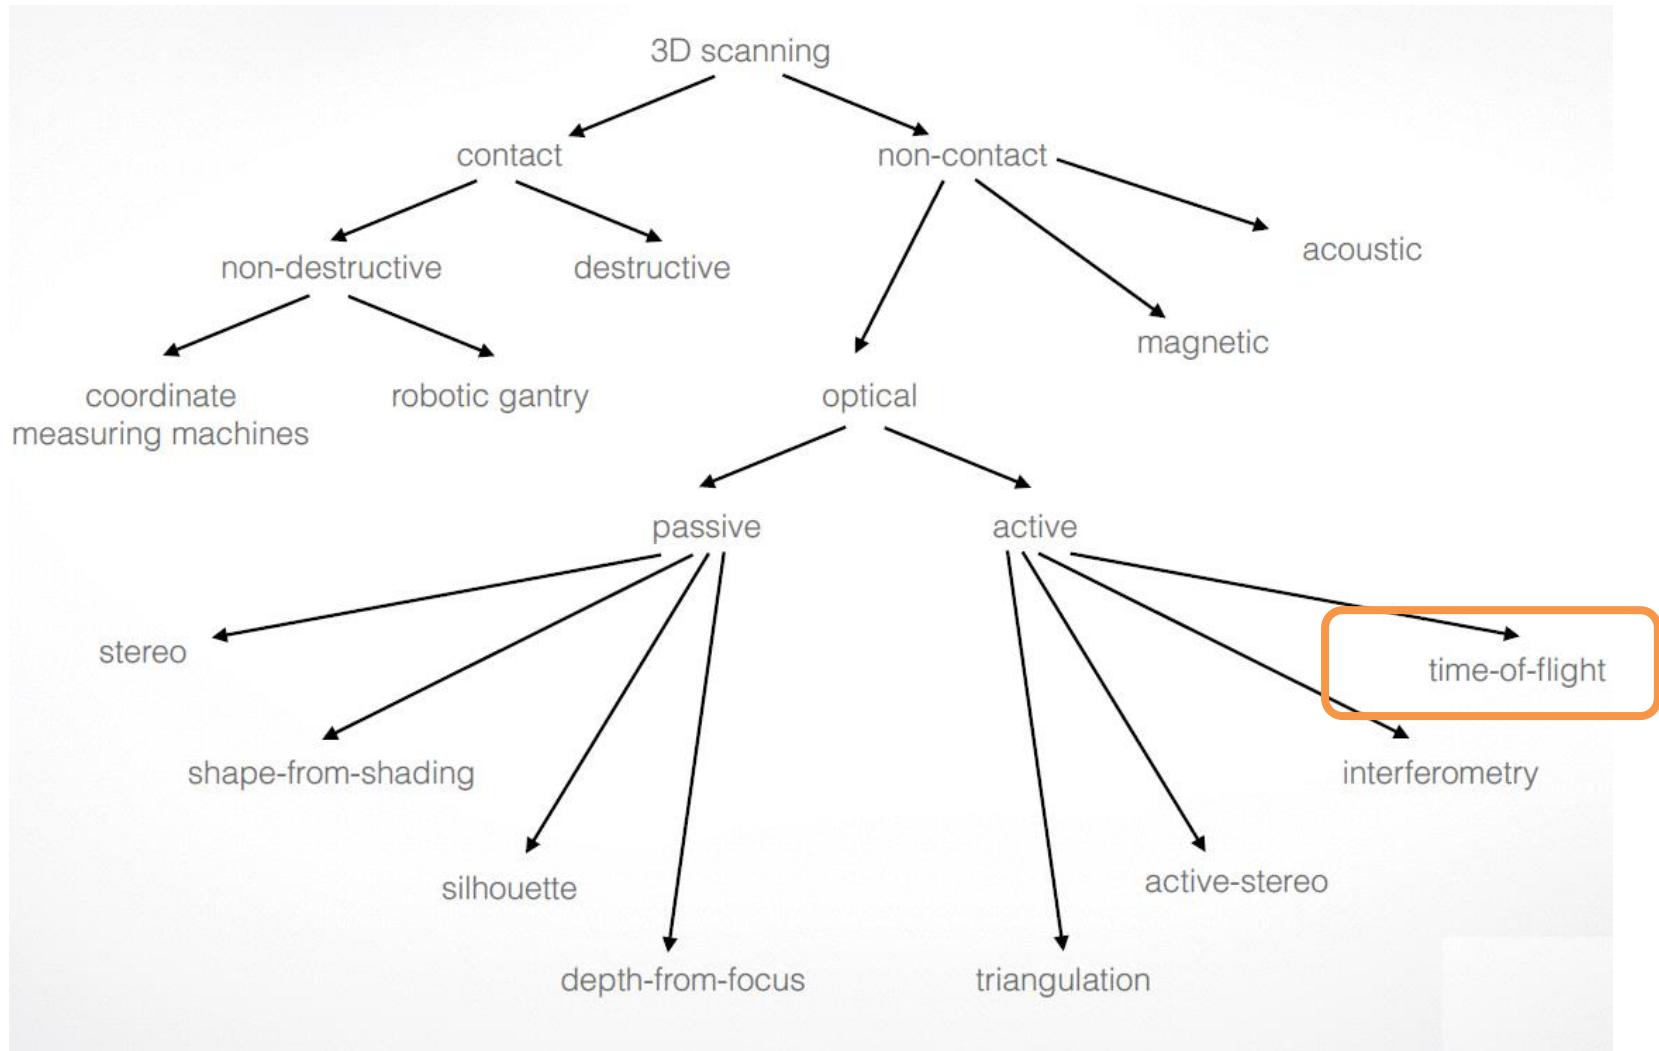
A taxonomy



Active acquisition systems

- based on **sensor** and **emitter** (controlled EM wave)
- influence of surface reflectance to emitted signal
- **correspondence** problem simplified (via known signal) → less computation (realtime?)
- examples (laser, structured light)
- high resolution and dense capture possible, even for texture poor regions
- more sensitive to surface reflection properties (mirrors?)

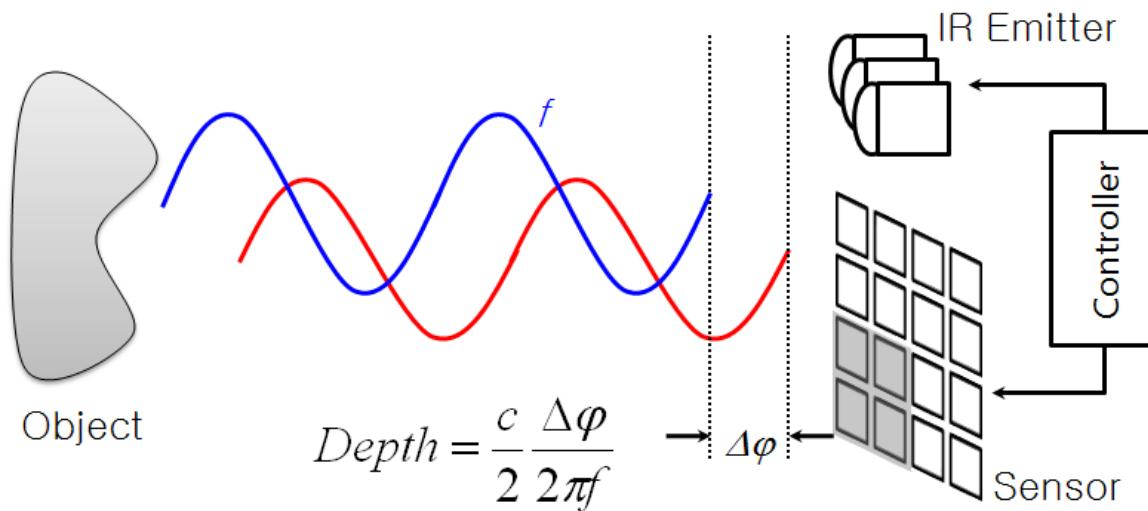
A taxonomy



Time of flight

Probe object by laser or infrared light!

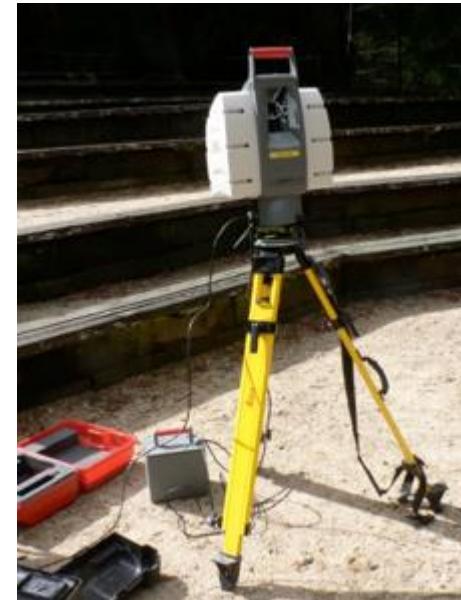
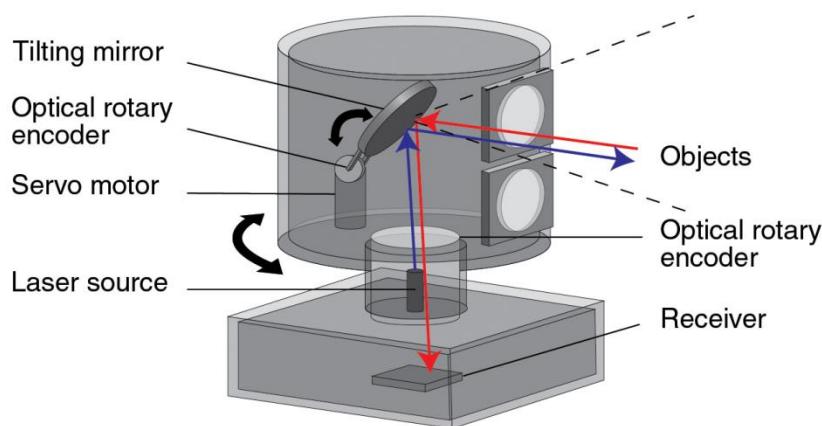
- Emit pulse of light, measure time till reflection from surface is seen by a detector
- Known speed of light & round-trip time allows to compute distance to surface



Time of flight

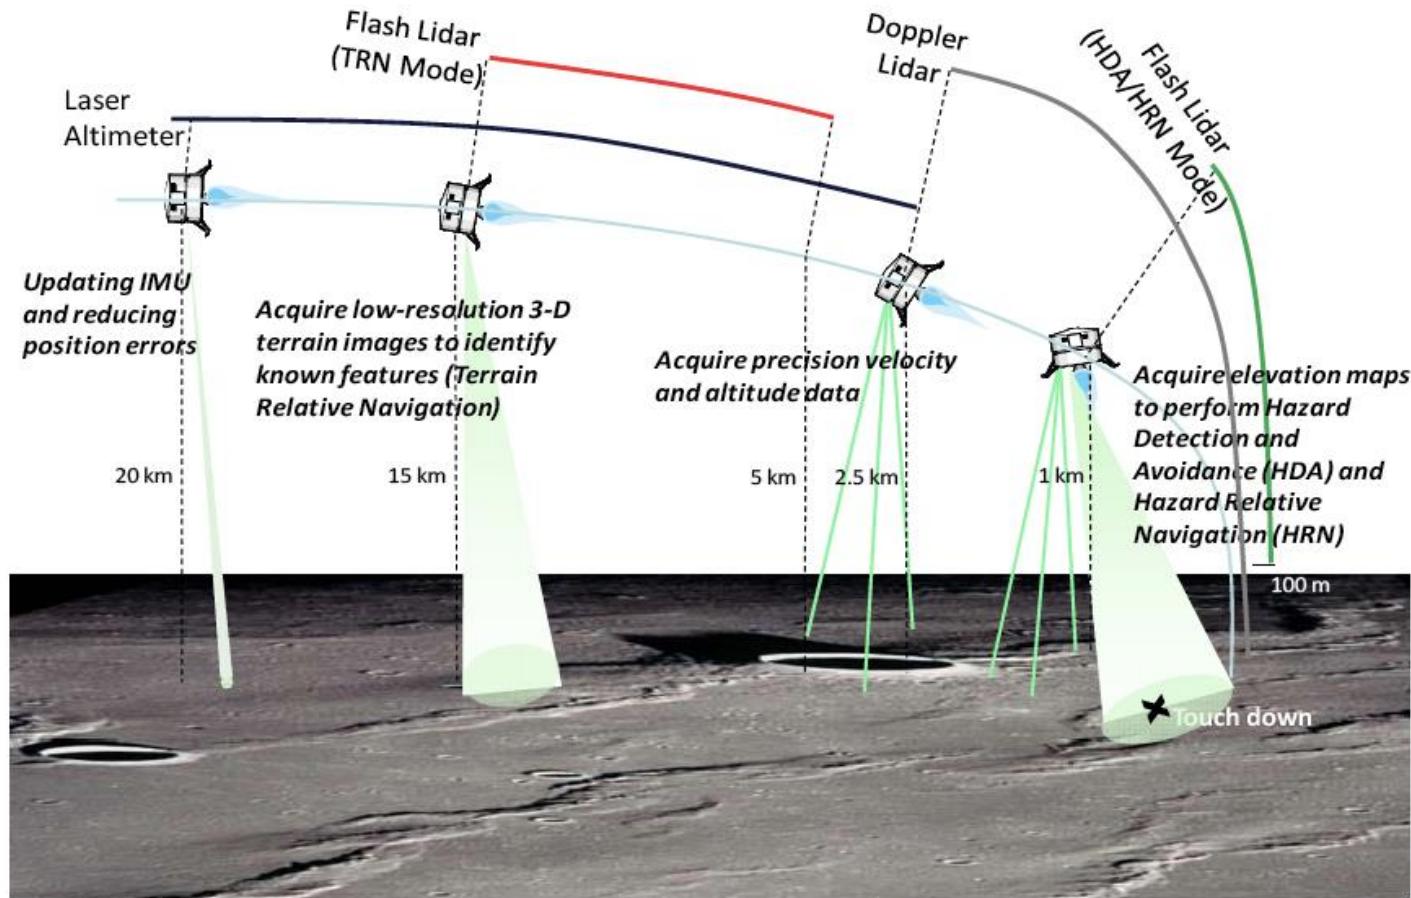
Laser LIDAR!

- Light Detection and Ranging
- Good for long distance scans
- 6mm accuracy at 50 m distance



Time of flight

Laser LIDAR for planetary exploration



Time of flight

Laser LIDAR for planetary exploration

Sensor	Function	Operational Altitude Range	Precision/Resolution
Flash Lidar	HDA/HRN	1000 m – 100 m	5 cm/40 cm
	TRN	15 km – 5 km	20 cm/6 m
	Altimetry ¹	20 km – 100 m	20 cm
Doppler Lidar	Velocimetry	2500 m – 10 m	1 cm/sec
	Altimetry	2500 m – 10 m	10 cm
Laser Altimeter	Altimetry	20 km – 100 m	20 cm
	TRN ¹	15 km – 5 km	20 cm

Time of flight

Infrared light!

- 176x144 pixels, up to 50 fps
- 30 cm to 5 m distance, 1 cm accuracy
- technology is improving drastically

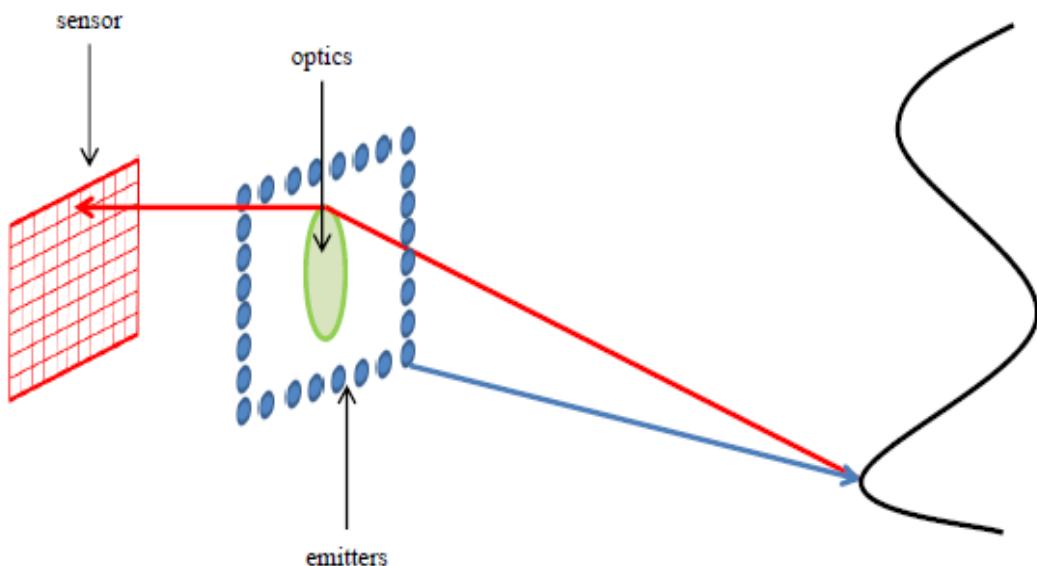


Fig. 2.7 ToF camera structure and signaling: propagation towards the scene (blue arrow), reflection (from the black surface on the right), back-propagation (red arrow) towards the camera through the optics (green) and reception (red sensor).

Matricial Time-of-Flight

Time of flight

Kinect One (= second gen Kinect)!

- 30 fps
- Depth map x/y resolution: 512 x 424
- z-resolution 1 mm & accuracy:
- <1.5 mm (depth < 50 cm)
- < 3.9 mm (depth < 180 cm)
- < 17.6 mm (depth < 450 cm)
- 1080 HD for RGB input
- uses Kinect2 SDK



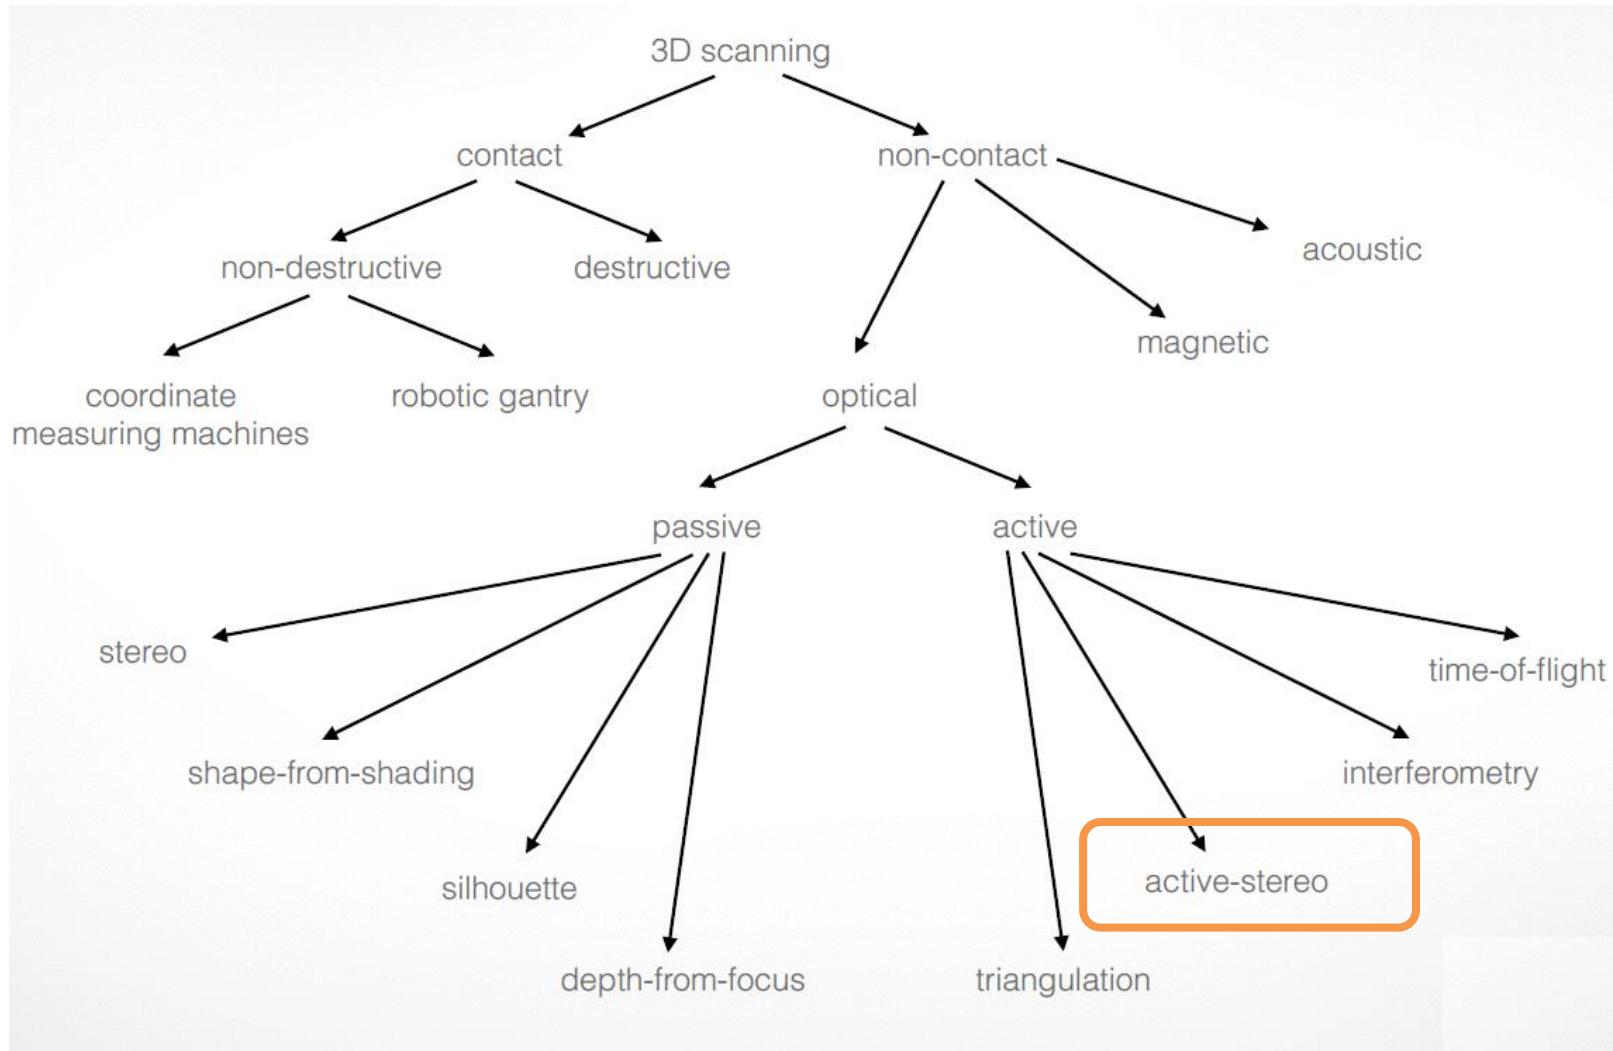
Time of flight

Intel RealSense L514

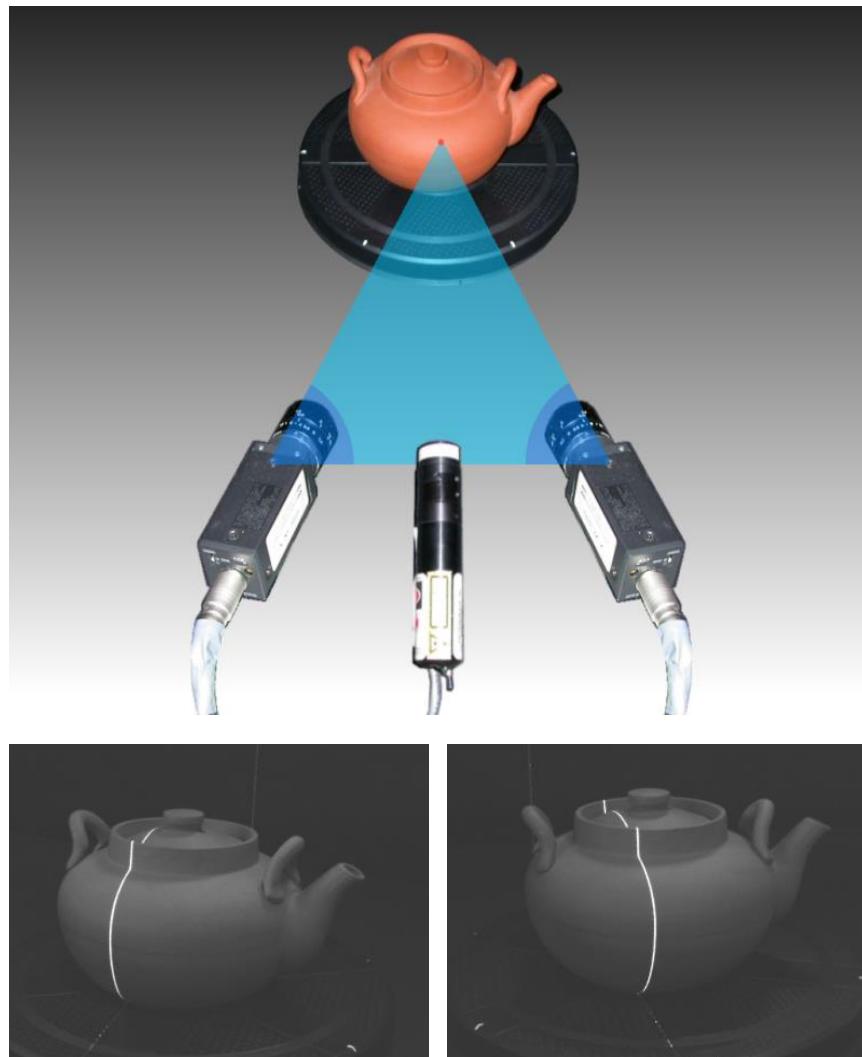
- 30 fps
- Depth map x/y resolution:
1024x768
- **Depth Accuracy:**
~5 mm to ~14 mm thru 9 m²
- 1920x1080 RGB input



A taxonomy



Active stereo

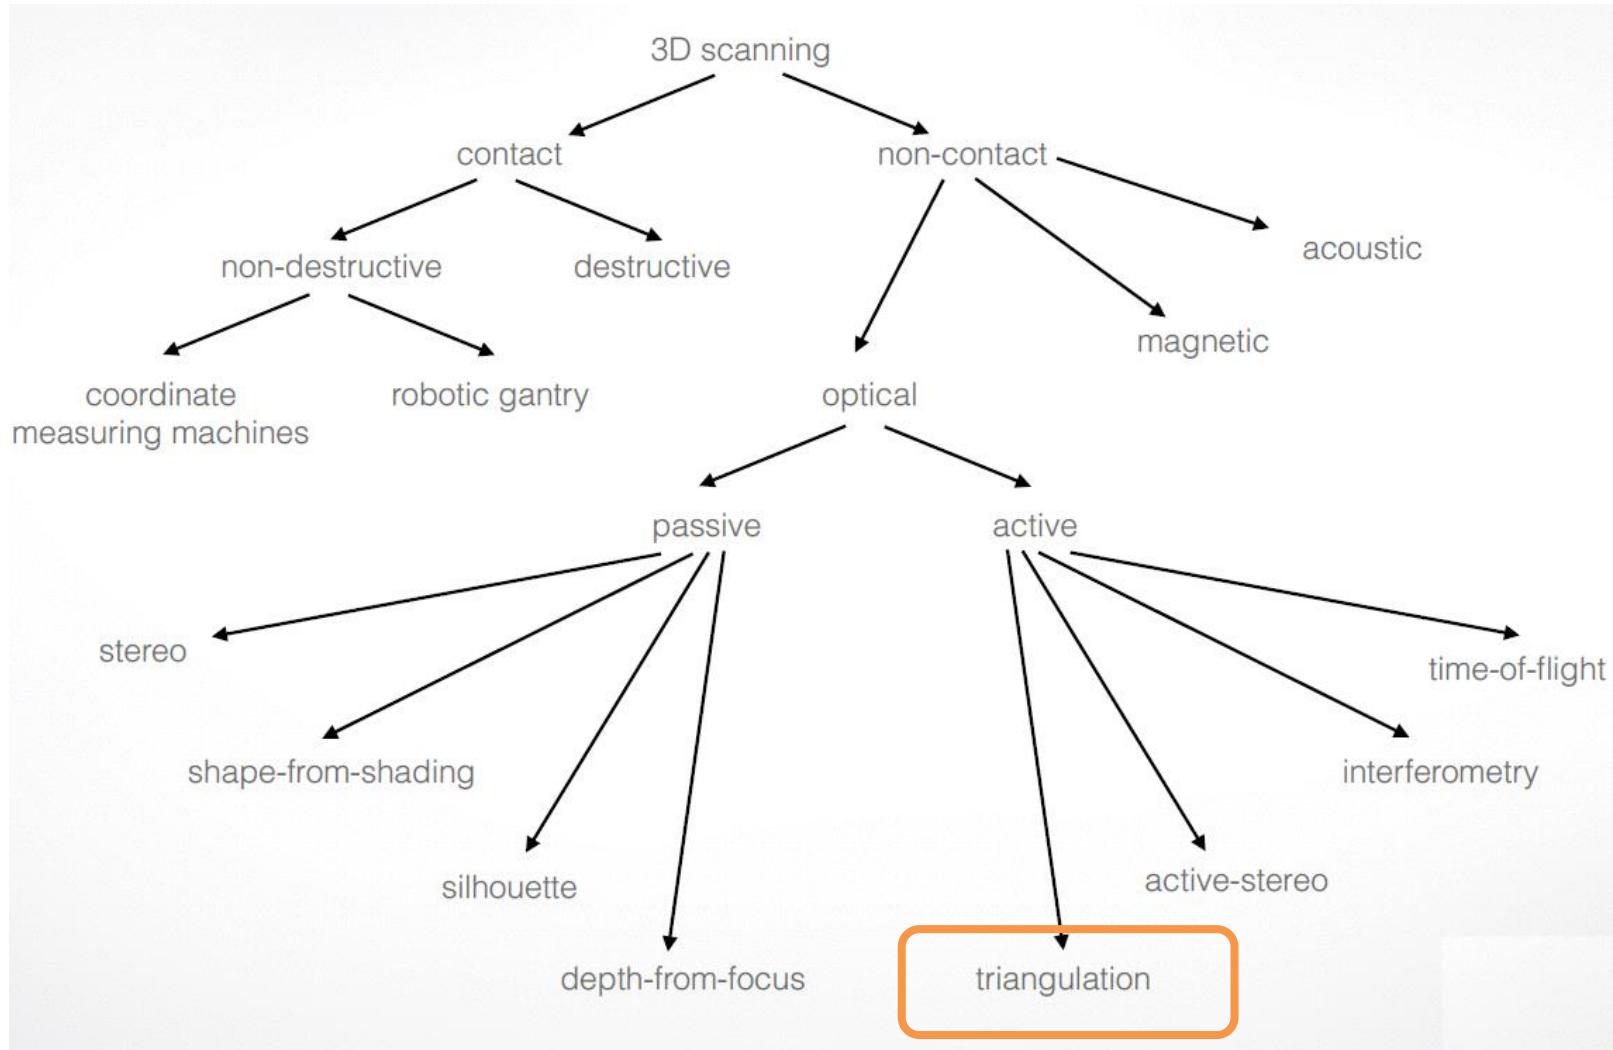


Active stereo

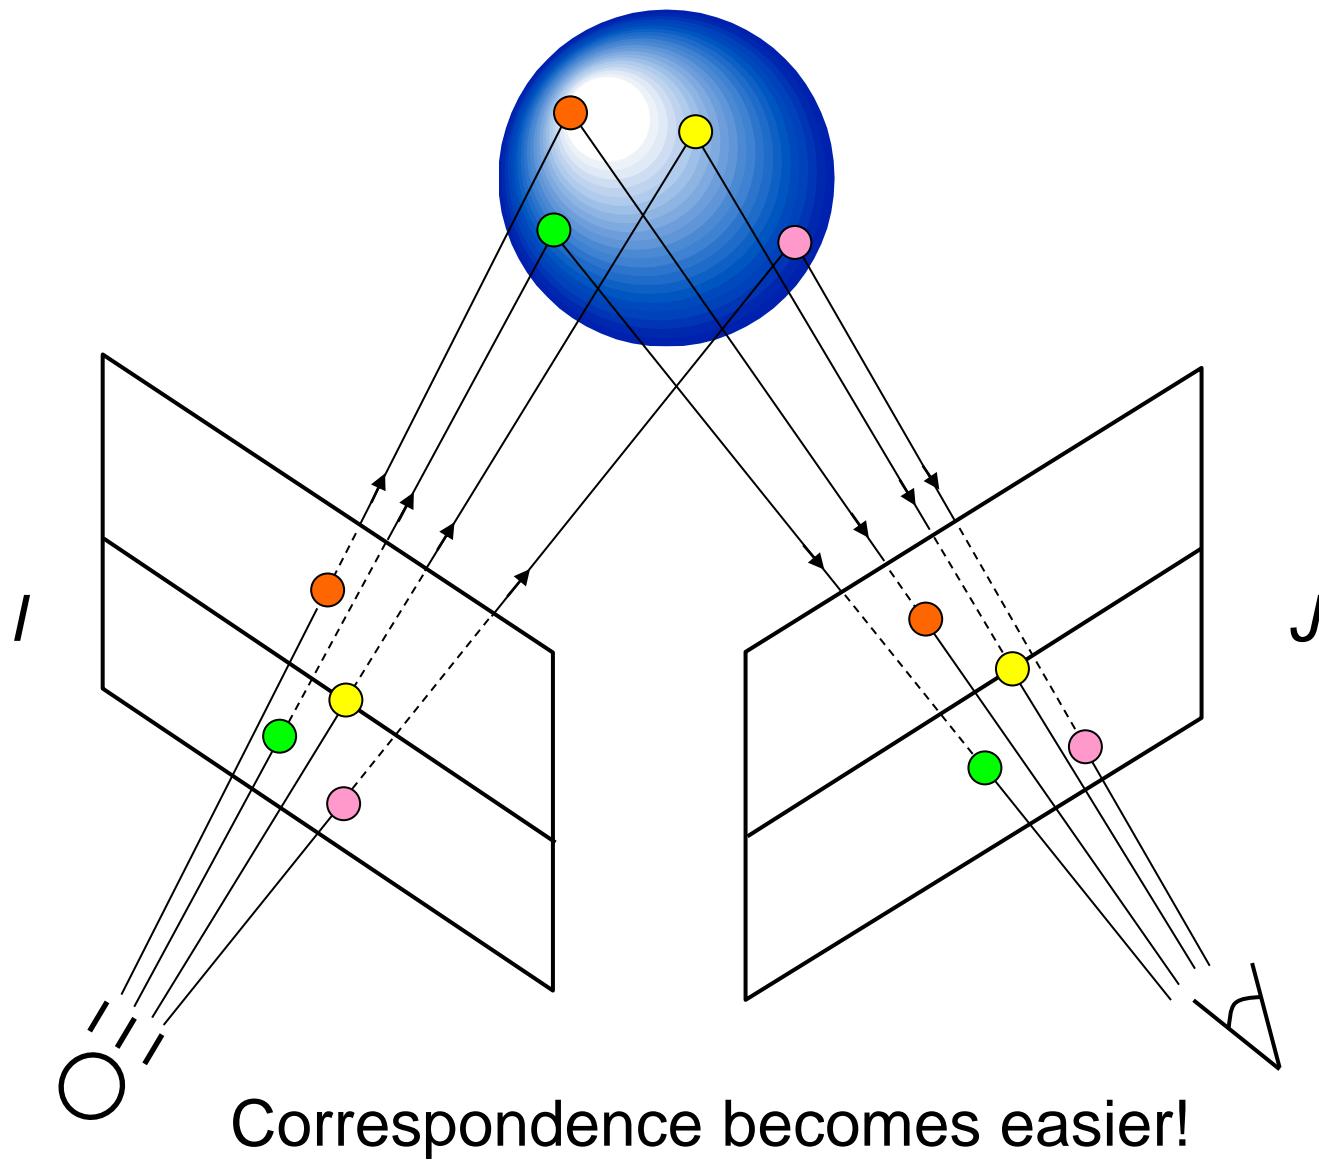


M. Waschbüsch et al. *Scalable 3D Video of Dynamic Scenes*. *The Visual Computer*, pp. 629-638, 2005

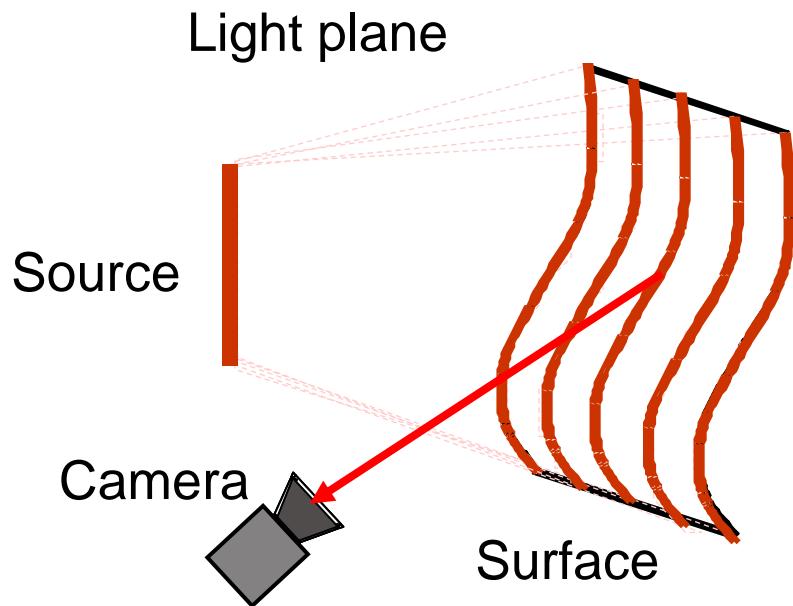
A taxonomy



Structured Light Triangulation

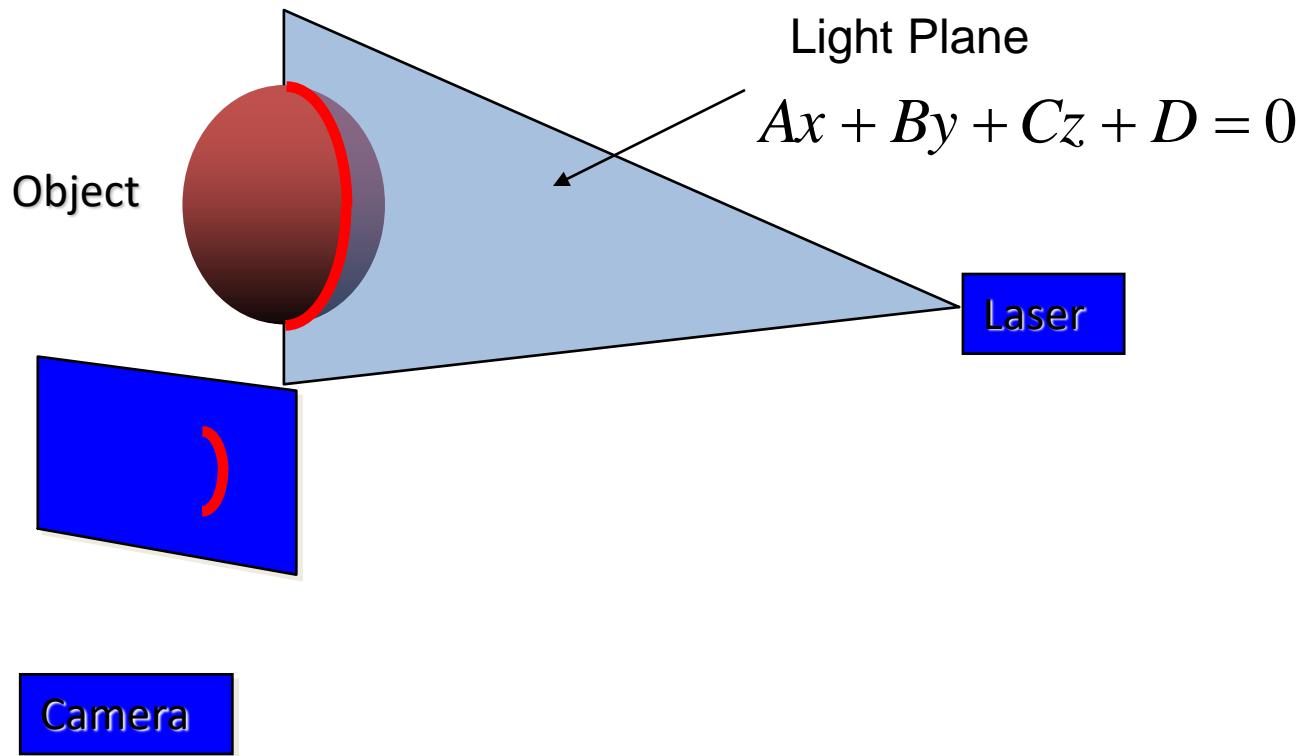


Light Stripe Scanning – Single Stripe



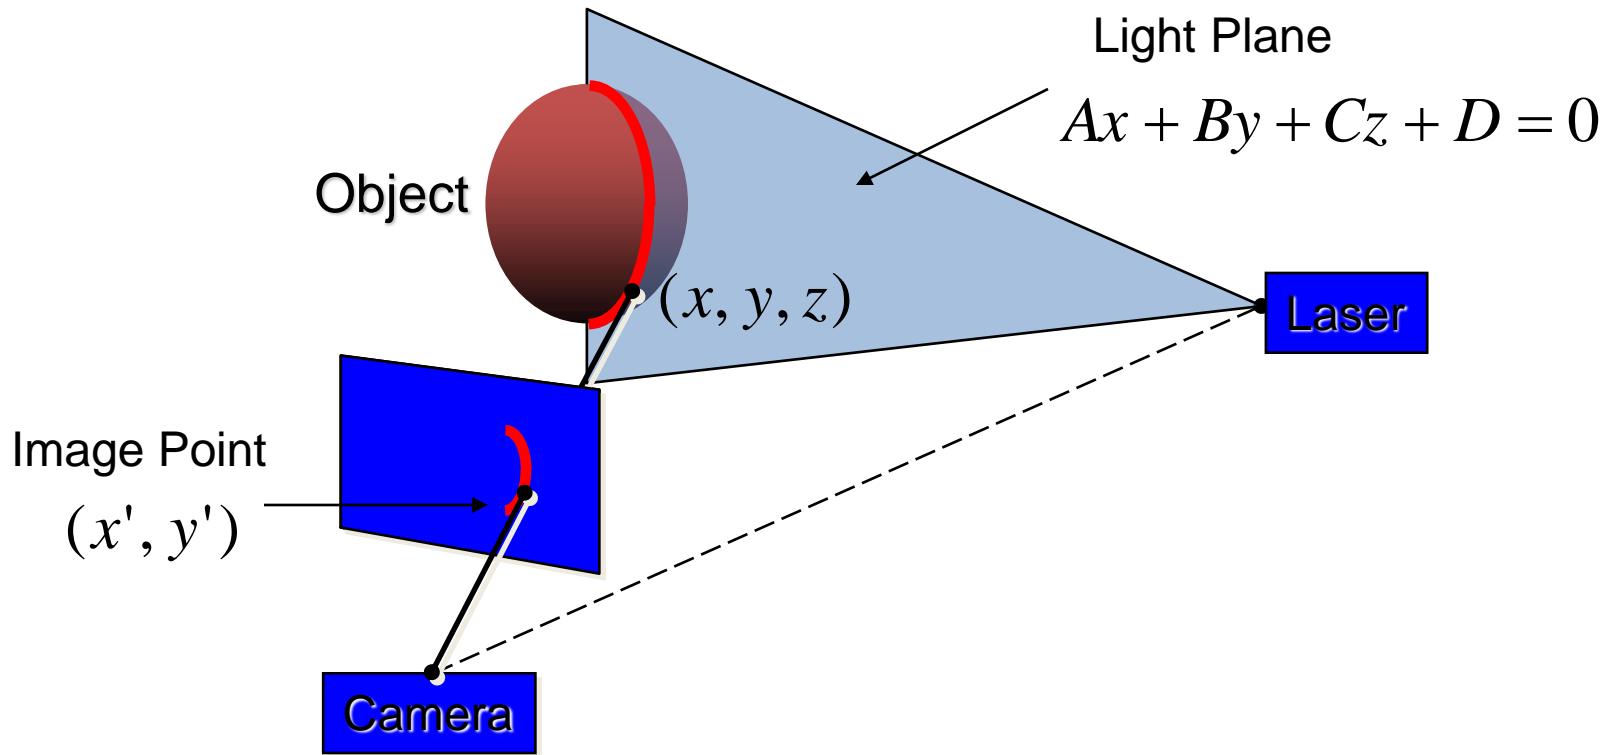
- Optical triangulation
 - Project a single stripe of laser light
 - Scan it across the surface of the object
 - This is a very precise version of structured light scanning
 - Good for high resolution 3D, but needs many images and takes time

Triangulation



- Project laser stripe onto object

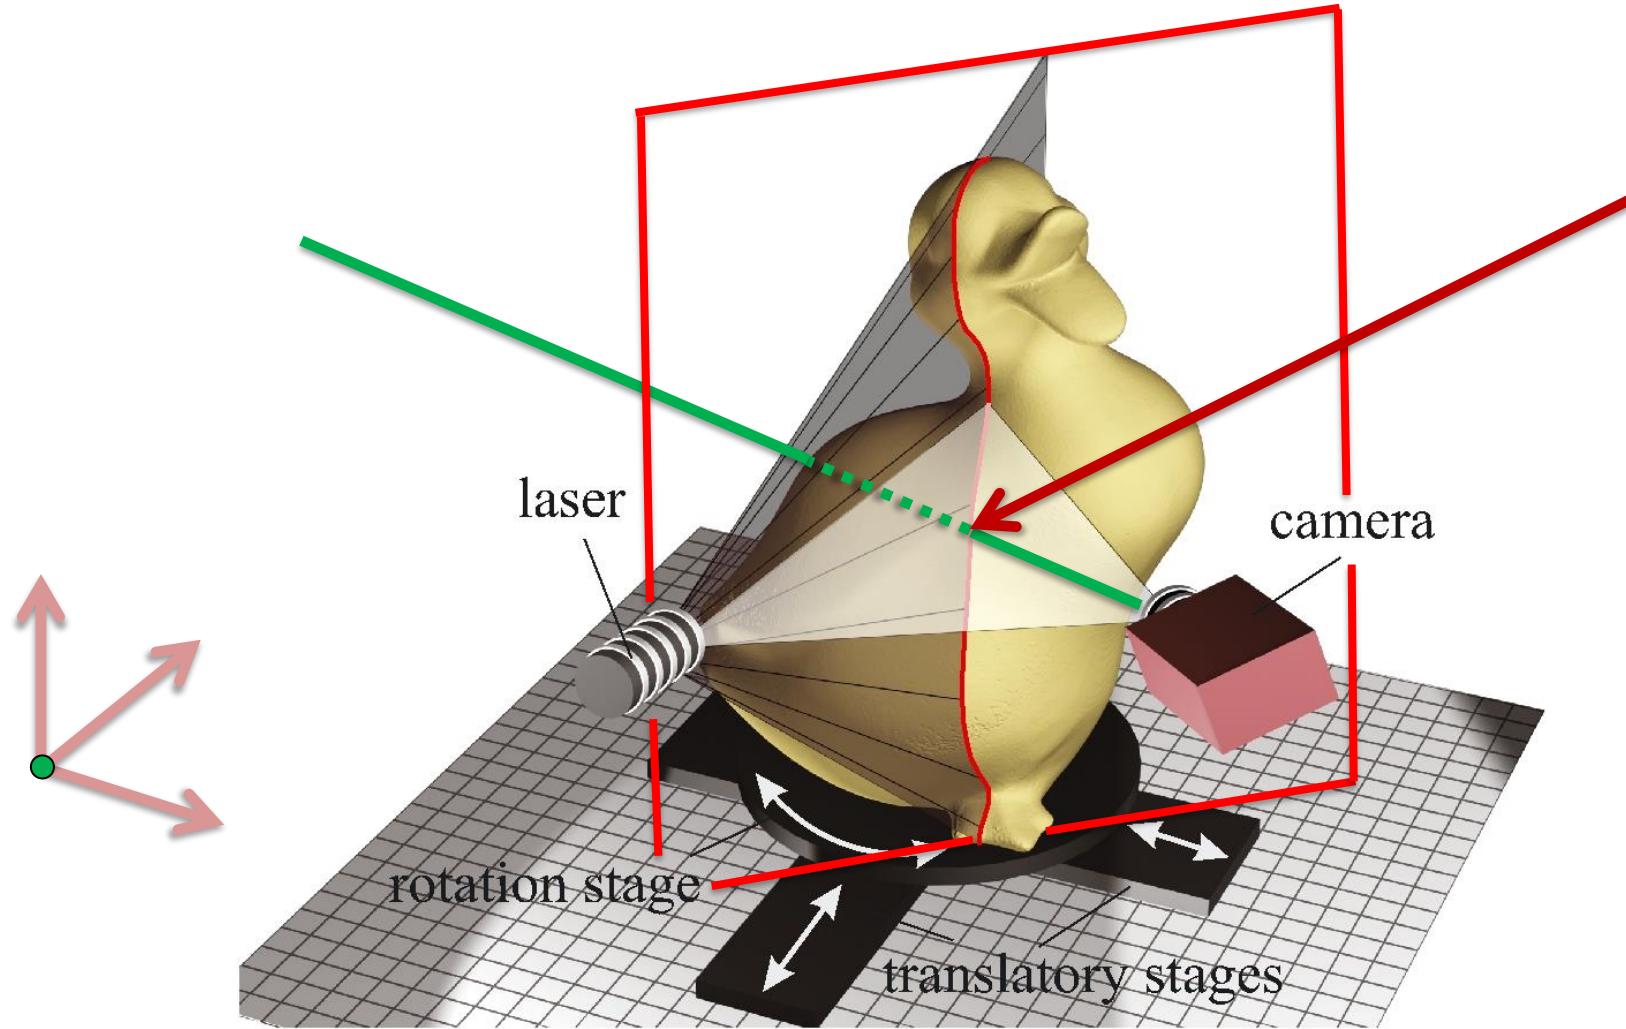
Triangulation



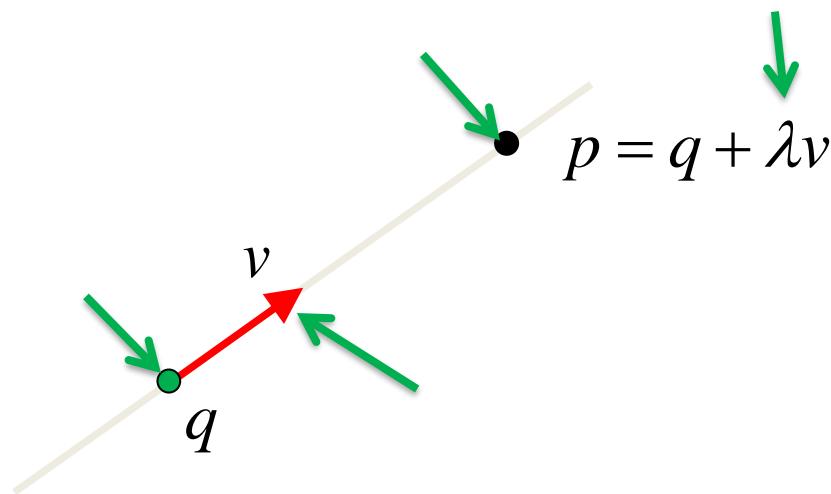
- Depth from ray-plane triangulation:
 - Intersect camera ray with light plane

$$x = x' z / f \quad z = \frac{-Df}{Ax' + By' + Cf}$$
$$y = y' z / f$$

3D Triangulation: plane-ray intersection



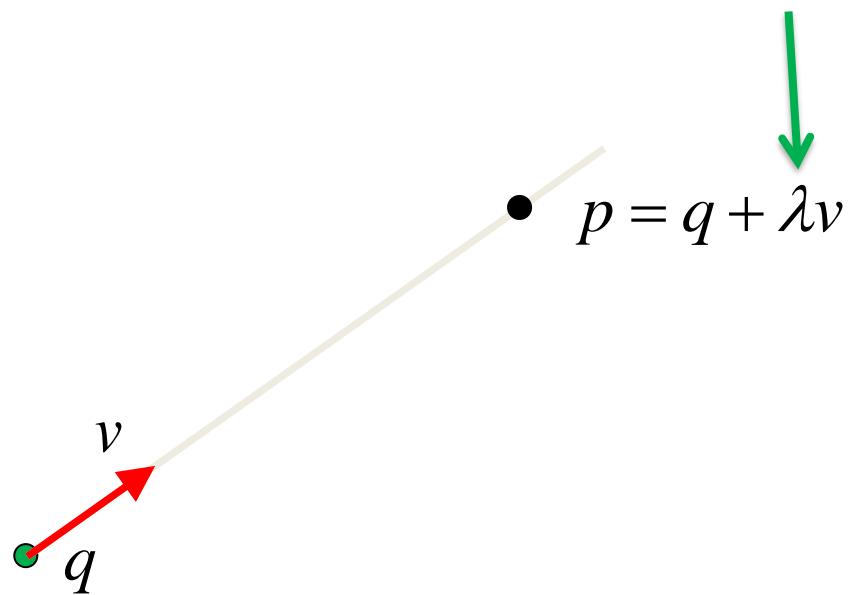
Representation of lines and rays



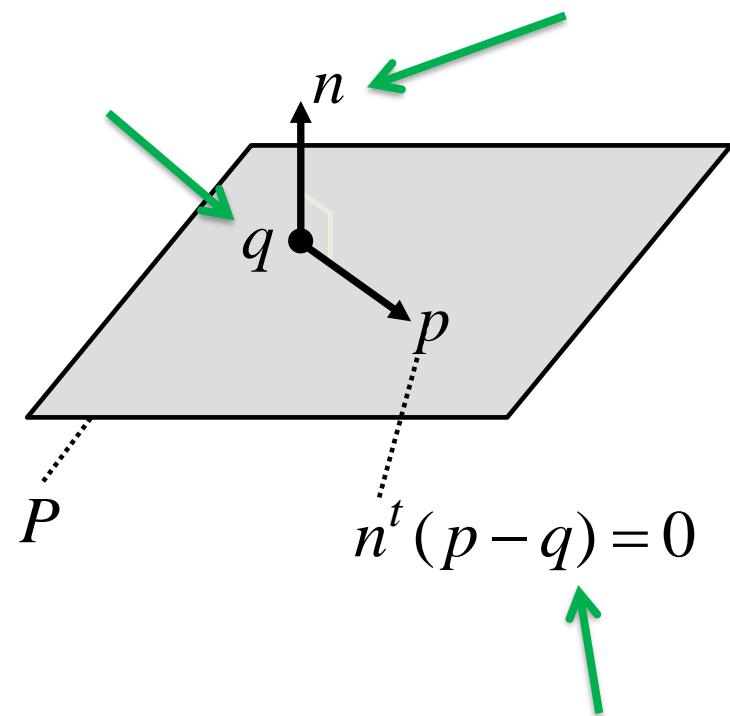
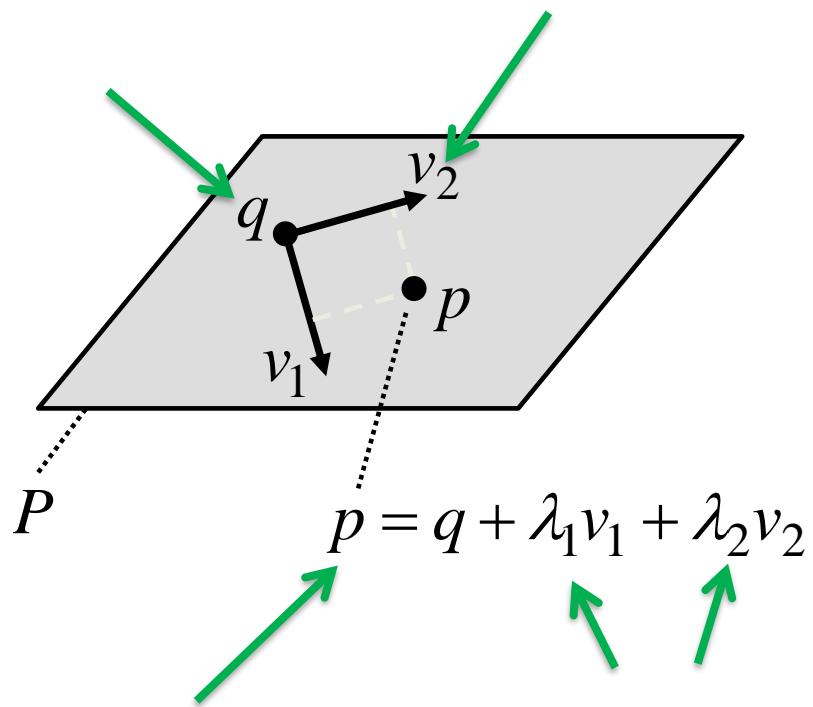
COMPUTER VISION

$$\underbrace{M}_{P} = C + \lambda \underbrace{\begin{bmatrix} Q \\ 0 \end{bmatrix}}_{\sim m}$$

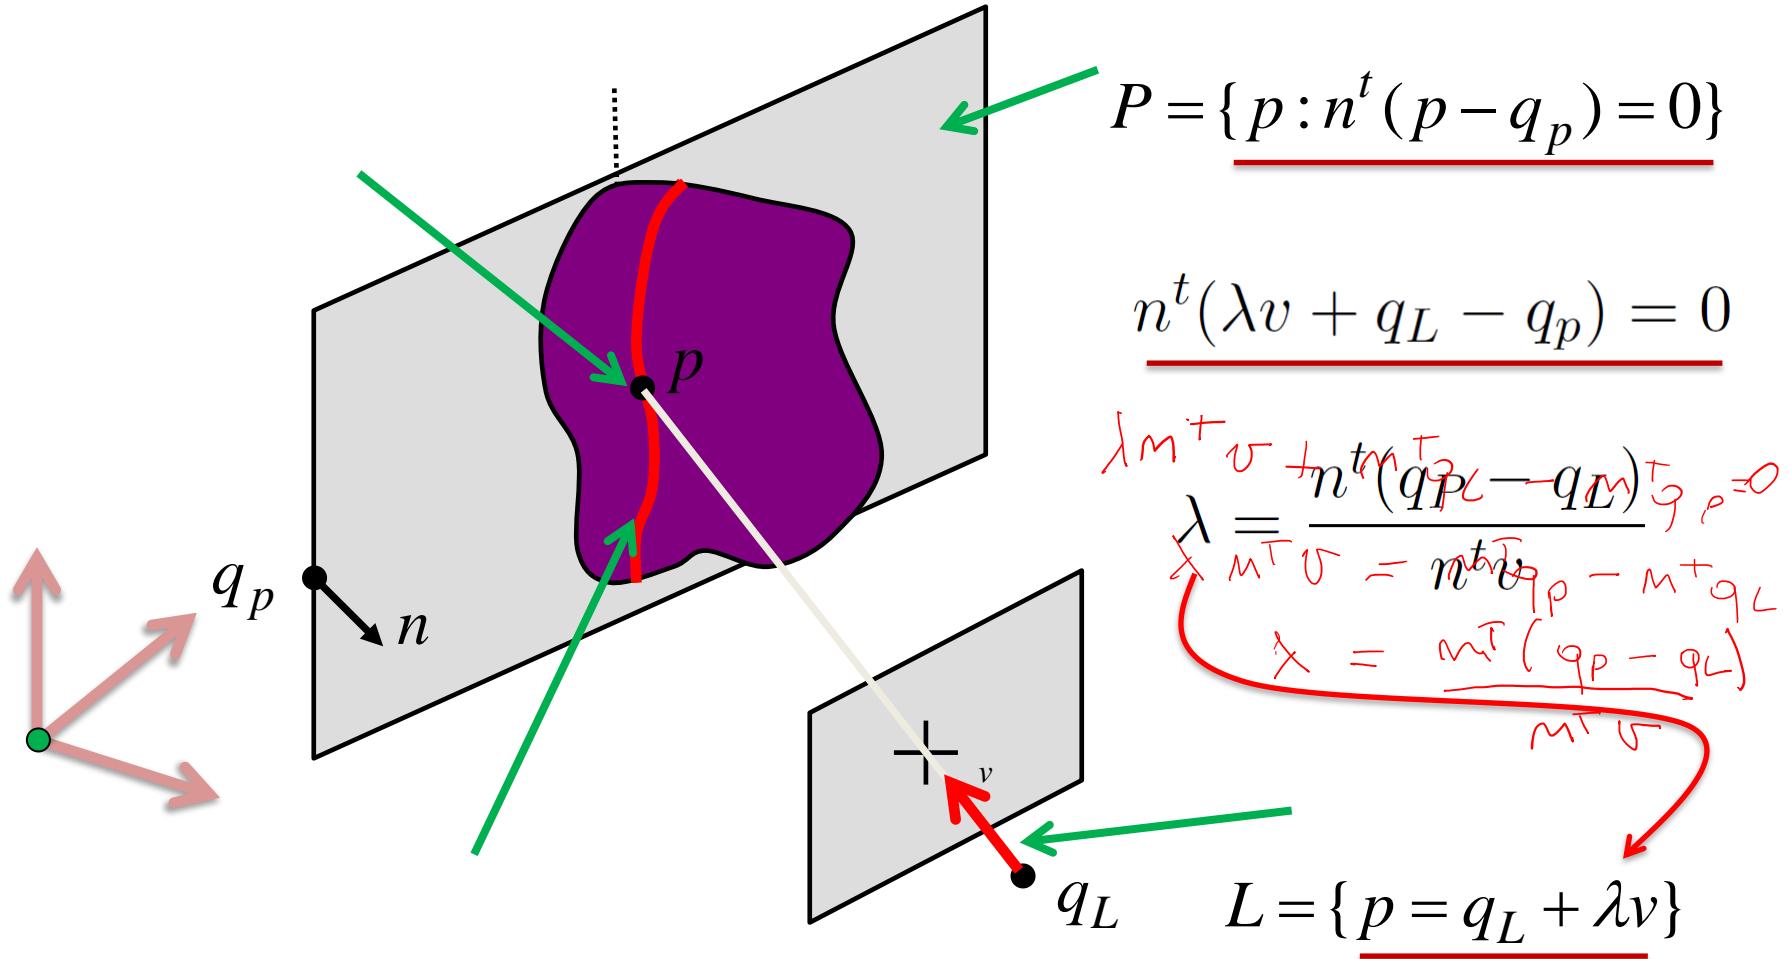
$$P = [Q \mid q]$$



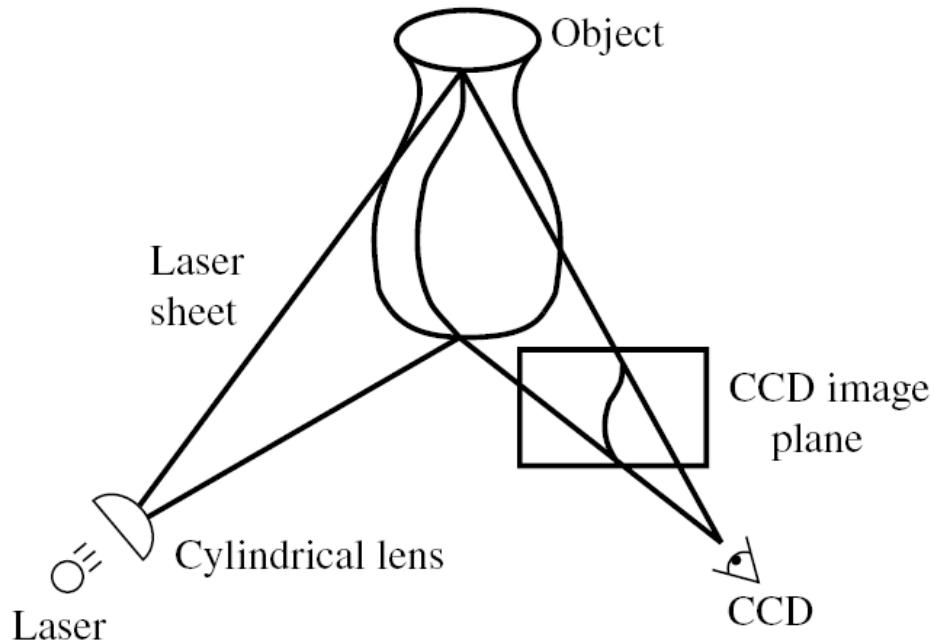
Representation of planes



Triangulation by line-plane intersection



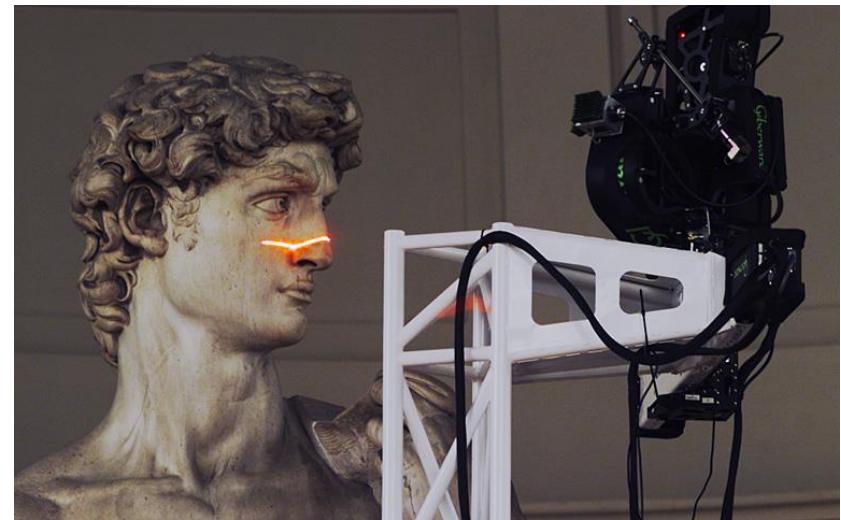
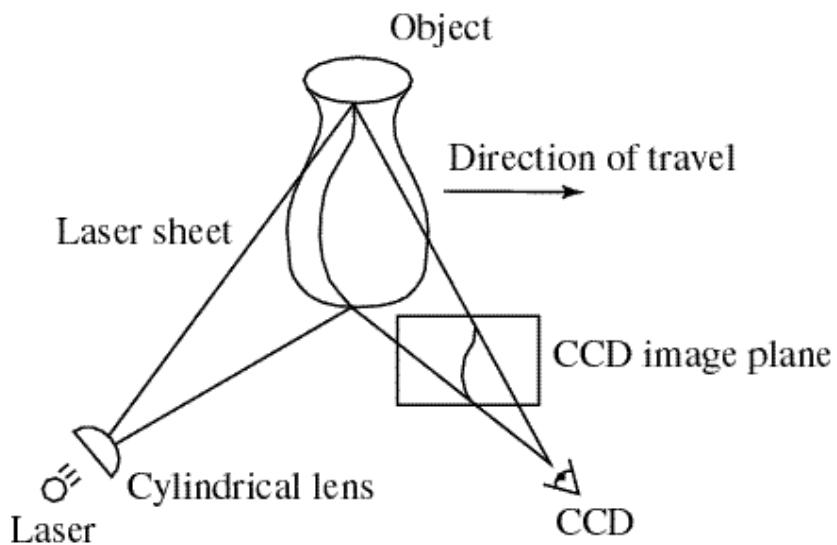
Example: Laser scanner



Cyberware® face and head scanner

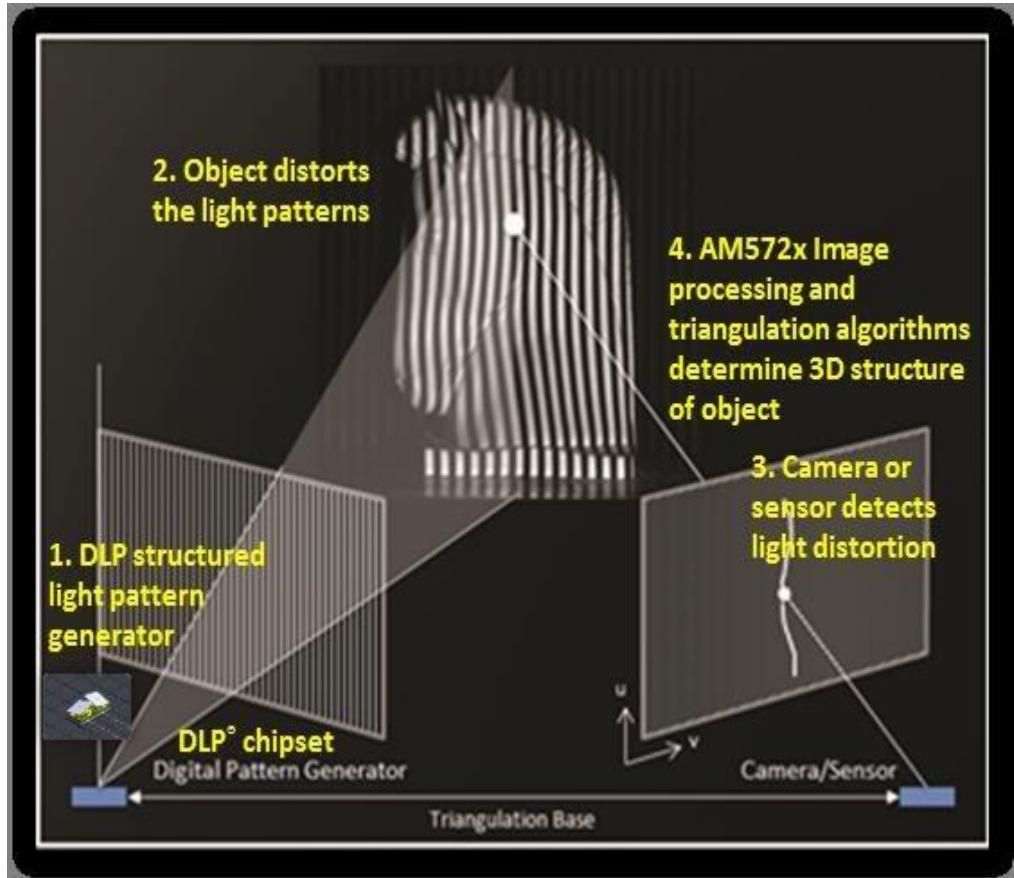
- + very accurate < 0.01 mm
- more than 10sec per scan

Example: Laser scanner



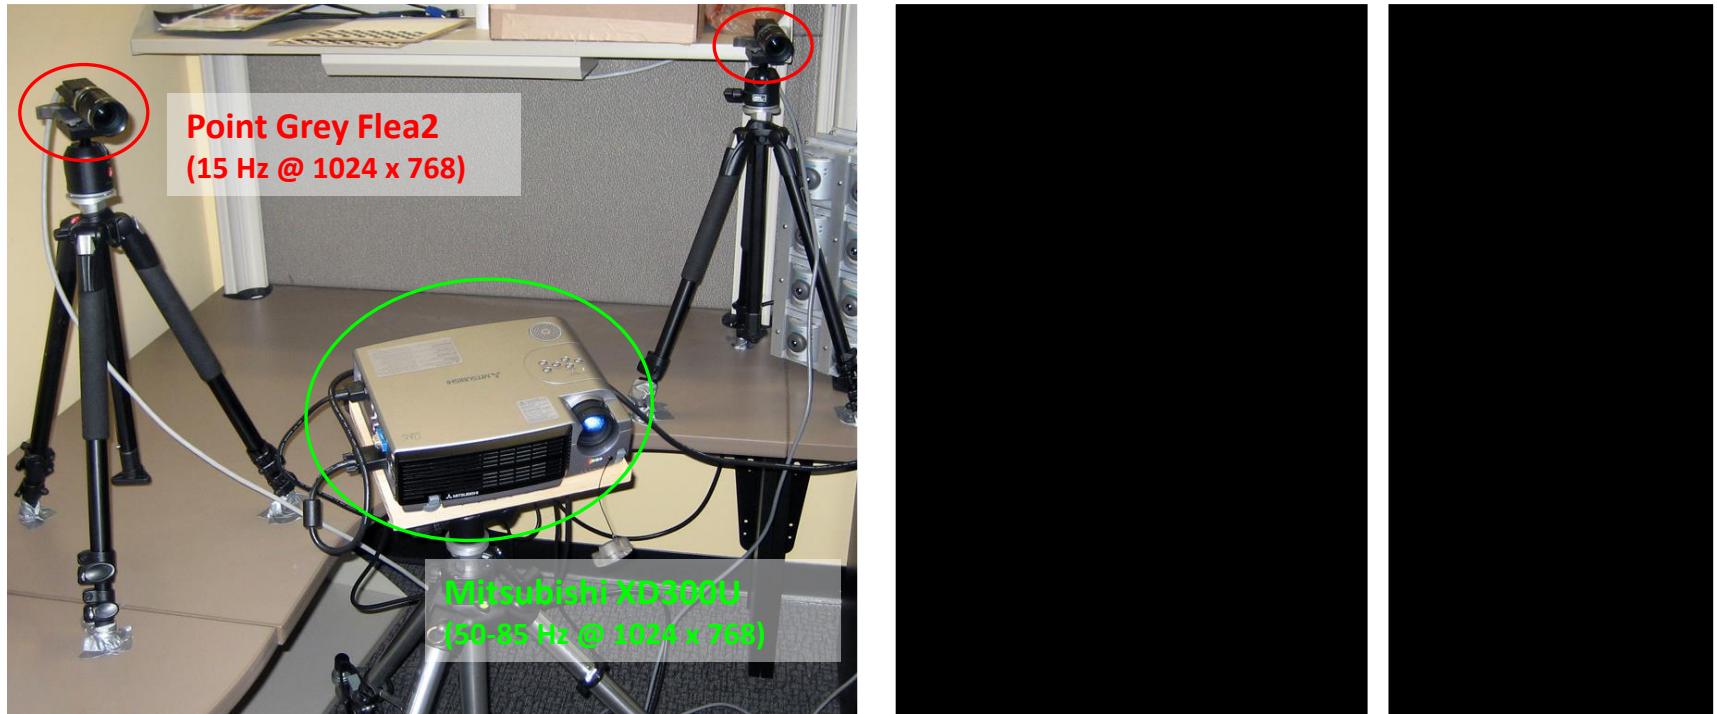
Digital Michelangelo Project
<http://graphics.stanford.edu/projects/mich/>

Lighting Patterns



- Any spatio-temporal pattern of light projected on a surface using a standard light projector.
- Cleverly illuminate the scene to extract correspondences
- Avoids problems of mechanical motion
- Allows faster acquisition projecting multiple patterns simultaneously

Structured Lighting: Swept-Planes Revisited

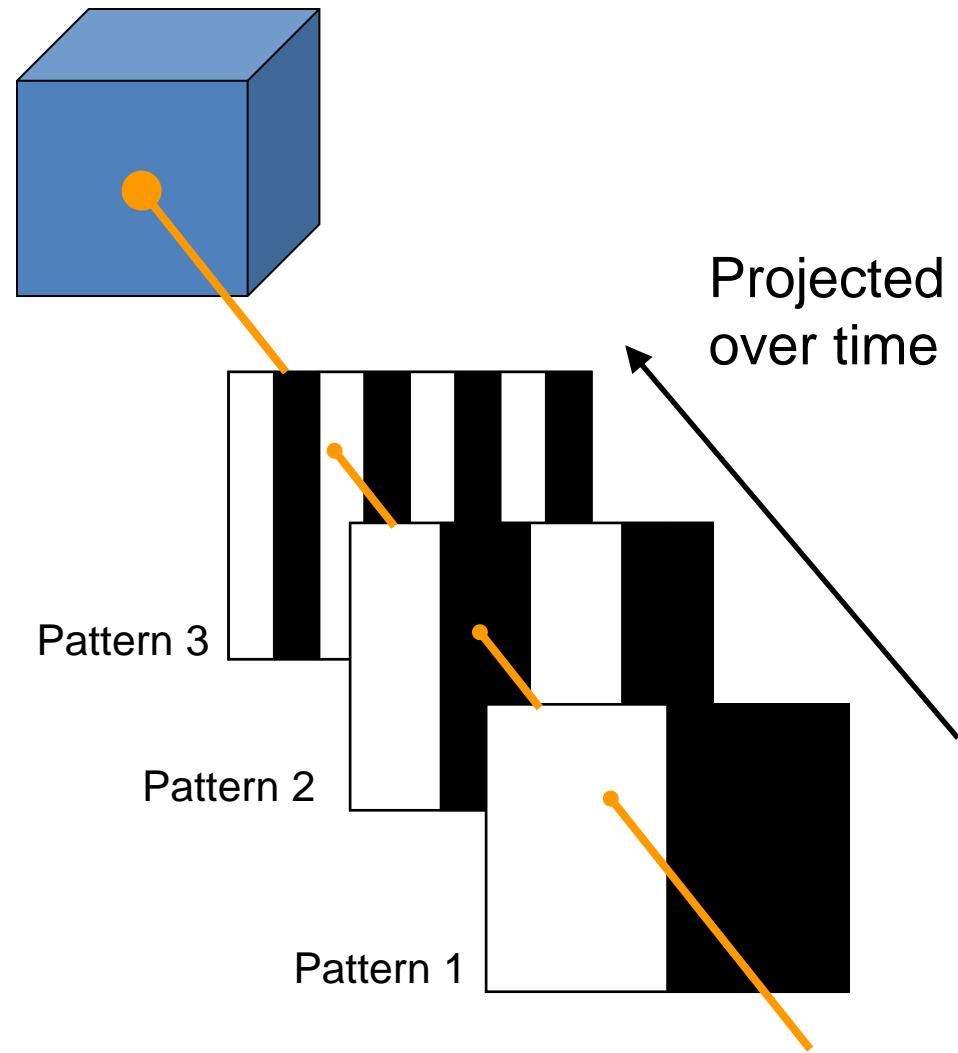


- Swept-plane scanning recovers 3D depth using ray-plane intersection
- Use a data projector to replace manually-swept laser/shadow planes
- How to assign correspondence from projector planes to camera pixels?
which stripe is which?
- Solution: Project a spatially- and temporally-encoded image sequence
- **What is the optimal image sequence to project?**

Binary Coding

Faster:

$2^n - 1$ stripes in n images.

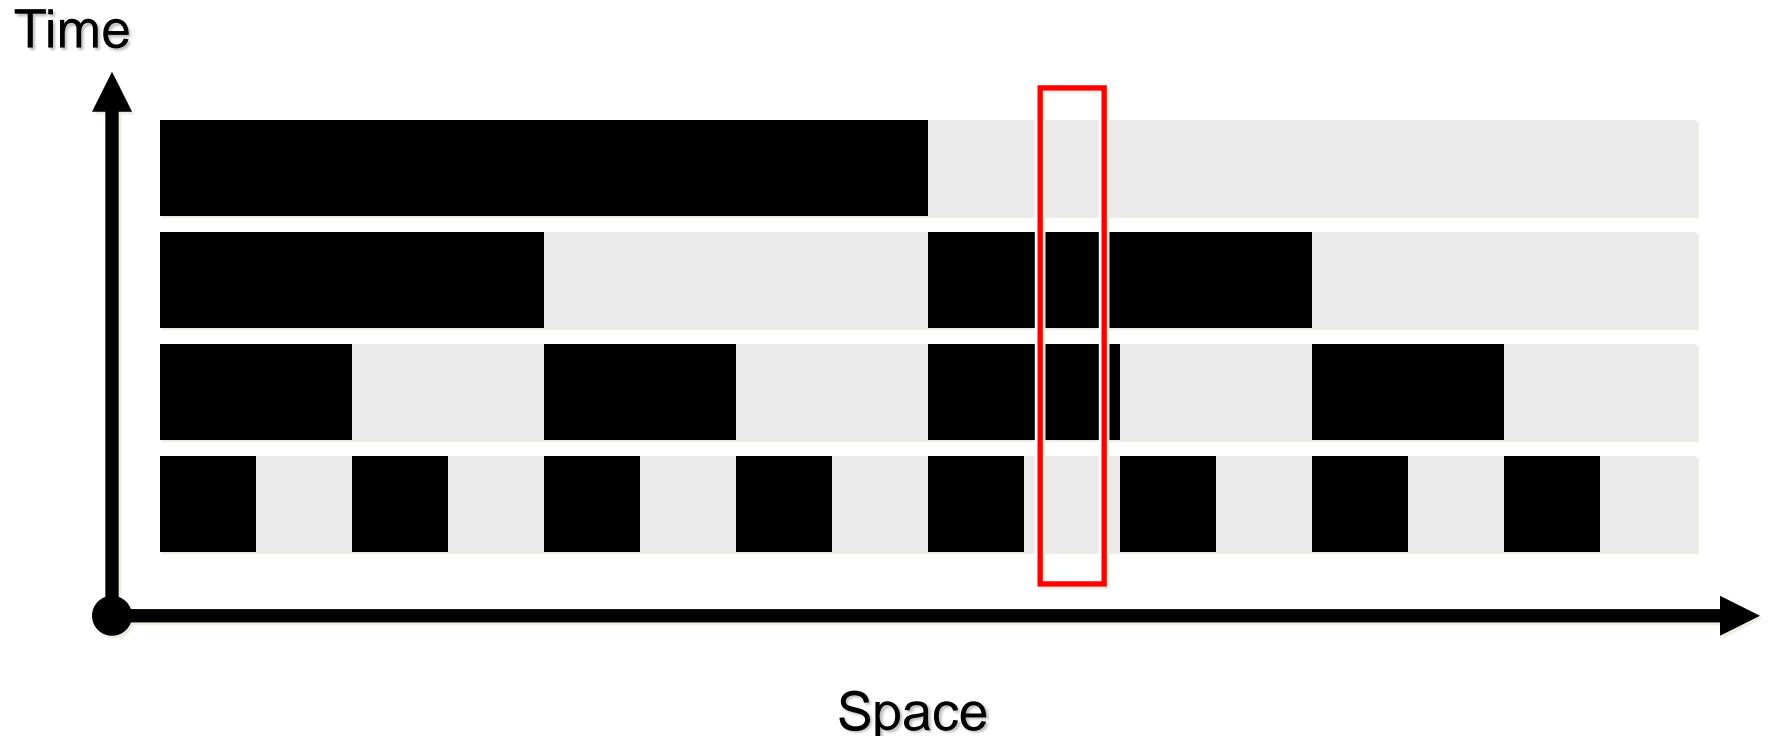


Example:

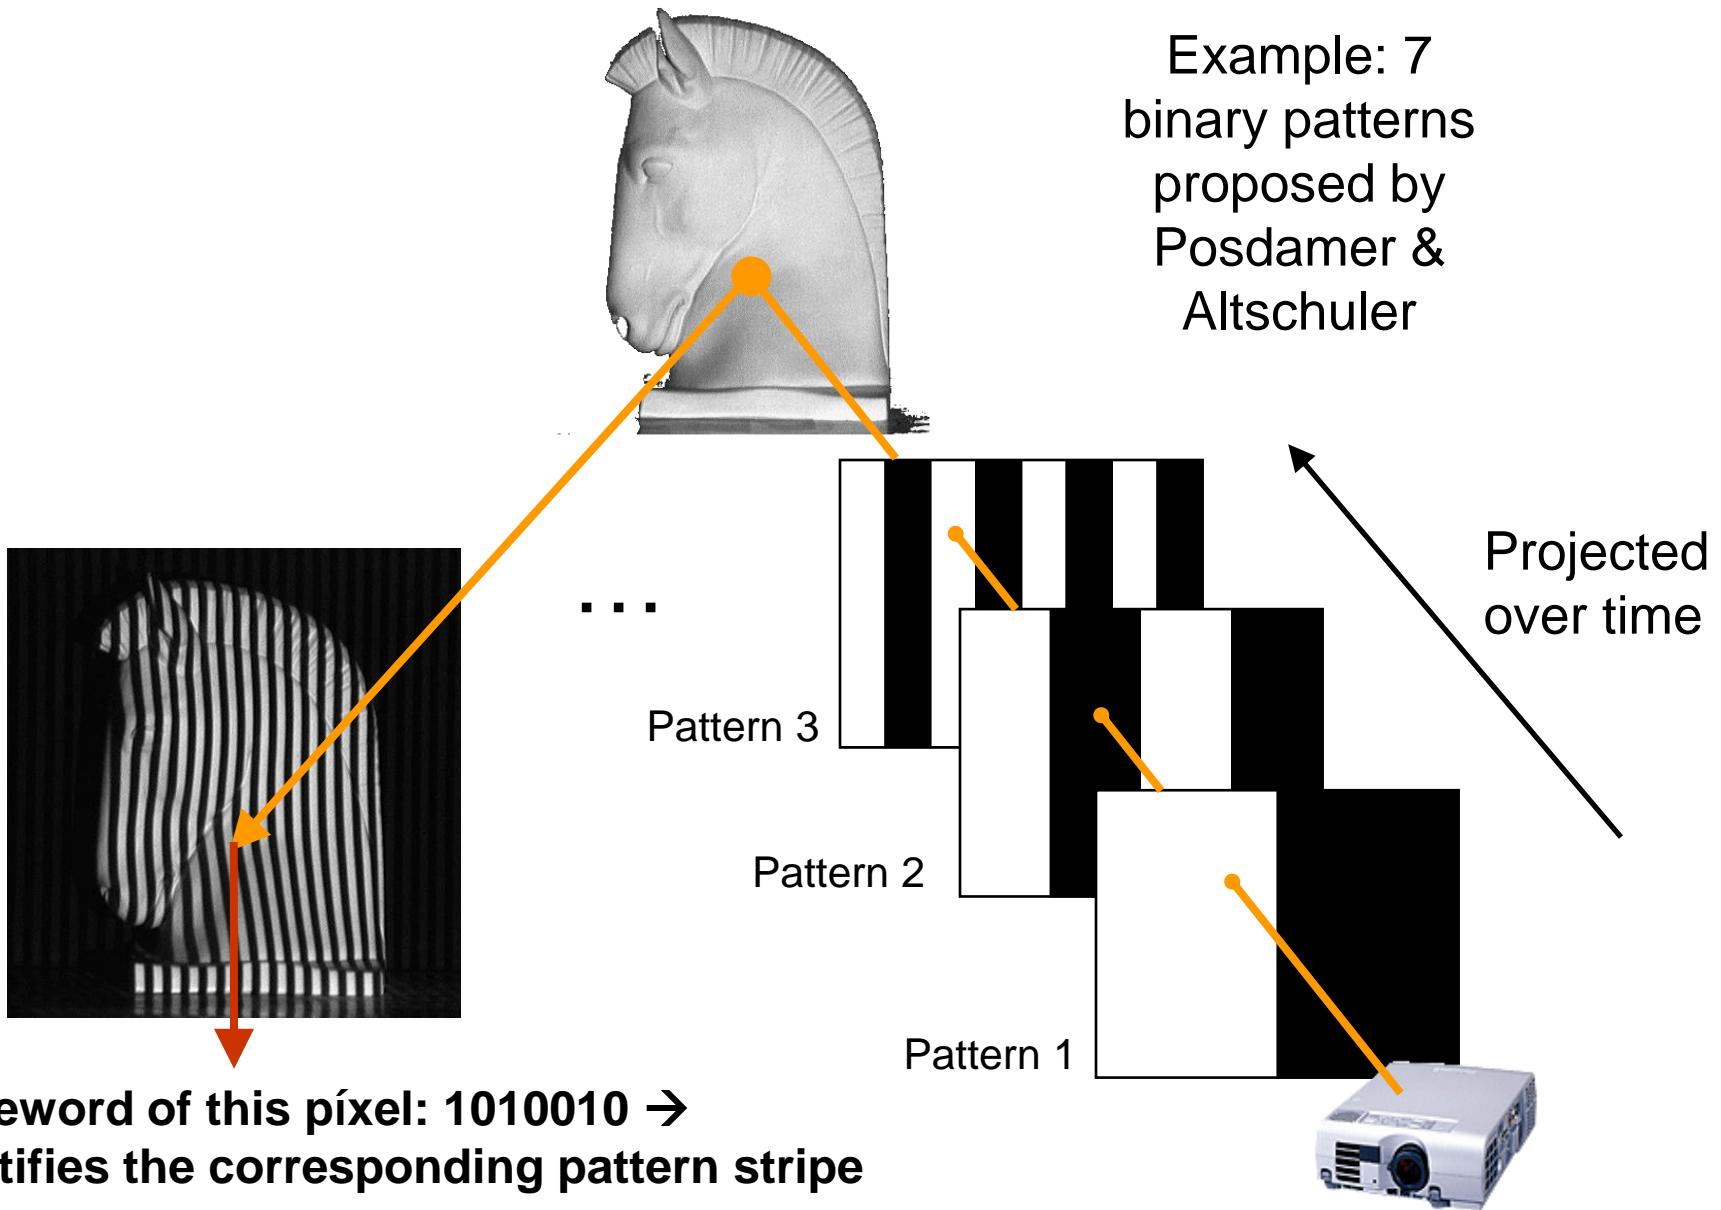
3 binary-encoded patterns
which allows the measuring
surface to be divided in 8 sub-
regions

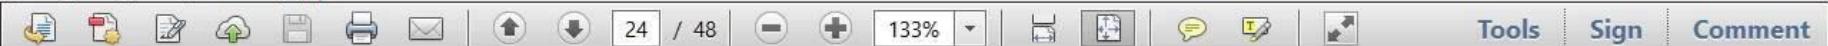
Binary Coding

- Assign each stripe a unique illumination code over time [Posdamer 82]



Binary Coding





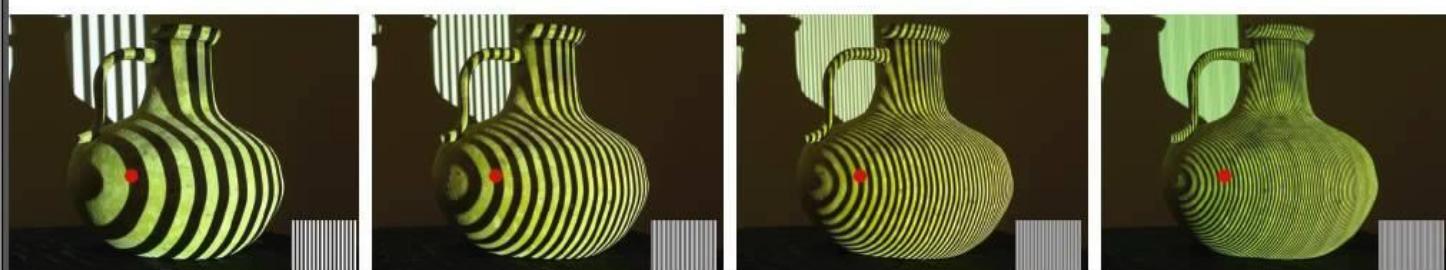
24 / 48

133%

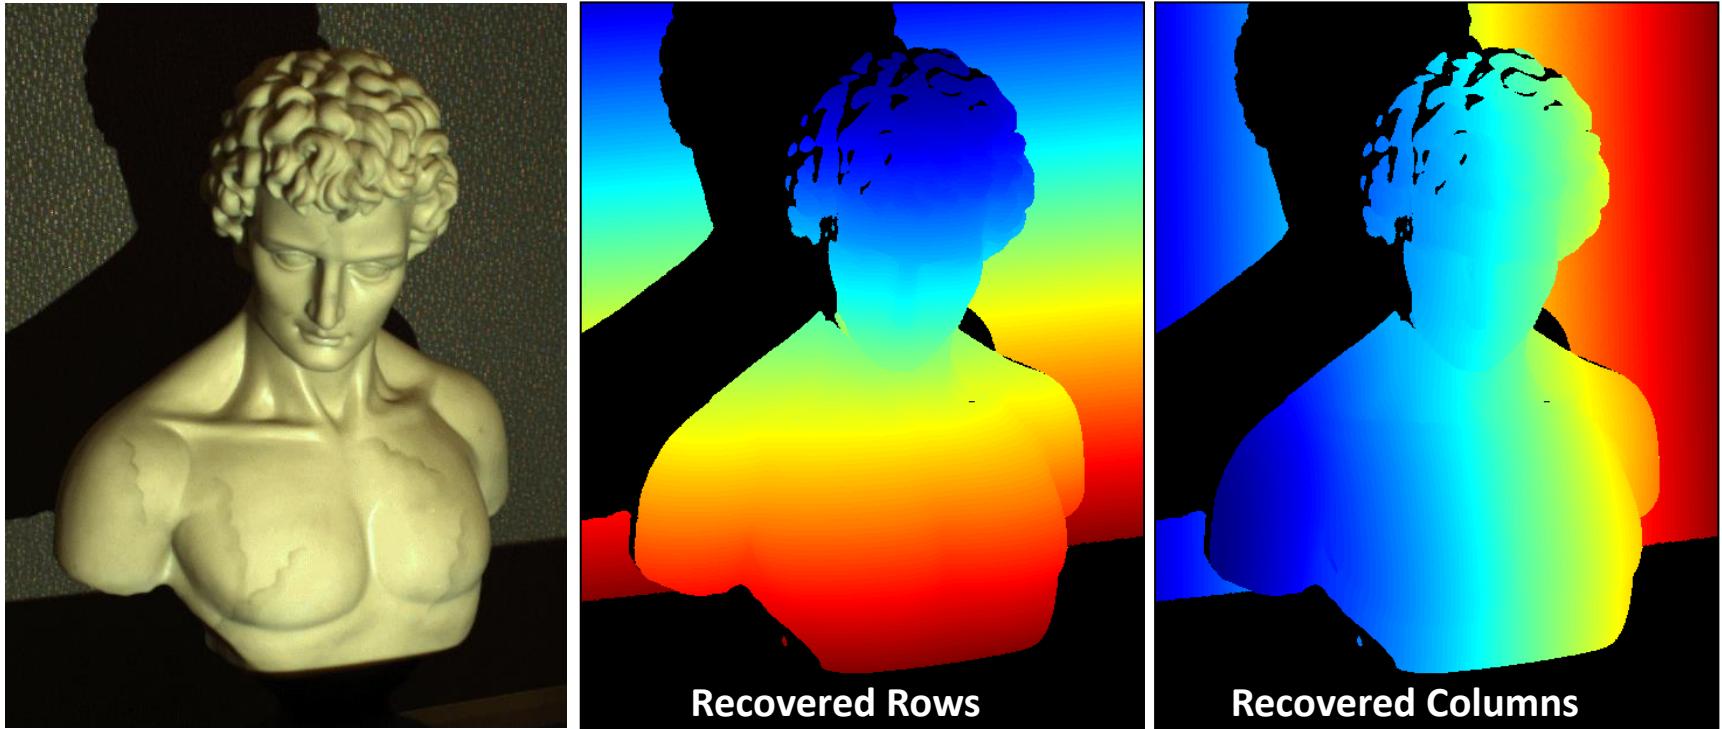
Tools Sign Comment



01010001 = stripe 81



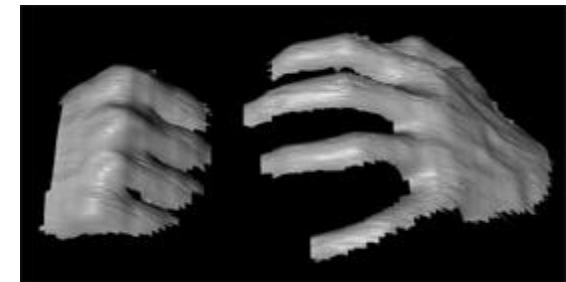
Gray Codes: Decoding Performance



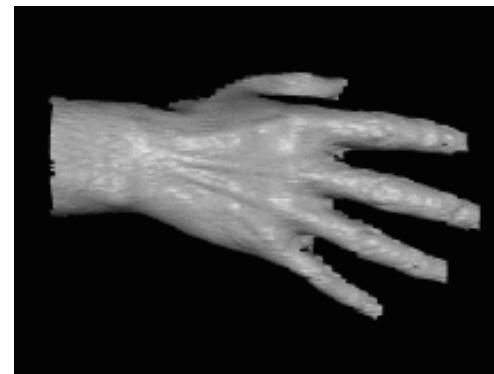
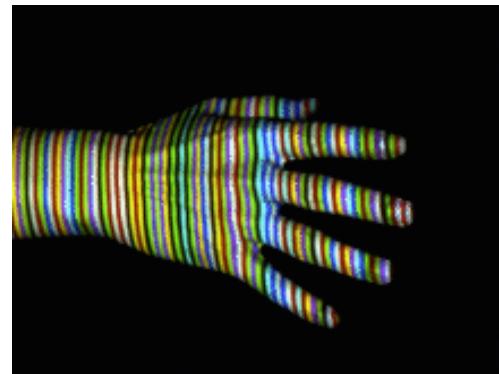
3D Reconstruction using Structured Light [Inokuchi 1984]

- Implemented using a total of 42 images (2 to measure dynamic range, 20 to encode rows, 20 to encode columns)
- Individual bits assigned by detecting if bit-plane (or its inverse) is brighter
- Decoding algorithm: Gray code → binary code → integer row/column index

More complex patterns



Works despite complex appearances



Works in real-time and on dynamic scenes

- Need very few images (one or two).
- But needs a more complex correspondence algorithm

Output of acquisition

- Range image (or depth image),
- Disparity map,
- Cloud of point,
- Distribution of normals,
- Camera positions,
- Color/Texture (i.e., Uv-map),
- Poligonal mesh,

Range image



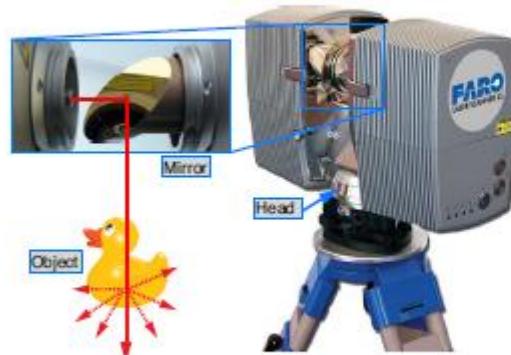
Intensity map



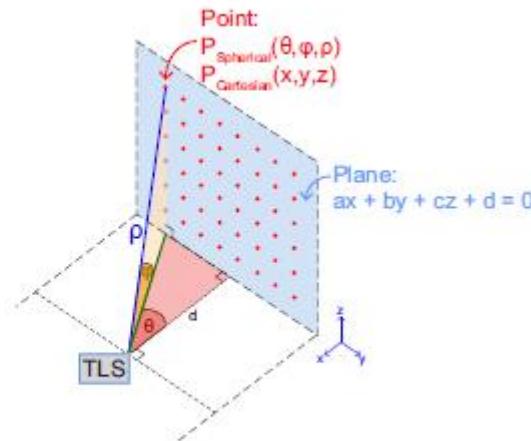
Range (of depth) map

Range map: each pixel encodes a depth value (i.e., the distance from sensor to the object in the scene). Lighter pixels corresponds to closer objects (and viceversa).

Range image

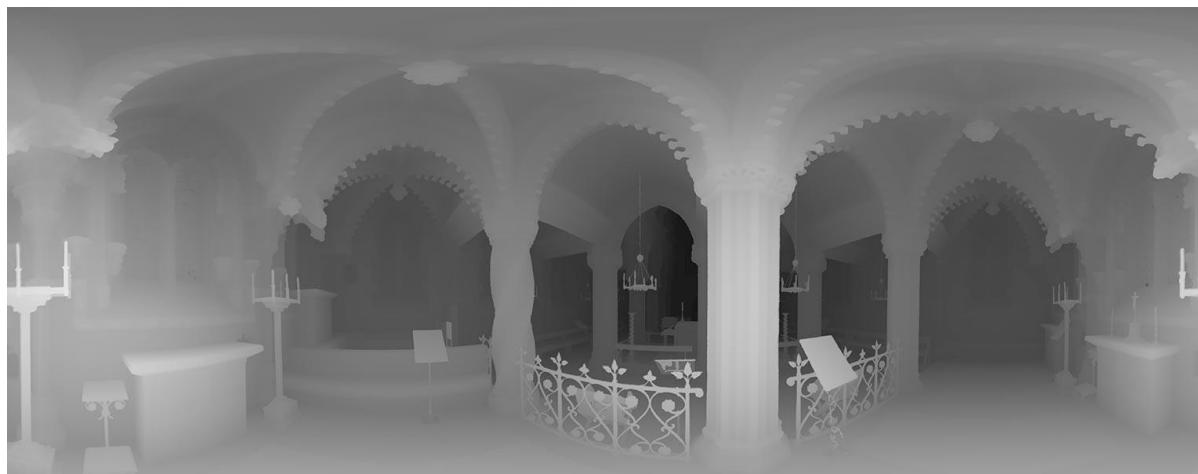


(a) FARO LS880 Phase based Terrestrial Laser Scanner

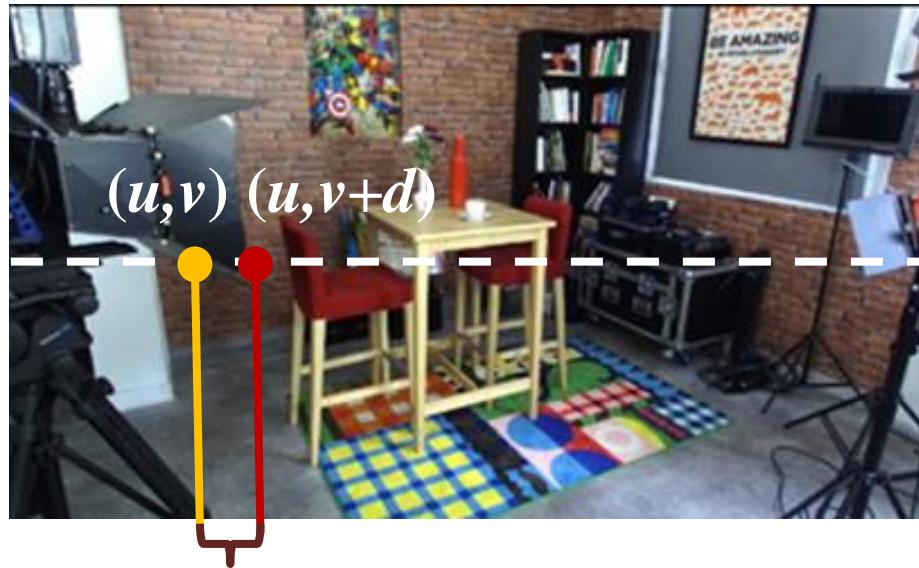


(b) Schematic representation of a point cloud acquisition. A point P is measured at the angular positions (θ, φ) , with a range ρ .

Spherical Scanning



Disparity map

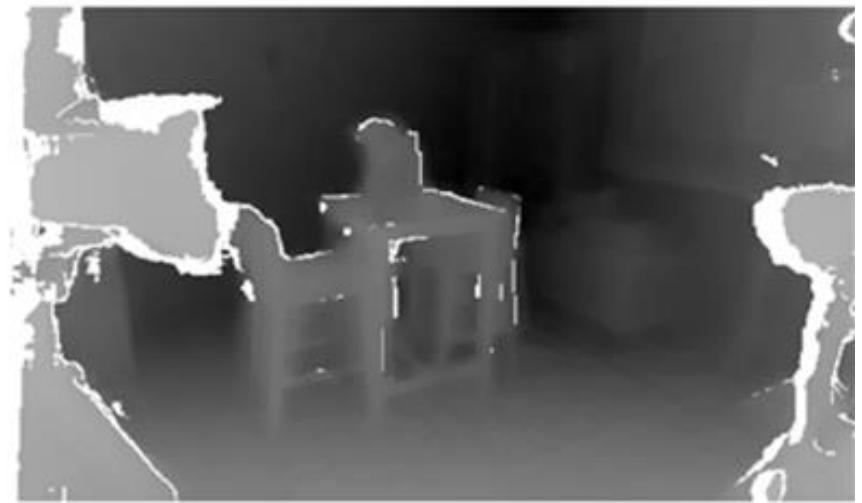


Disparity value d

Disparity map



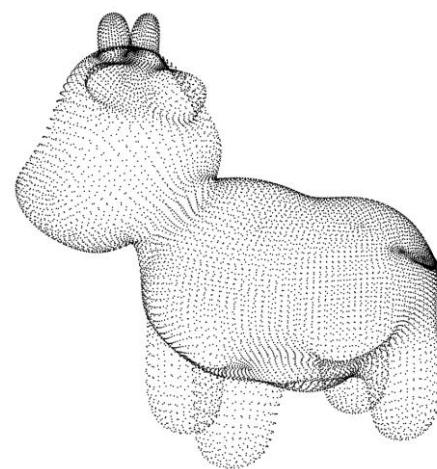
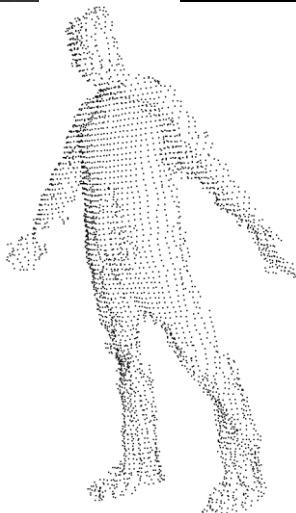
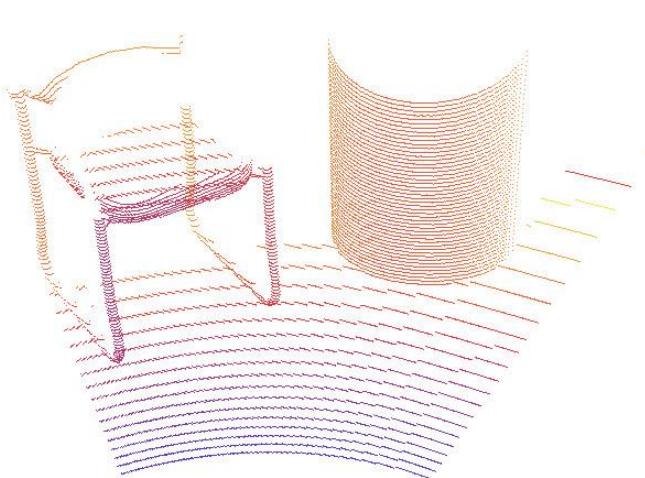
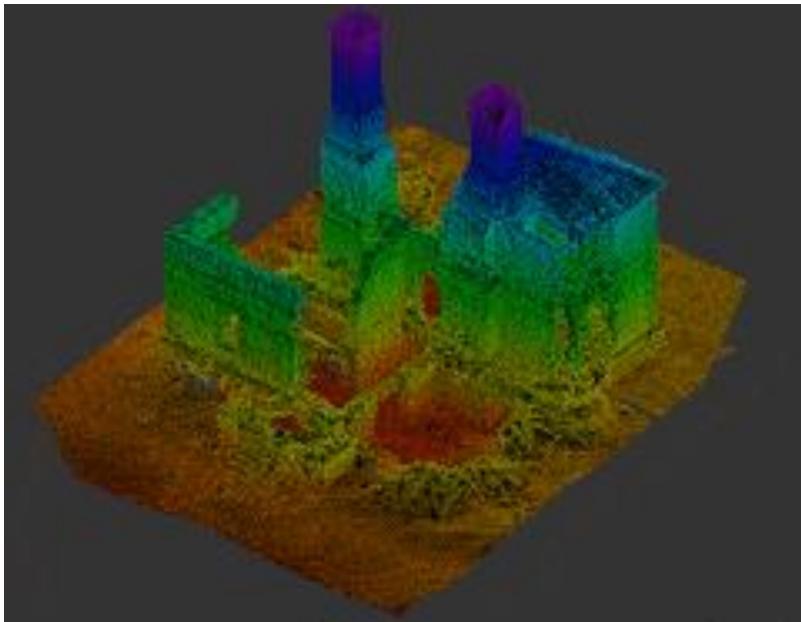
Intensity map



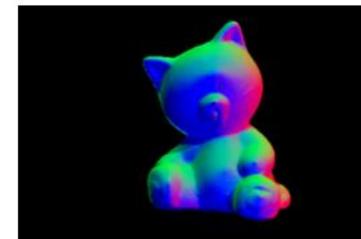
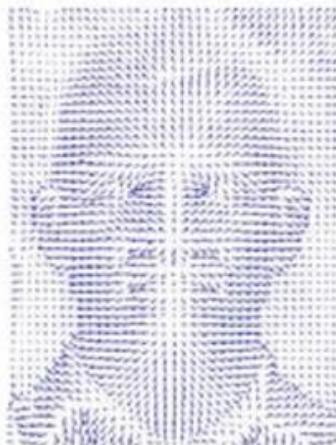
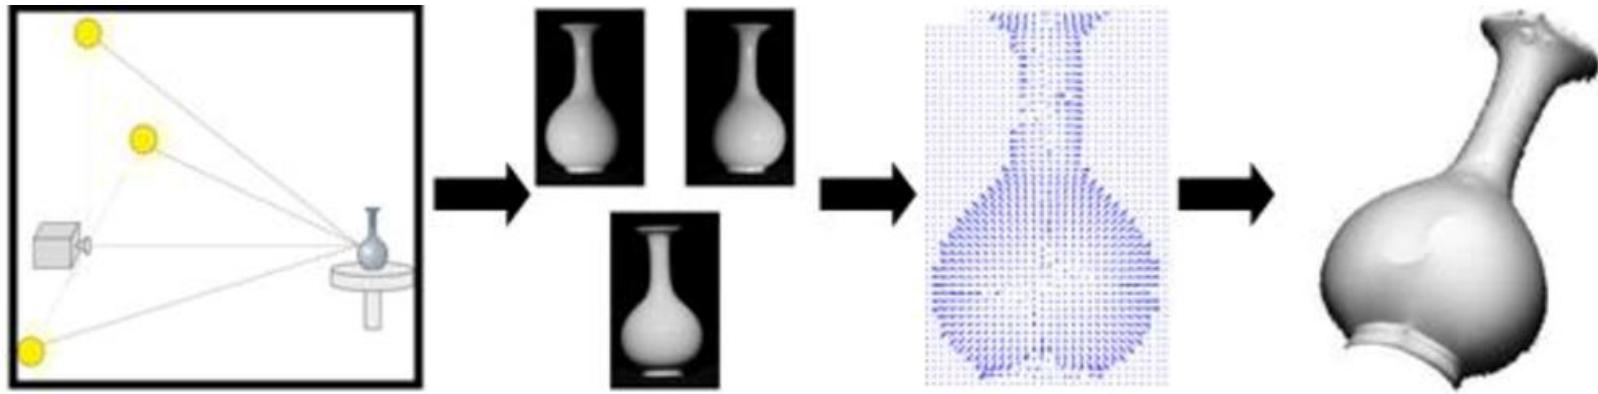
Disparity map

Disparity map: each pixel encodes a disparity value. Lighter pixels corresponds to closer objects (and viceversa).

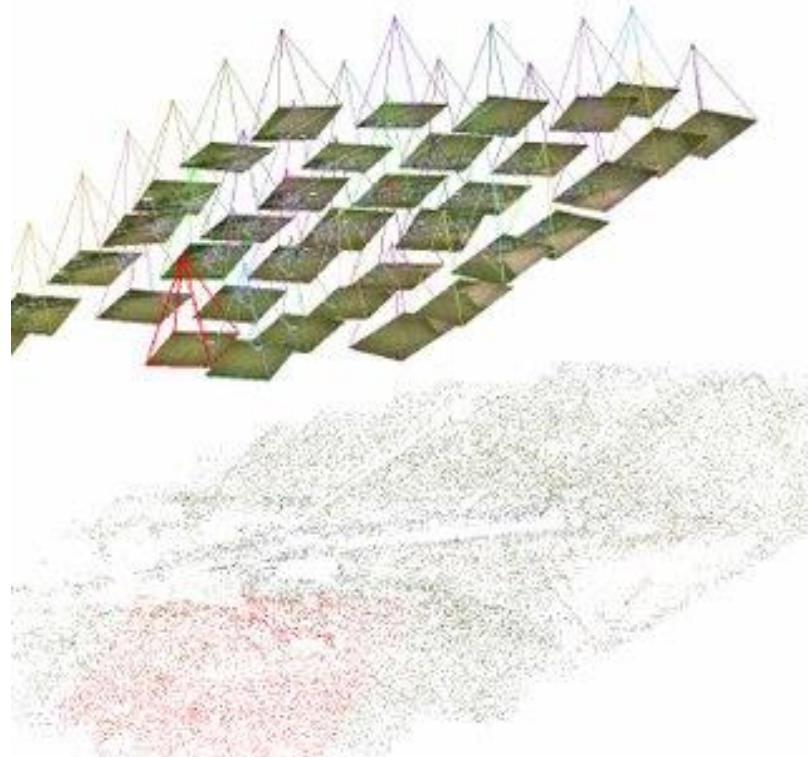
Cloud of points



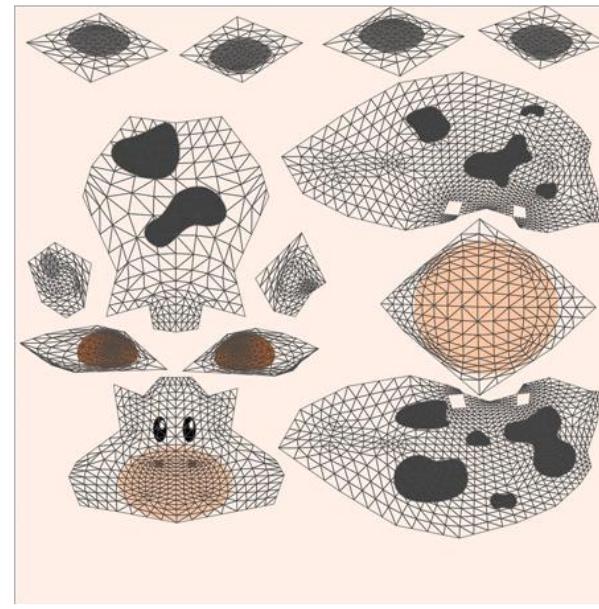
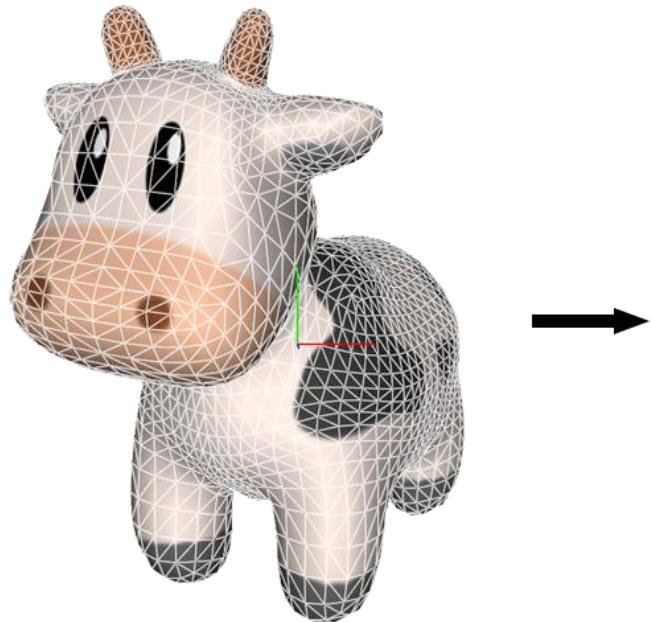
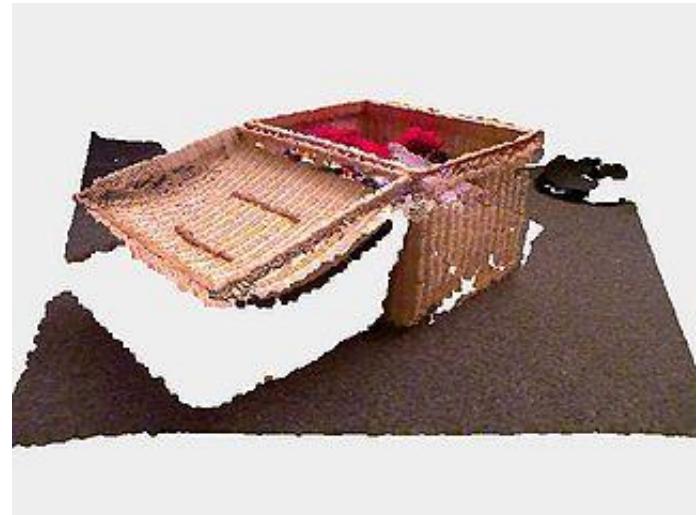
Distribution of normals



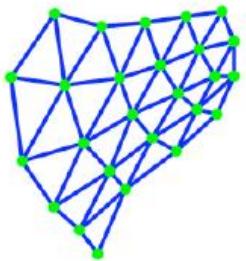
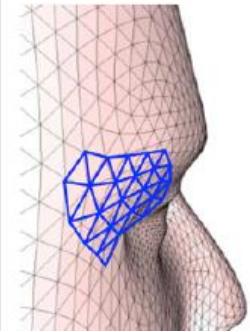
Camera positions



Color and texture



Polygonal meshes



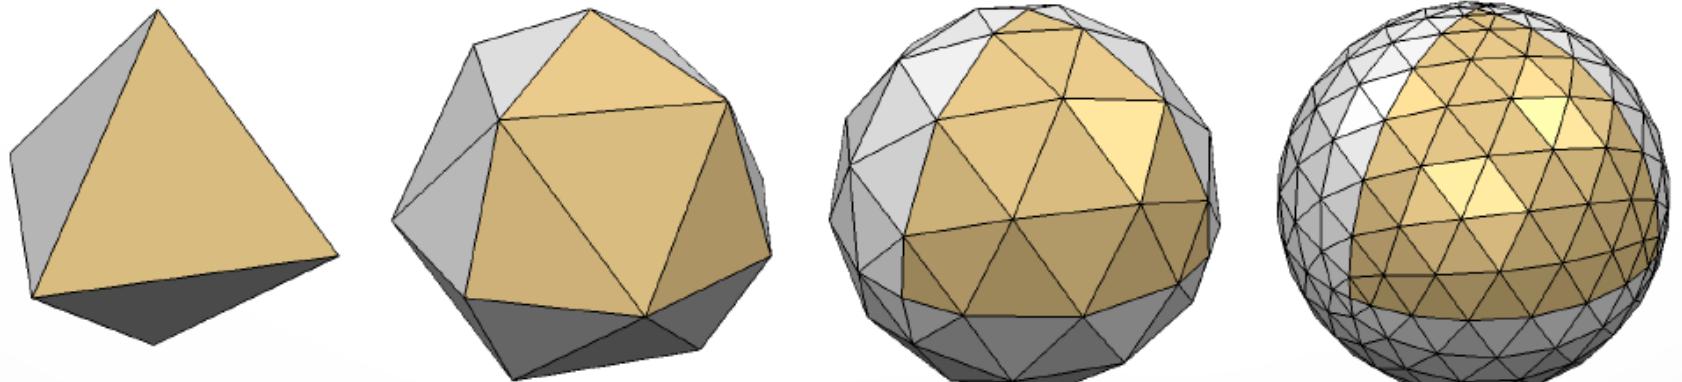
surface mesh: set of vertices, edges and faces (polygons) defining a polyhedral surface in embedded in 3D (discrete approximation of a shape)

- Polygonal meshes are the most used representation for 3D shapes

Polygonal meshes

- Polygonal meshes are the most used representation for 3D shapes:

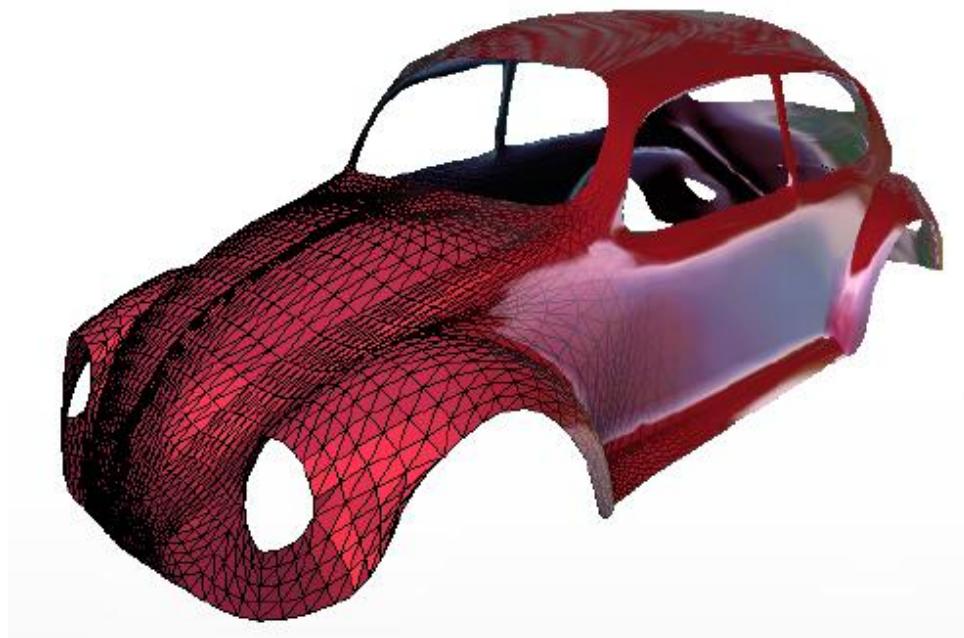
“Piecewise linear approximation --> error is $O(h^2)$ ”



Error inversely proportional to #faces

Polygonal meshes

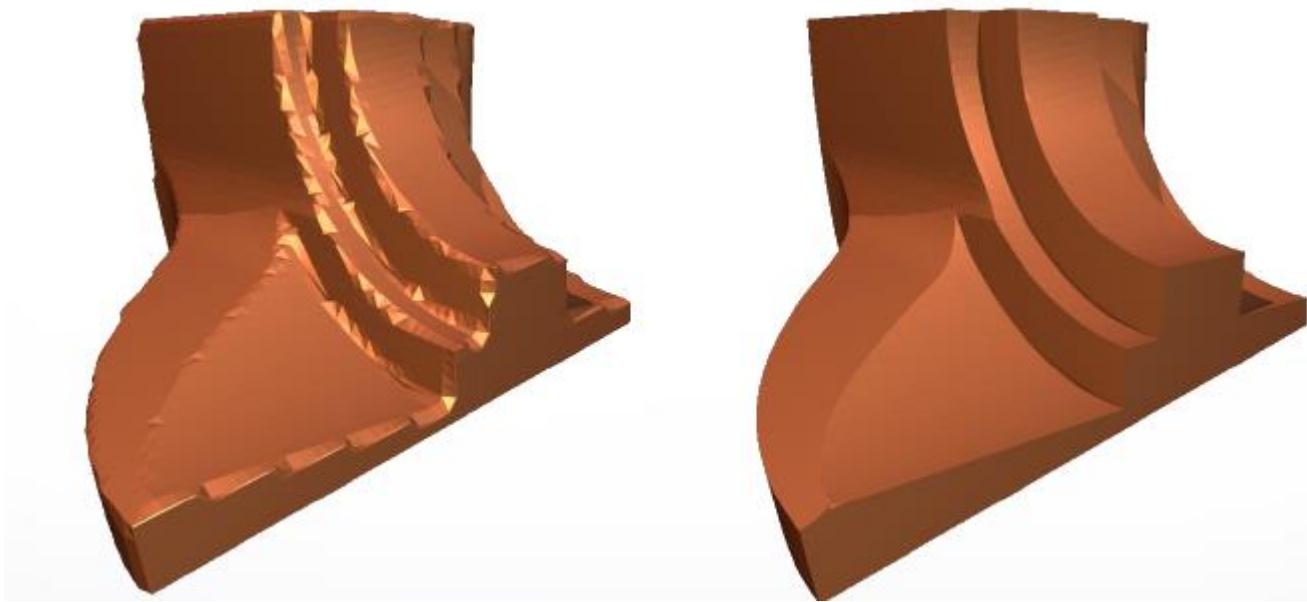
- Polygonal meshes are the most used representation for 3D shapes:
“Arbitrary topology surfaces”



Allow subdivision for smoothness (today)

Polygonal meshes

- Polygonal meshes are the most used representation for 3D shapes:
“Piecewise smooth surfaces”

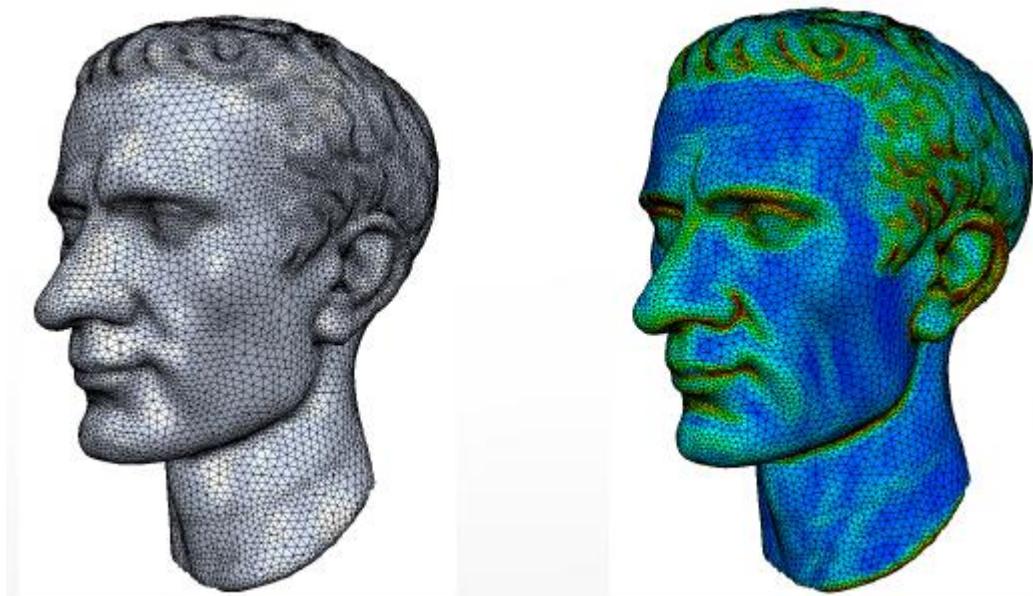


No need to have special rules for joining different patches

Polygonal meshes

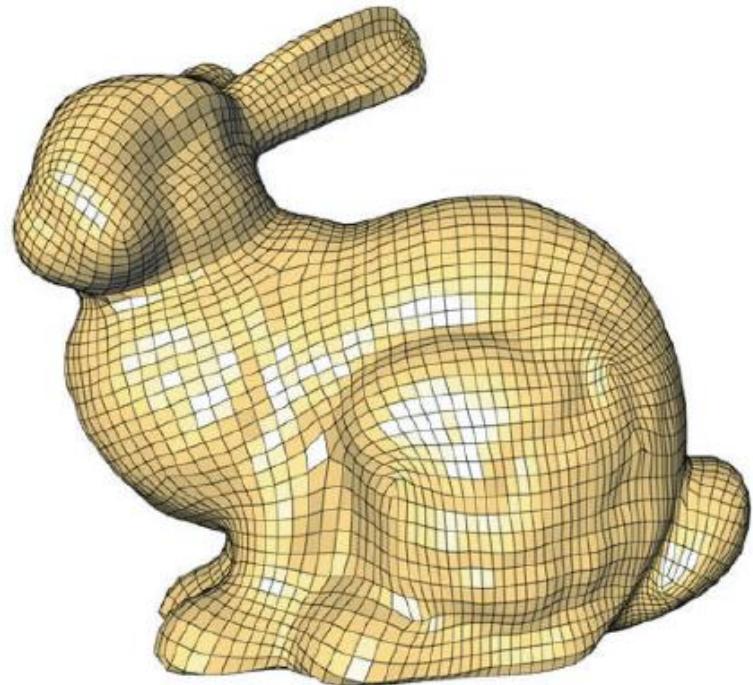
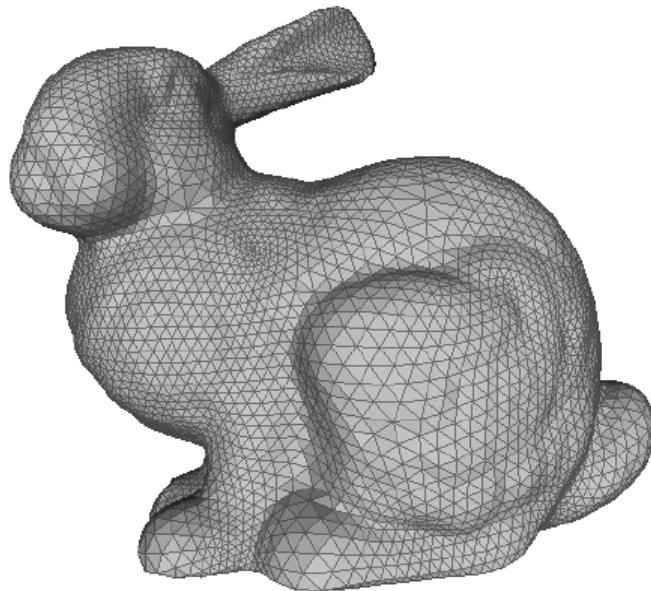
- Polygonal meshes are the most used representation for 3D shapes:

“Adaptive sampling”



Can add resolution only where necessary

Triangular meshes



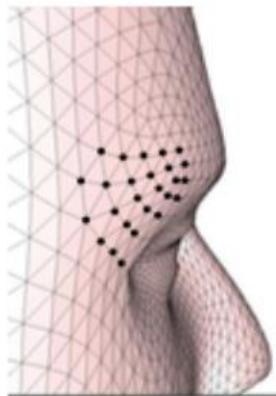
- Simplest piecewise linear element
- Easy to reconstruct from point clouds
- Easy to represent

- Quad meshes often used in animation
- Typically require some handtuning in reconstruction
- Can provide more flexibility for deformation

Triangular mesh

- Two main **components**:

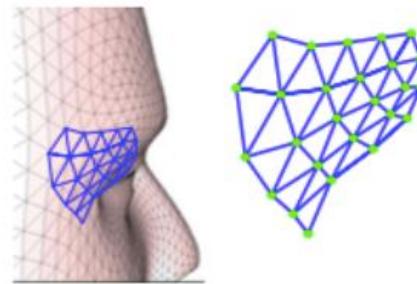
Geometry



vertex
coordinates

geometric structure

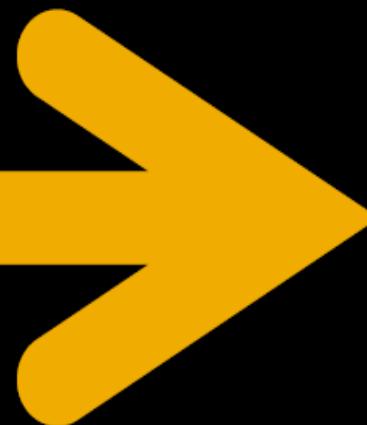
"Connectivity": the underlying triangulation



incidence relations
between triangles,
vertices and edges

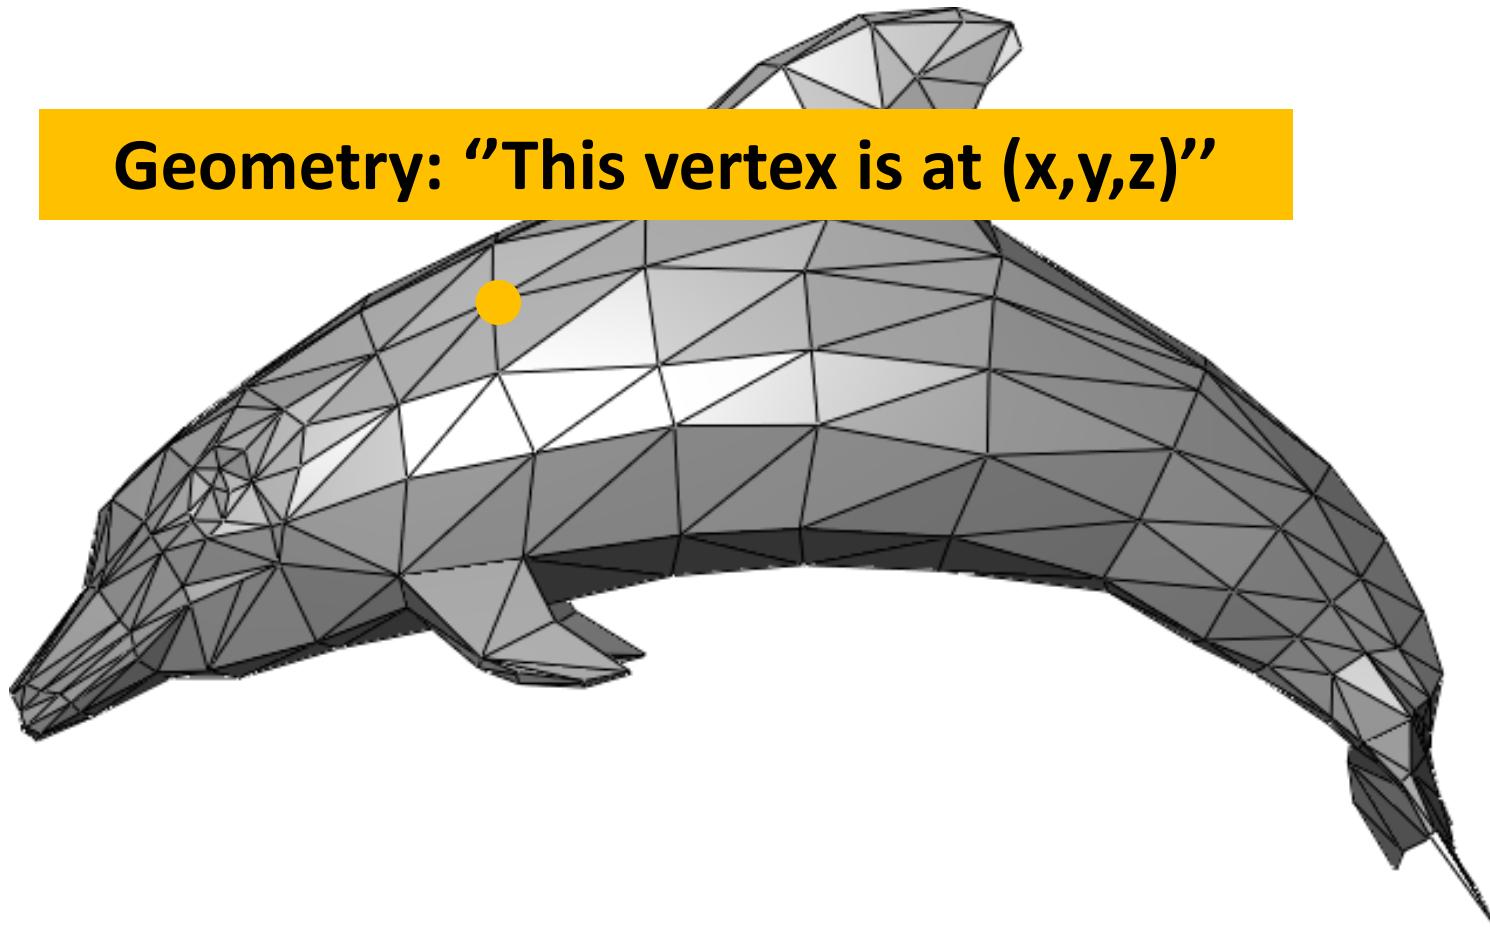
combinatorial structure

Topology [tuh-pol-uh-jee]:
The study of geometric
properties that remain
invariant under certain
transformations

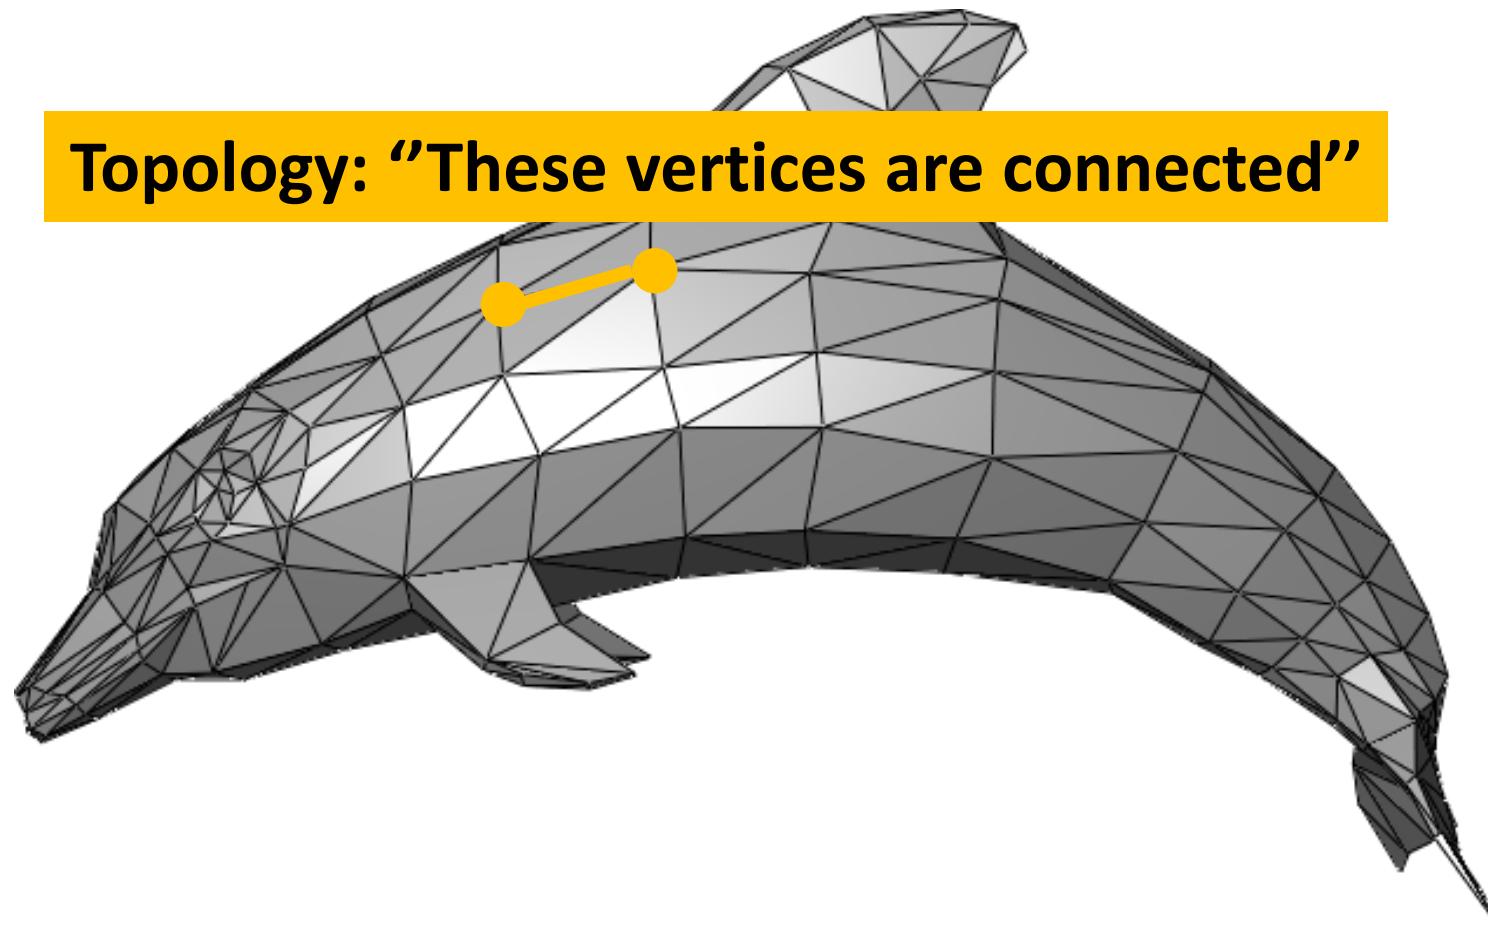


Geometry vs. Topology

Geometry: “This vertex is at (x,y,z) ”



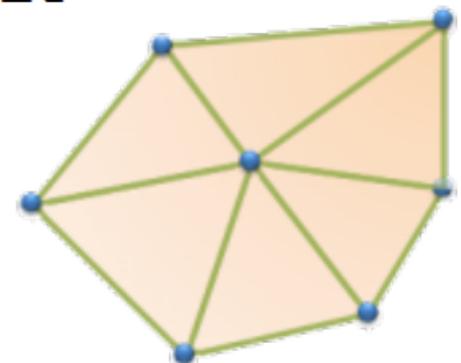
Geometry vs. Topology



Two components

- Geometry: vertex positions

$$\mathcal{P} = \{p_1, p_2, \dots, p_n\}, \quad p_i \in \mathbb{R}^3$$



- Connectivity:

- Vertices: $\mathcal{V} = \{v_1, v_2, \dots, v_n\}$

- Edges: $\mathcal{E} = \{e_1, e_2, \dots, e_m\}, \quad e_i \in \mathcal{V} \times \mathcal{V}$

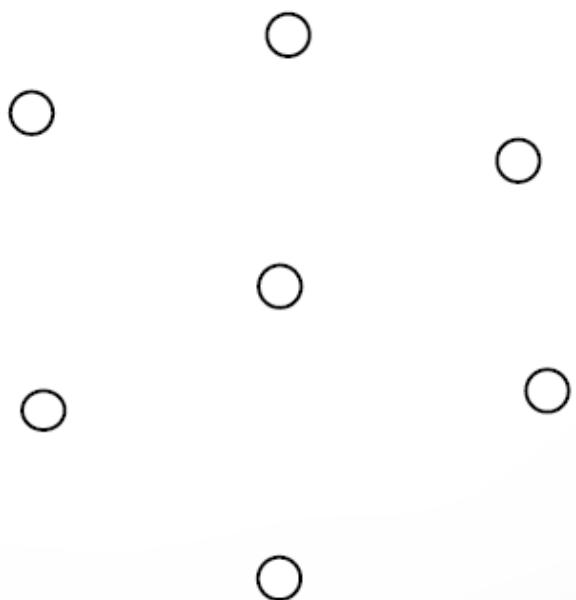
- Faces: $\mathcal{F} = \{f_1, f_2, \dots, f_k\}, \quad f_i \in \mathcal{V} \times \mathcal{V} \times \mathcal{V}$

Two components

$$\mathcal{M} = (\{\mathbf{v}_i\}, \{e_j\}, \{f_k\})$$



geometry $\mathbf{v}_i \in \mathbb{R}^3$



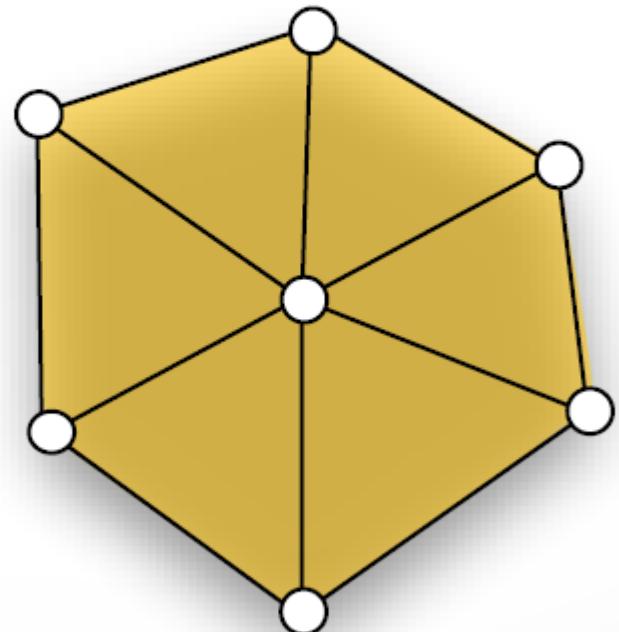
Two components

$$\mathcal{M} = (\{\mathbf{v}_i\}, \{e_j\}, \{f_k\})$$

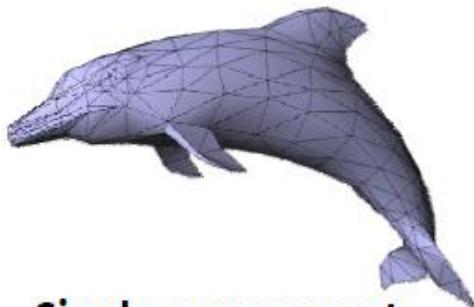


geometry $\mathbf{v}_i \in \mathbb{R}^3$

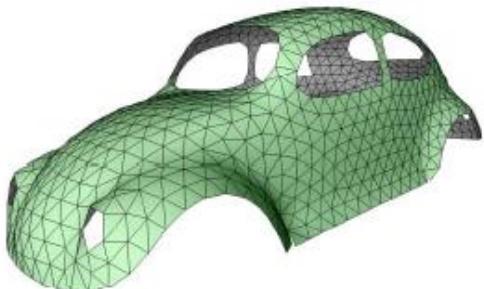
topology $e_i, f_i \subset \mathbb{R}^3$



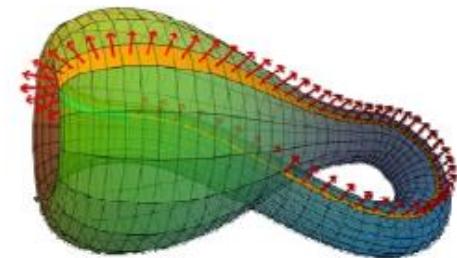
Mesh Zoo



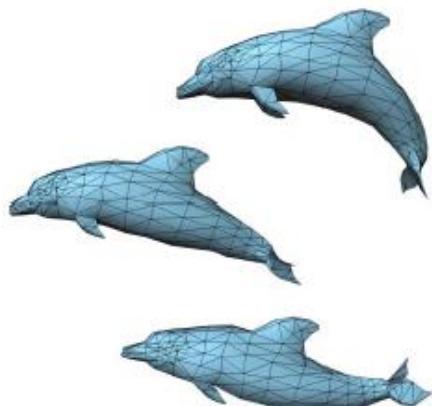
Single component,
closed, triangular,
orientable manifold



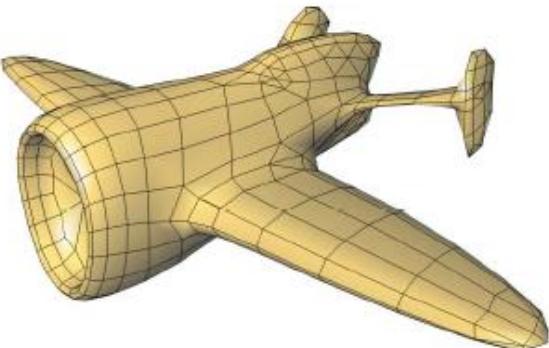
With boundaries



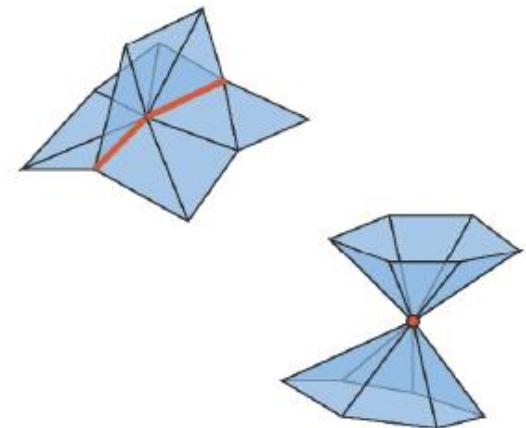
Not orientable



Multiple components



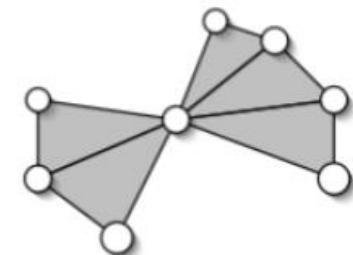
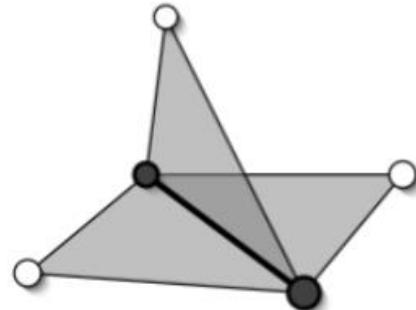
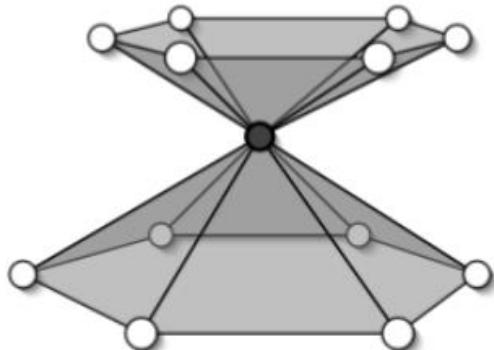
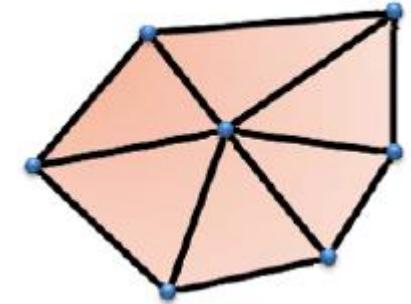
Not only triangles



Non manifold

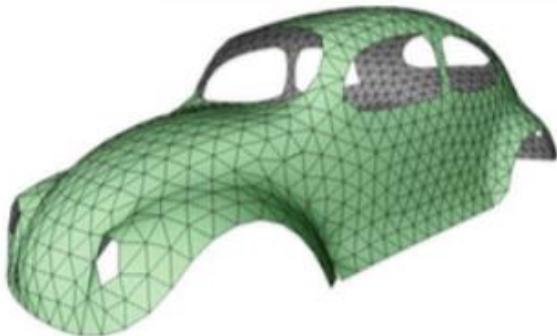
Triangular manifold mesh

- Each Edge is adjacent to at most 2 faces
- Each vertex has a disk-shaped neighborhood

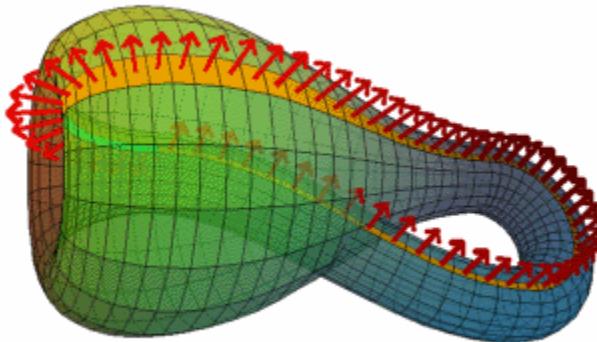


Non-manifold

With boundaries, Not orientable



With boundaries

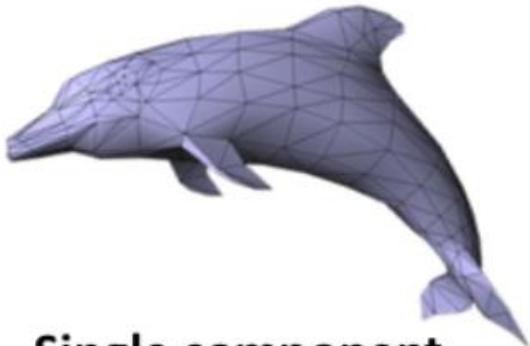


Not orientable

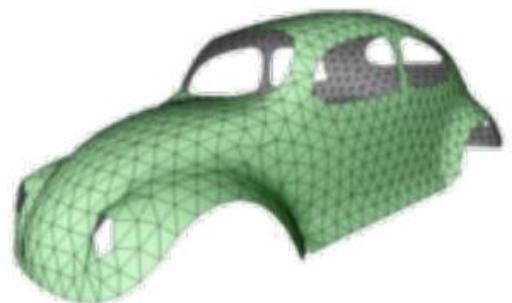
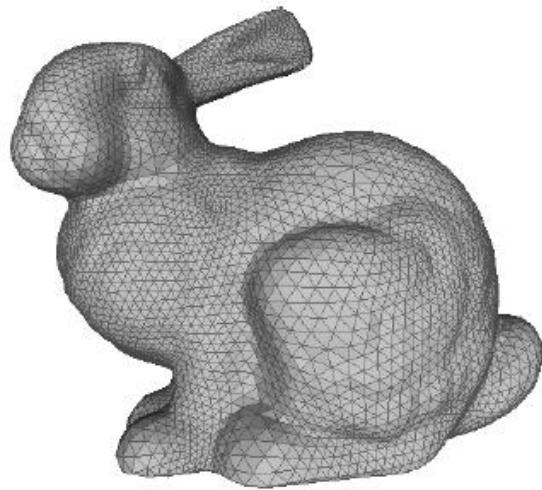
Boundary edge: adjacent to exactly one face

Orientable surface: possible to assign a consistent normal orientation (e.g. outward)

Single component



**Single component,
closed, triangular,
orientable manifold**



Possibly with
boundaries

Closed (without boundaries), Single component, 2D-manifold, is often called **Watertight mesh**

File Formats

SITUATION:
THERE ARE
14 COMPETING
STANDARDS.

14?! RIDICULOUS!
WE NEED TO DEVELOP
ONE UNIVERSAL STANDARD
THAT COVERS EVERYONE'S
USE CASES.

YEAH!



SOON:

SITUATION:
THERE ARE
15 COMPETING
STANDARDS.

File Formats

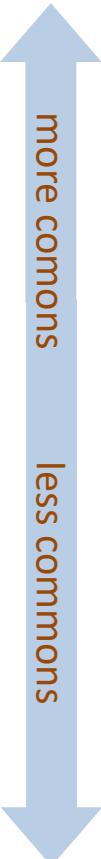
- 3DS - 3D Studio Max file format
- OBJ - Another file format for 3D objects
- MA, MB - Maya file formats
- 3DX - Rinoceros file format
- BLEND - Blender file format
- DAE - COLLADA file format (Khronos)
- FBX - Autodesk interchange file format
- X - Direct X object
- SMD - good for animations (by Valve)
- MD3 - quake 3 vertex animations
- DEM - Digital Elevation Models
- DXF - (exchange format, Autodesk's AutoCAD)
- FIG - Used by REND386/AVRIL
- FLT - MultGen Inc.'s OpenFlight format
- HDF - Hierarchical Data Format
- IGES - Initial Graphics Exchange Specification
- IV - Open Inventor File Format Info
- LWO, LWB & LWS - Lightwave 3D file formats
- MAZ - Used by Division's dVS/dVISE
- MGF - Materials and Geometry Format
- MSDL - Manchester Scene Description Language
- 3DML - by Flatland inc.
- C4D – Cinema 4D file format
- SLDPTR - SolidWork "part"
- WINGS - Wings3D object
- NFF - Used by Sense8's WorldToolKit
- SKP - Google sketch up
- KMZ - Google Earth model
- OFF - A general 3D mesh Object File Format
- OOGL - Object Oriented Graphics Library
- PLG - Used by REND386/AVRIL
- POV - “persistence of vision” ray-tracer
- QD3D - Apple's QuickDraw 3D Metafile format
- TDDD - for Imagine & Turbo Silver ray-tracers
- NFF & ENFF - (Extended) Neutral File Format
- VIZ - Used by Division's dVS/dVISE
- VRML, VRML97 - Virtual Reality Modeling Language (RIP)
- X3D - tentato successore di VRML
- PLY - introdotto by Cyberware – tipic. dati range scan
- DICOM - Dalla casa omonima – tipic. dati volum, es CAT scan
- Renderman - per l'omonimo visualizzatore
- RWX - RenderWare Object
- Z3D - ZModeler File format
- etc, etc, etc...

File formats for 3D meshes

(Tower of Babel!)

- Different Types
 - Usually triangular or polygonal meshes
 - Sometimes contains also scene description
 - Cameras, model-view matrices, hierarchies...
 - Even with Polygonal mesh many variations
 - Direct, Indexed, Half edges
 - Per face, per vertex attributes
- Some are open standards, others are commercial soft
 - Some are easy to parse, others very difficult
 - Some of them pretend to be very *generic*: «3D esperanto»
- Format
 - ASCII o binary
 - Sometimes also compressed

File formats for 3D meshes



.OBJ (wavefront)

- ☺ Most common
- ☺ indexed, normals , uv-mapping
- ☹ no colors (only per face material index)
- ☹ no skinning or animation

.SMD ()

- ☺ Skeletal animations + skinning
- ☺ normals , uv-mapping
- ☹ no indexed
- ☹ no colors

.MD3 (Quake, IDsoft)

- ☺ vertex animations, normals
- ☹ no colors

.PLY (cyberware)

- ☺ extendable
- ☹ “diffuse in academic world”

simple

.3DS ()

- ☺ color on faces/vertices, uv-mapping, indexed, materials, textures...
- ☹ no: normals
- ☹ Maximum number of vertices(64K)

.COLLADA ()

- ☺ super complete!
- ☺ designed to be extendible
- ☺ open standard
- ☹ difficult to parse

.FBX ()

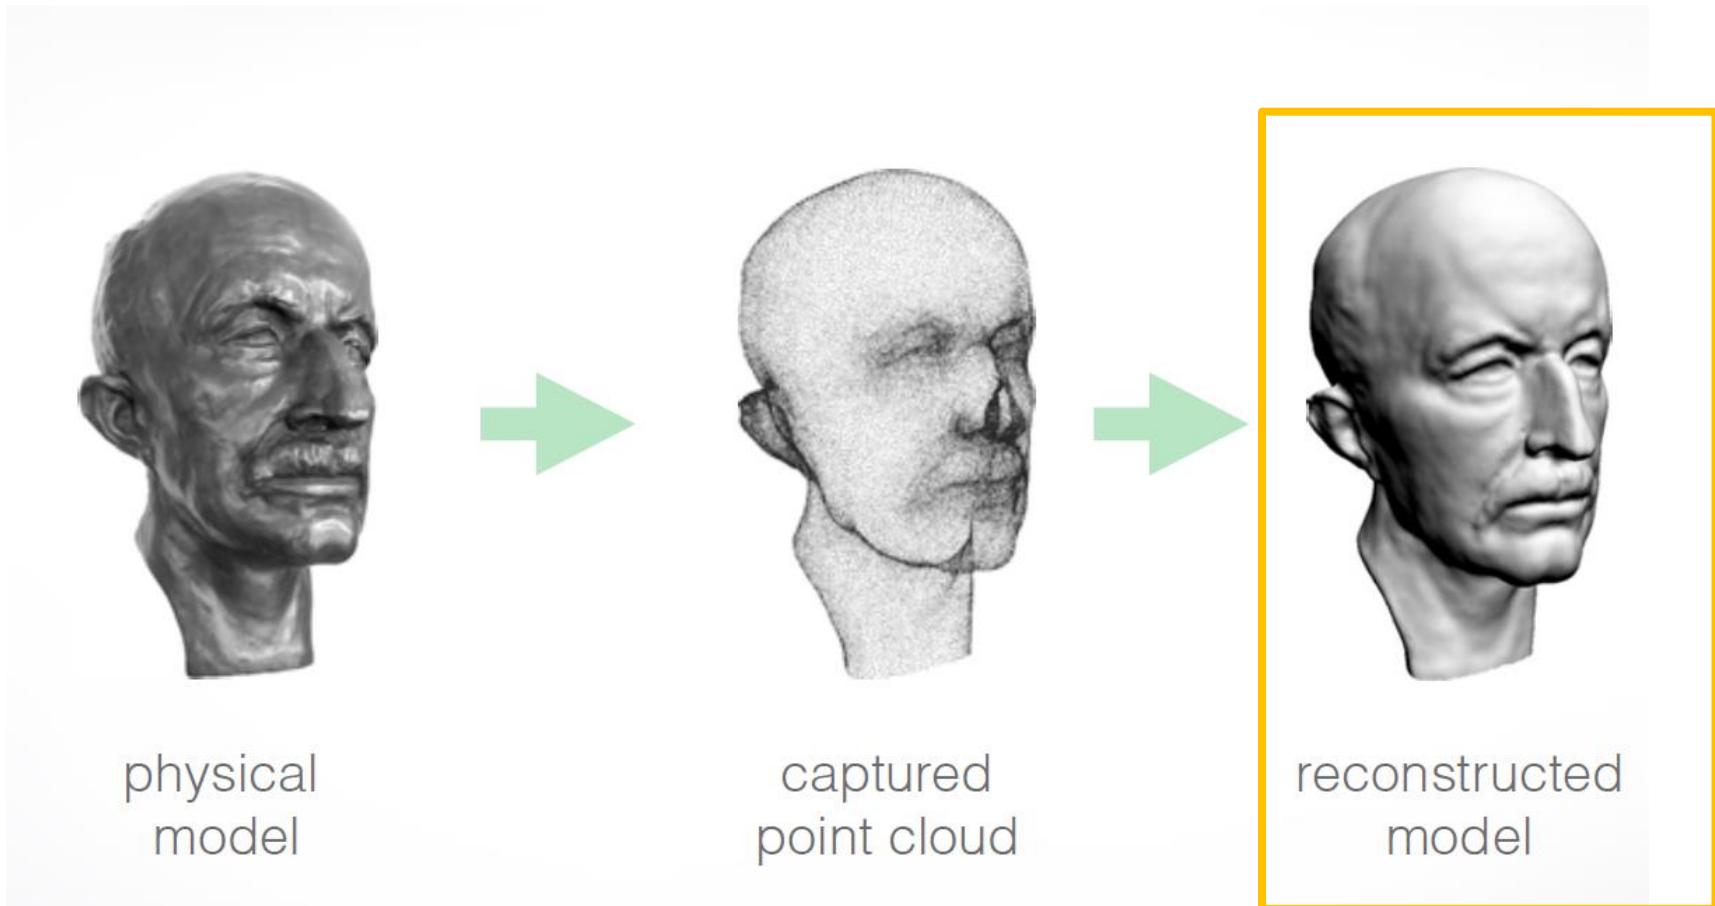
- ☺ complete with animation
- ☹ difficult to parse

.MA / .MB ()

- ☺ complete with animation
- ☹ difficult to parse

complex

Surface reconstruction



Polygonal mesh

Input

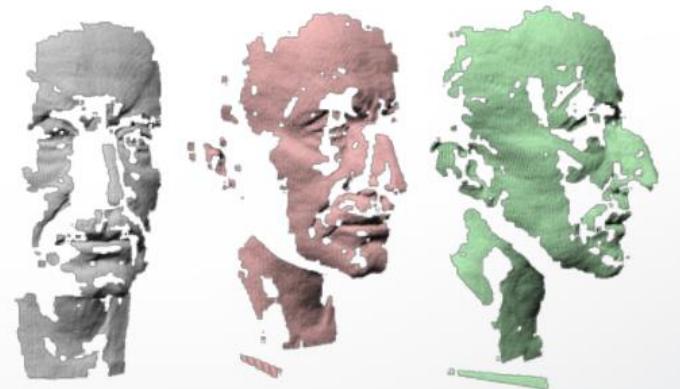
Set of irregular sample points

- with or without normals
- examples: multi-view stereo, union of range scan vertices



Set of range scans

- each scan is a regular quad or tri-mesh
- normal vectors can be obtained through local connectivity



Surface reconstruction

- Two approaches:

Explicit

Local surface
connectivity estimation

Point interpolation

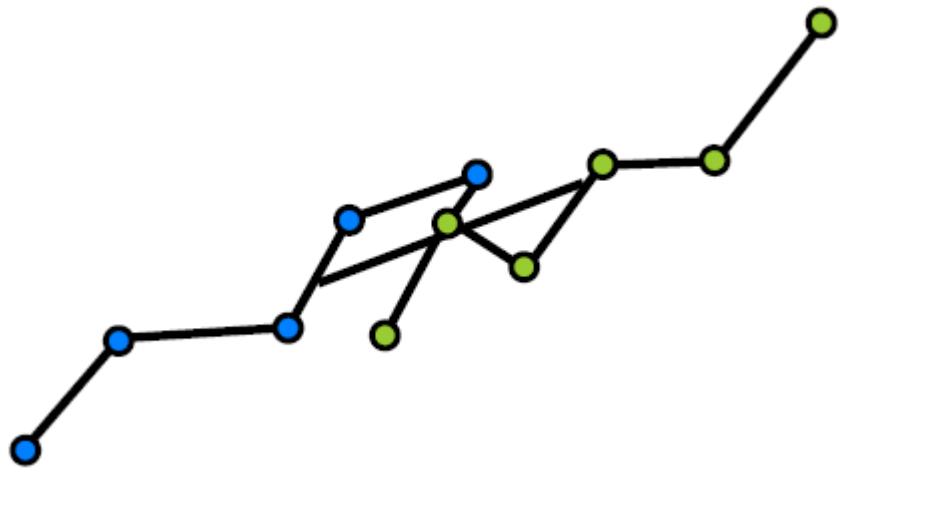
Implicit

Signed distance function
estimation

Mesh approximation

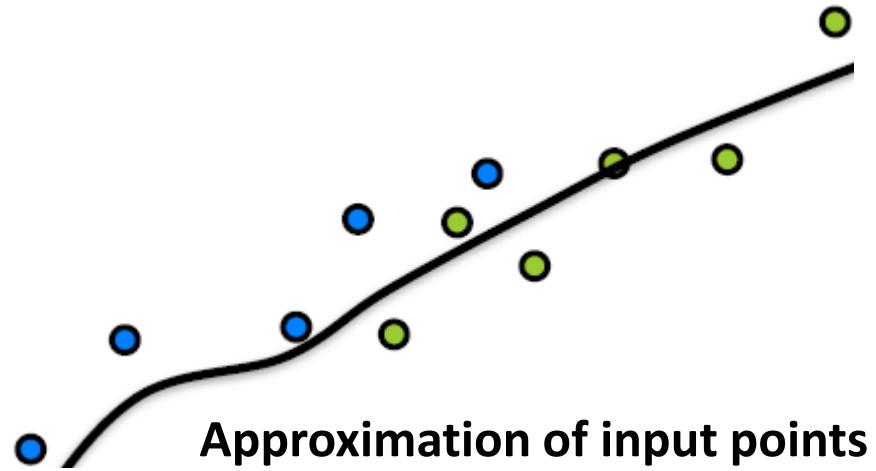
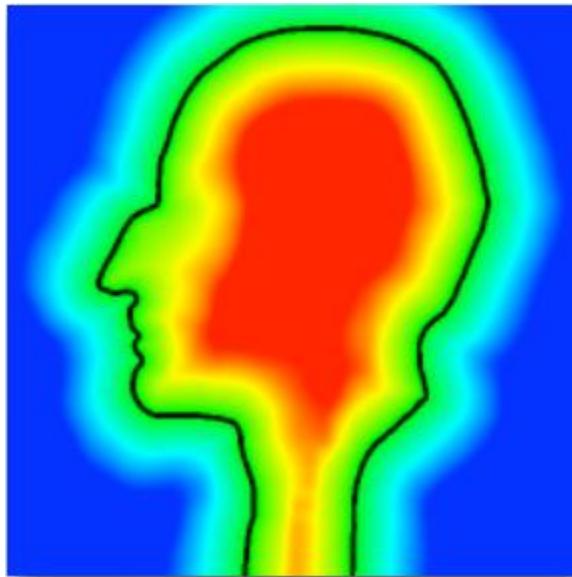
Explicit methods

- Connect sample points by triangles
- Exact interpolation of sample points
- Bad for noisy or misaligned data
- Can lead to holes or non-manifold situations



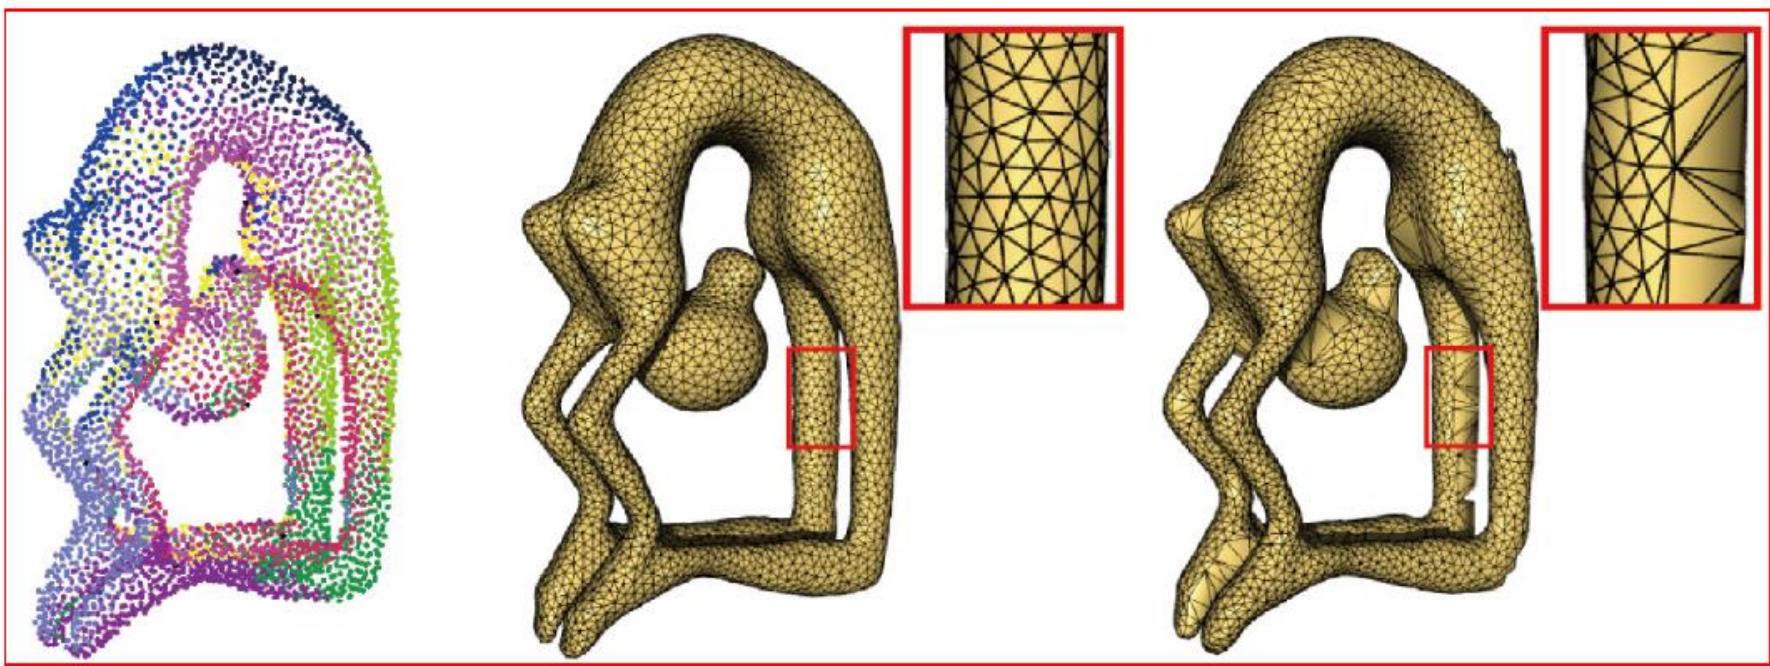
Implicit methods

1. estimate a signed distance function(SDF);
2. extract 0-level set mesh using Marching Cubes



Output is a **Watertight manifold** by construction !

Explicit vs Implicit



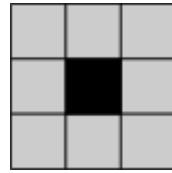
Input

Implicit

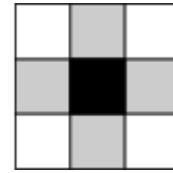
Explicit

Esplicit method

- Mesh reconstruction from **range image**,
 - **Idea:** points are on a regular grid where the connectivity can be inherited from the pixel neighbourhood



8-neigh



4-neigh

- Zippering range scans,
 - **Idea:** “Zipper” several scans to one single model



Mesh from range image

- Points are on a regular grid where the connectivity can be inherited from the pixel neighbourhood,

BUT...

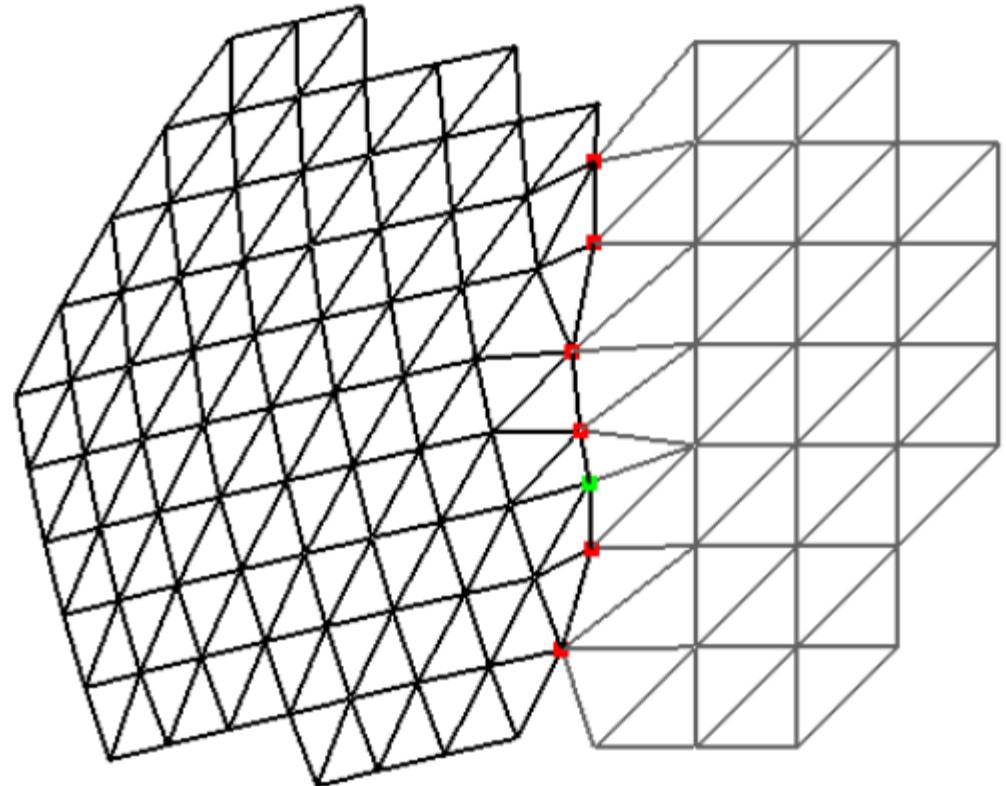
- 1) Not all the pixels on the range image are the projection of a point on the 3D space!
 - a binary mask can be used to define valid points,
- 2) Nearby pixels should not correspond to nearby points on the 3D space!
 - a robust strategy can be used to remove long edges.

See Matlab script available from lab section!

Zippering range scans

- “Zipper” several scans to one single model

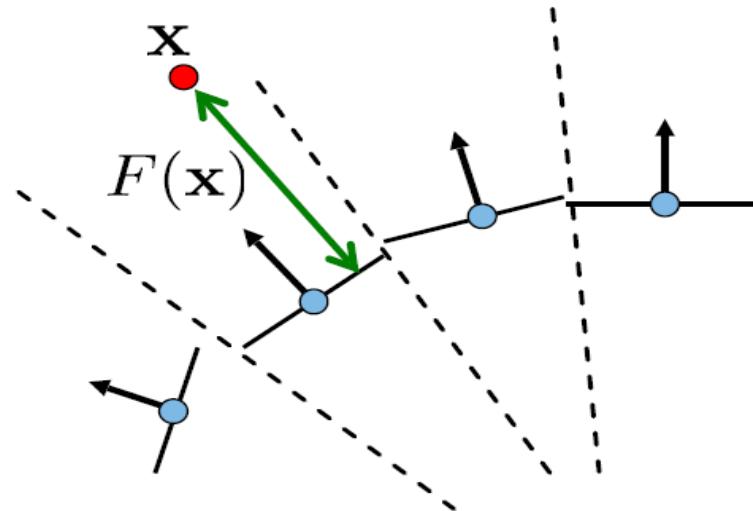
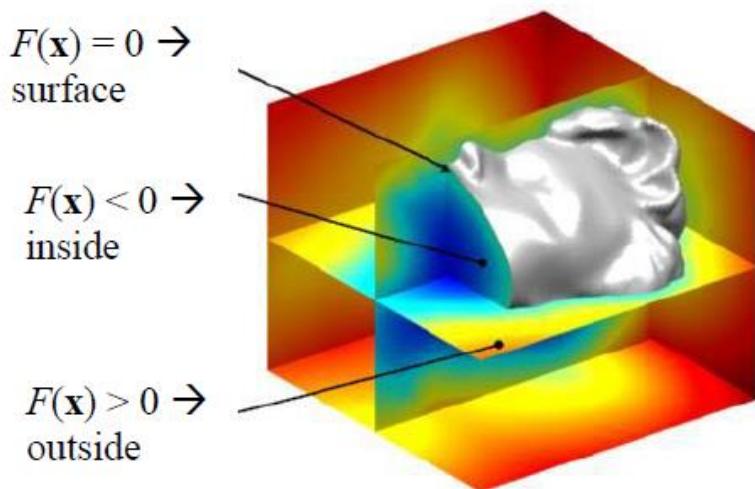
1. Project & insert boundary vertices
2. Intersect boundary edges
3. Discard overlap region
4. Locally optimize triangulation



Not much used in practice!

Implicit methods

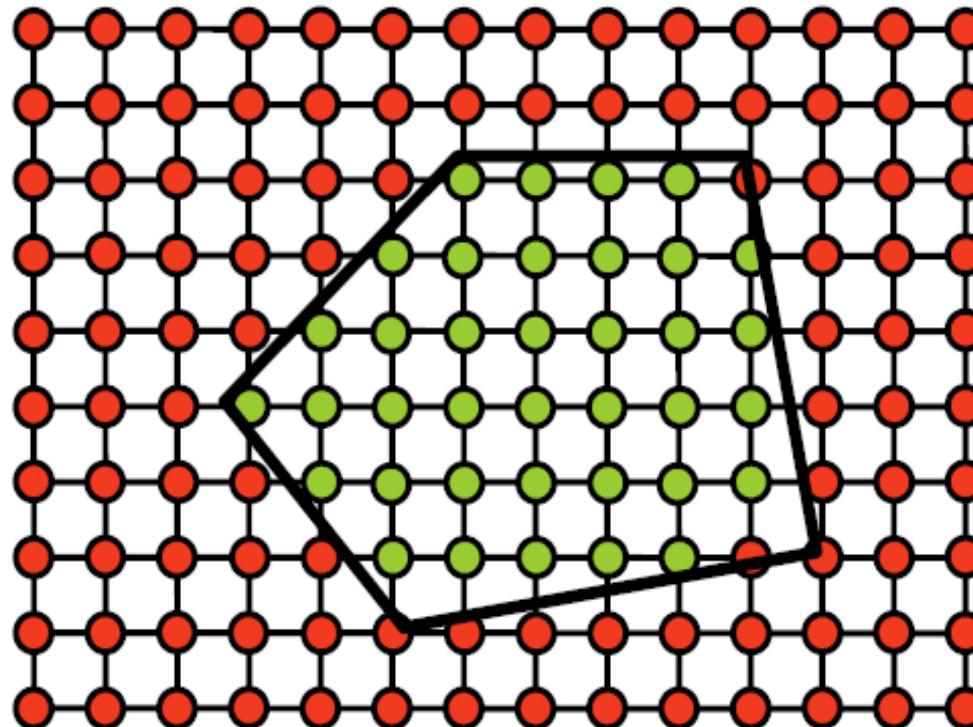
- Several methods for signed distance computation (SDF)
 - **Marching Cube**: classical method
 - **Poisson method**: the currently most used method



E.g.: signed distance to the tangent plane of the closest point

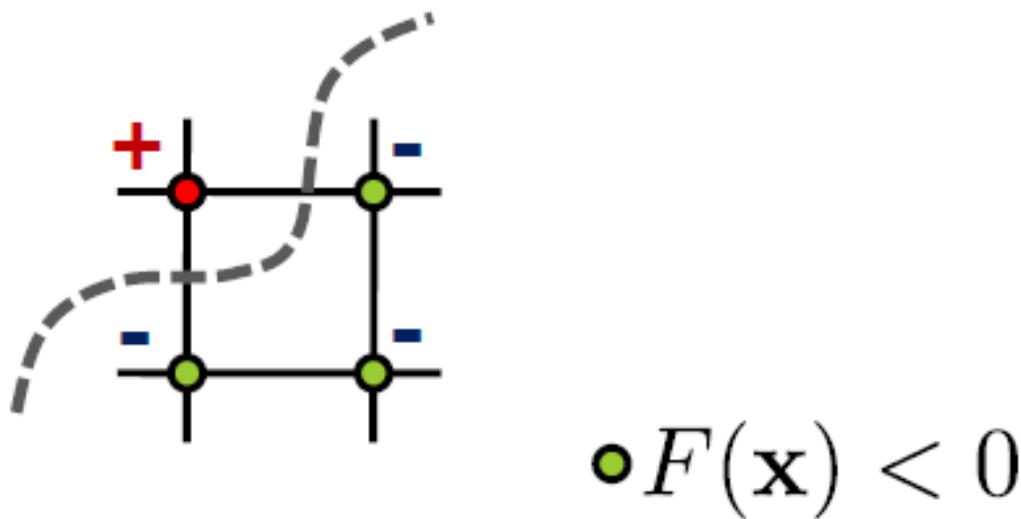
Marching cube

- Idea: sample the SDF



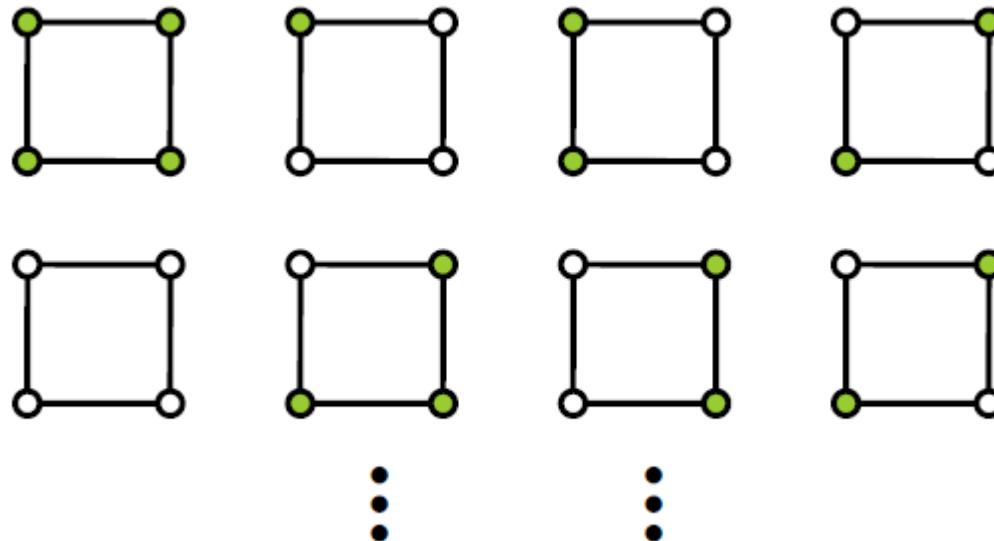
Marching cube

- Idea: sample the SDF



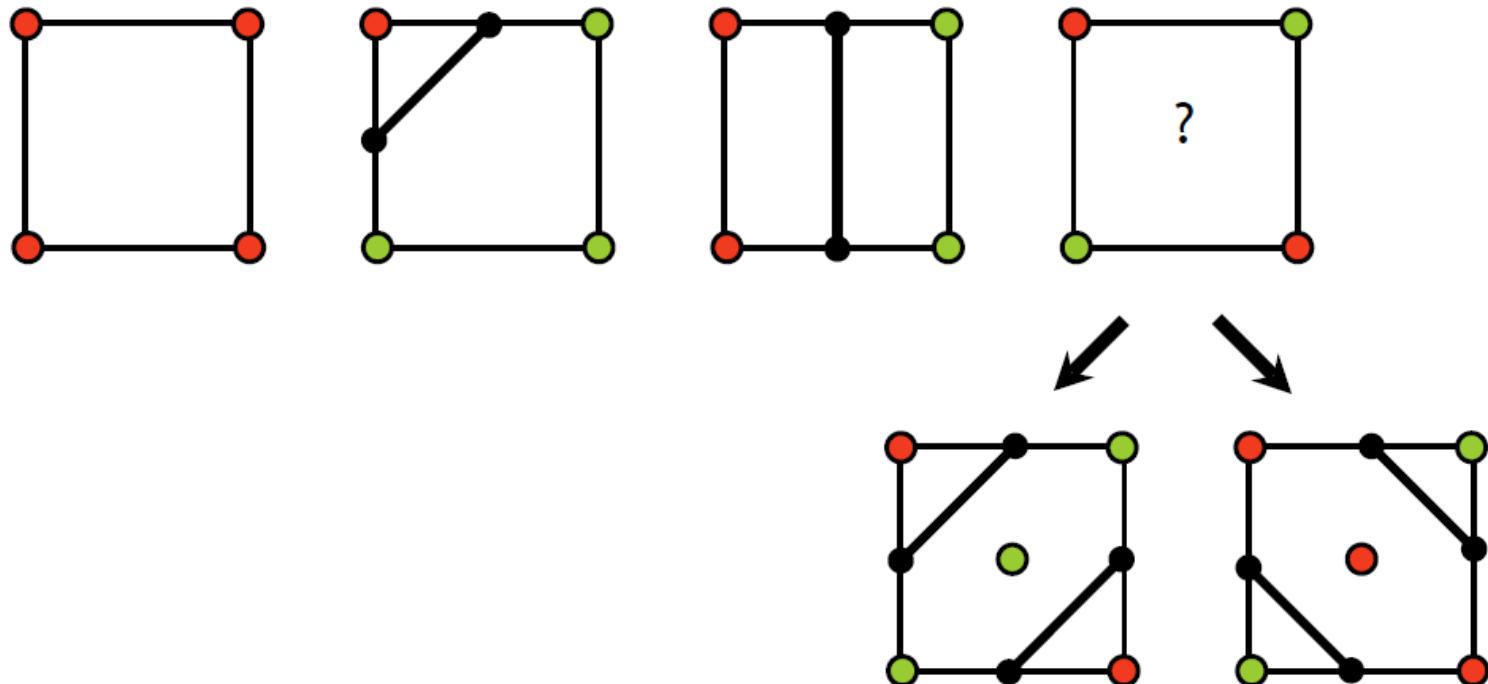
Marching square

- $2^4 = 16$ different configurations in 2D
- 4 equivalence classes (up to rotational and reflection symmetry + complement)

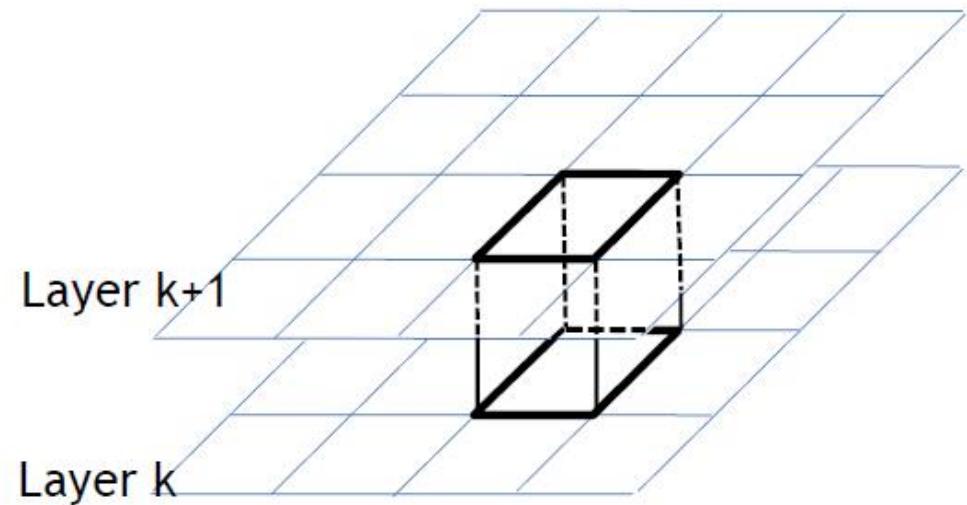
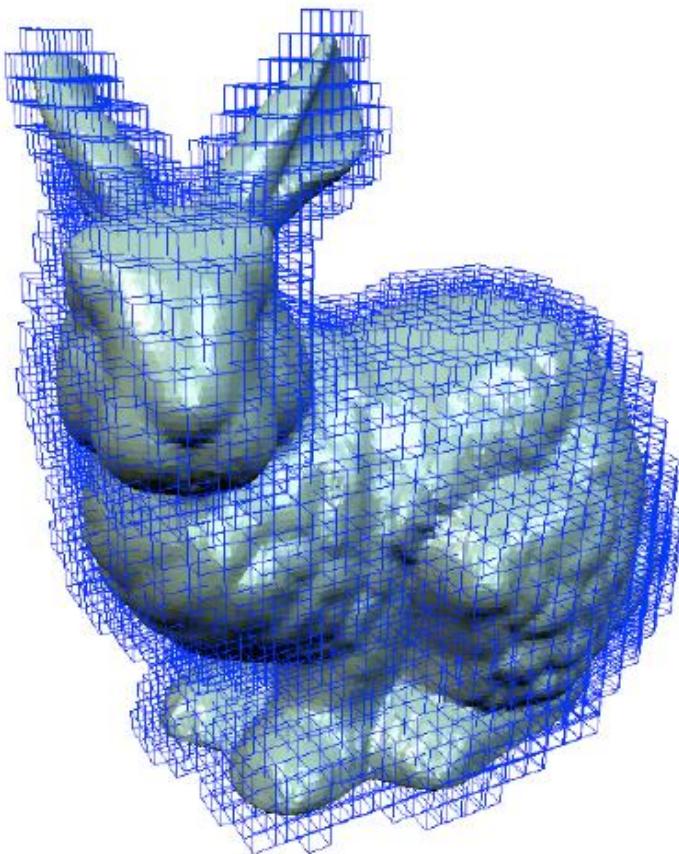


Marching square

- 4 equivalence classes (up to rotational and reflection symmetry + complement)



Marching cube

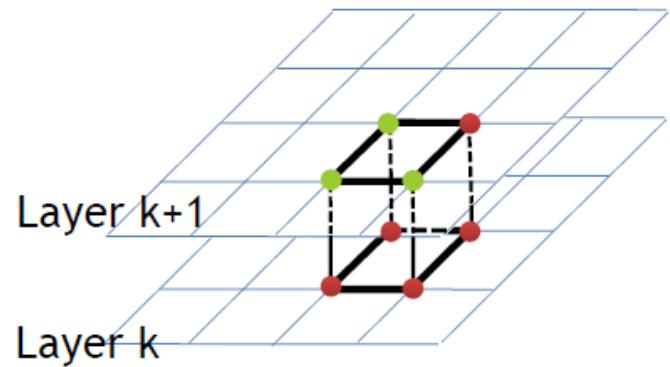


3D case!

Marching cube

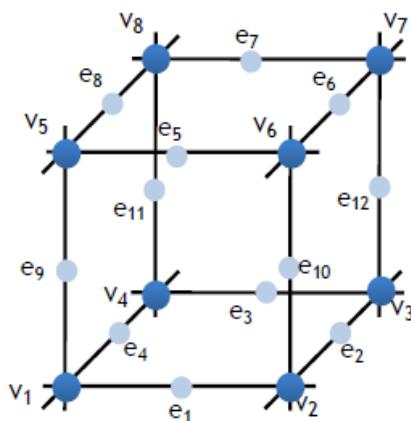
Marching Cubes (Lorensen and Cline 1987)

1. Load 4 layers of the grid into memory
2. Create a cube whose vertices lie on the two middle layers
3. Classify the vertices of the cube according to the implicit function (inside, outside or on the surface)

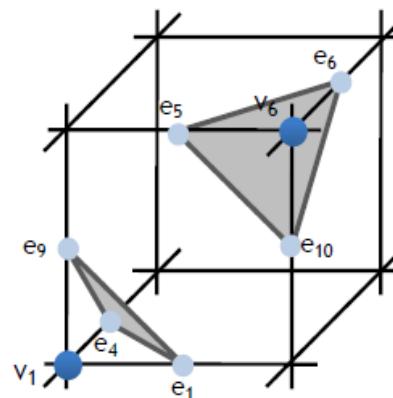


Marching cube

Compute case index. We have $2^8 = 256$ cases (0/1 for each of the eight vertices) – can store as 8 bit (1 byte) index.



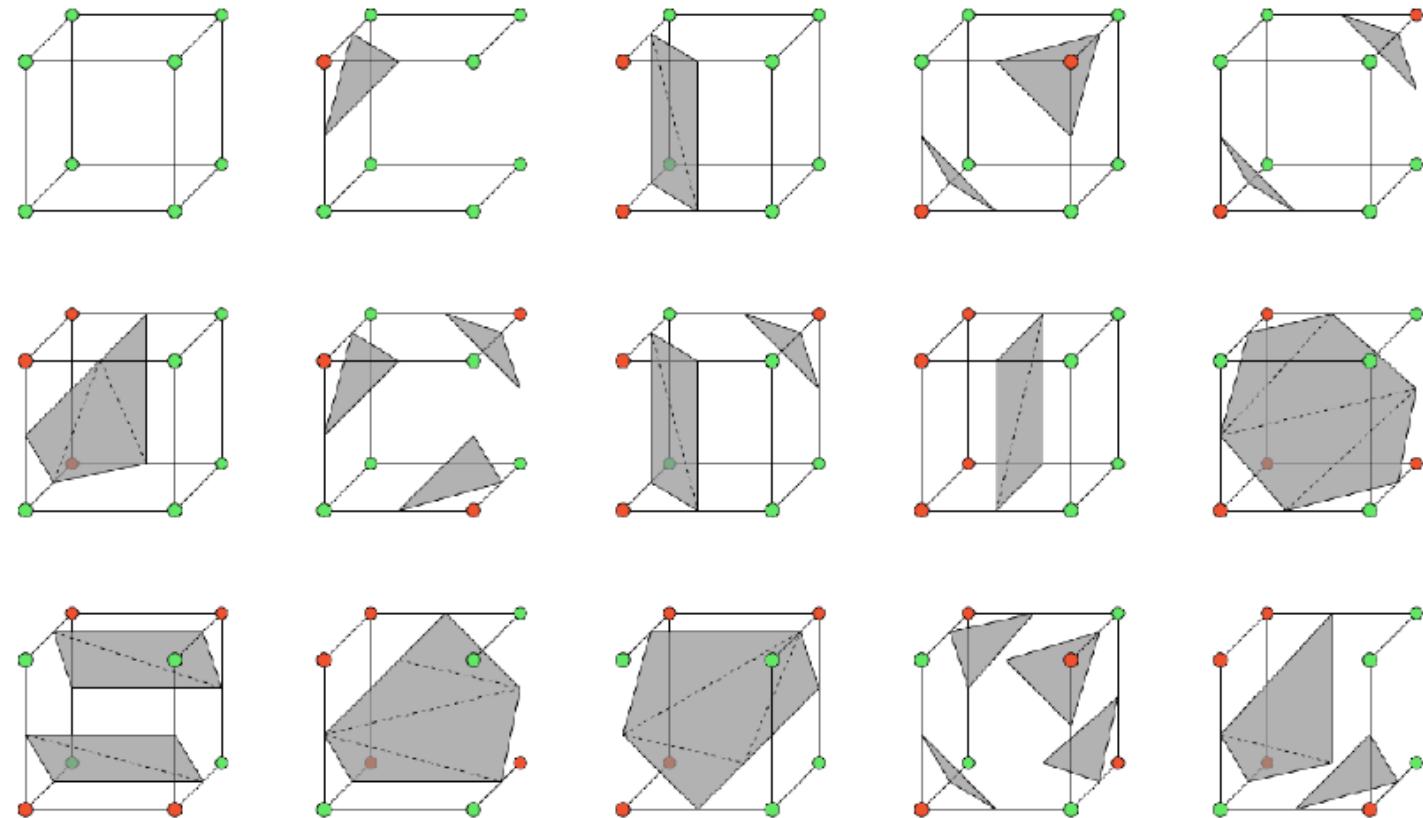
$$\text{index} = [v_1 \ v_2 \ v_3 \ v_4 \ v_5 \ v_6 \ v_7 \ v_8]$$



$$\text{index} = [0 \ 0 \ 1 \ 0 \ 0 \ 0 \ 0 \ 1] = 33$$

Marching cube

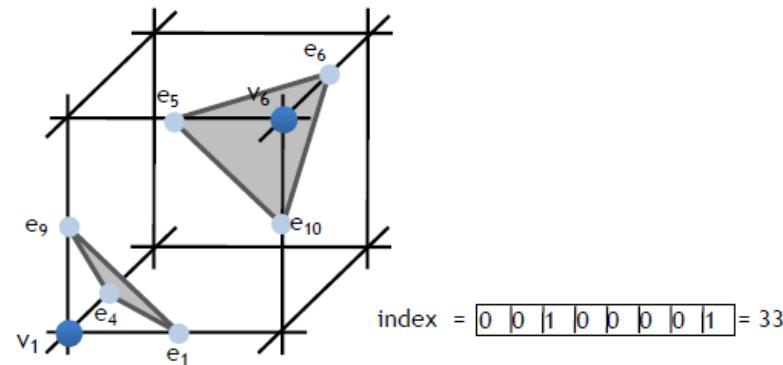
Unique cases (by rotation, reflection and complement)



Marching cube

Using the case index, retrieve the connectivity in the look-up table

- Example: the entry for index 33 in the look-up table indicates that the cut edges are $e_1; e_4; e_5; e_6; e_9$ and e_{10} ; the output triangles are $(e_1; e_9; e_4)$ and $(e_5; e_{10}; e_6)$.

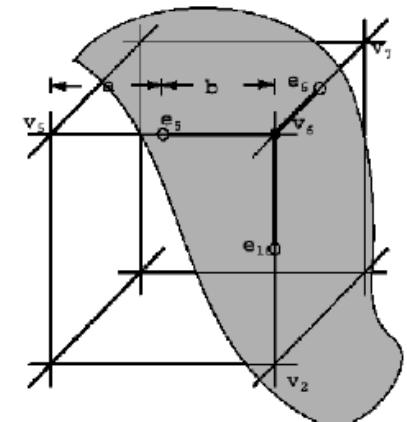


Marching cube

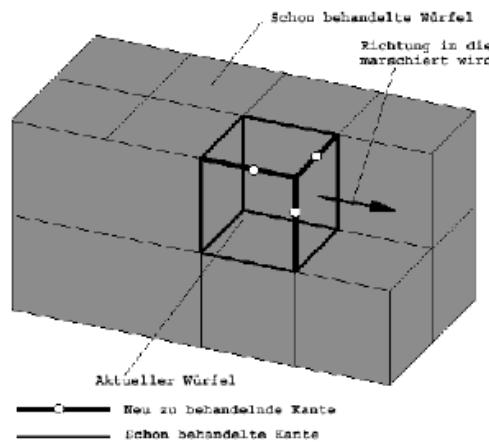
Compute the position of the cut vertices by linear interpolation:

$$\mathbf{v}_s = t\mathbf{v}_a + (1 - t)\mathbf{v}_b$$

$$t = \frac{F(\mathbf{v}_b)}{F(\mathbf{v}_b) - F(\mathbf{v}_a)}$$



Move to the next cube

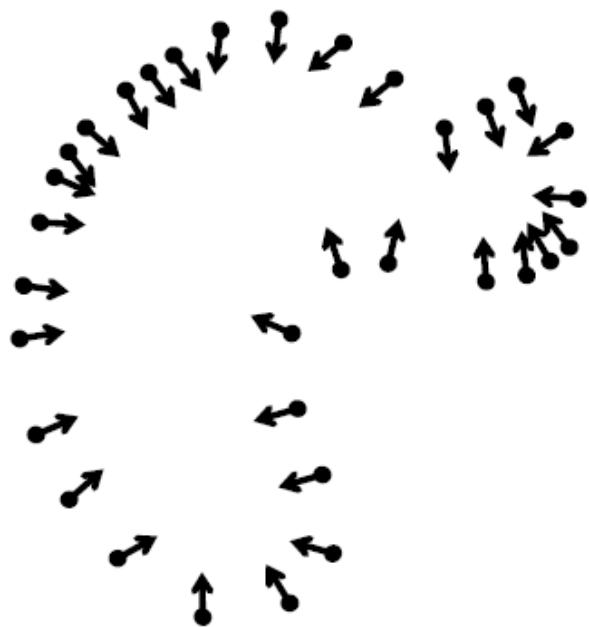


Poisson method

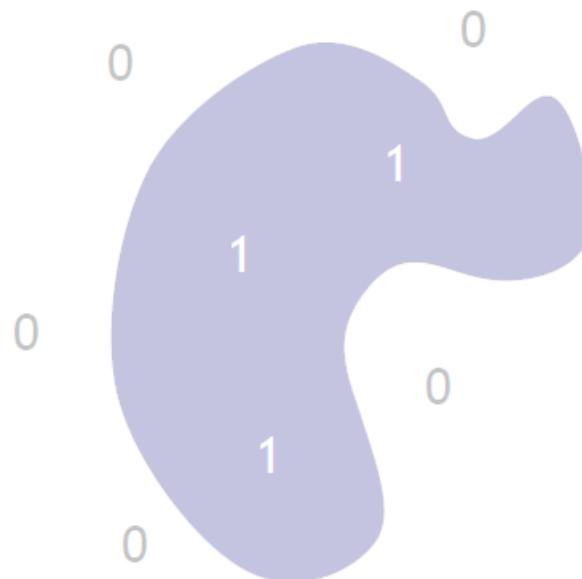
- Global fitting of an **indicator function** using Partial Differential Equation (PDE),



Poisson method



Oriented points

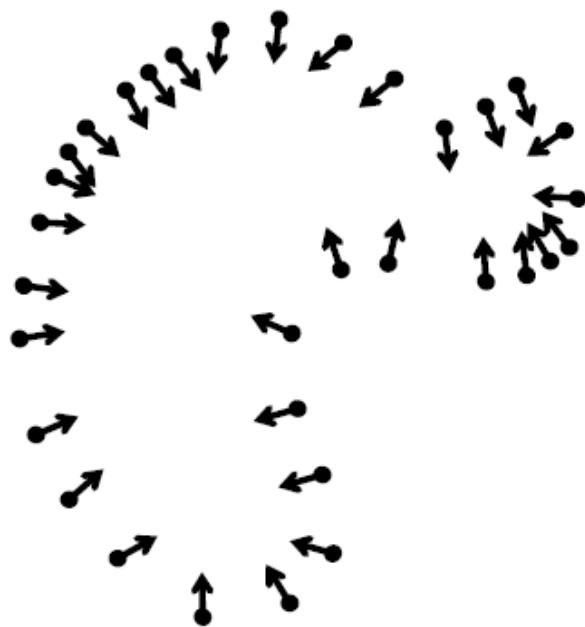


Indicator function

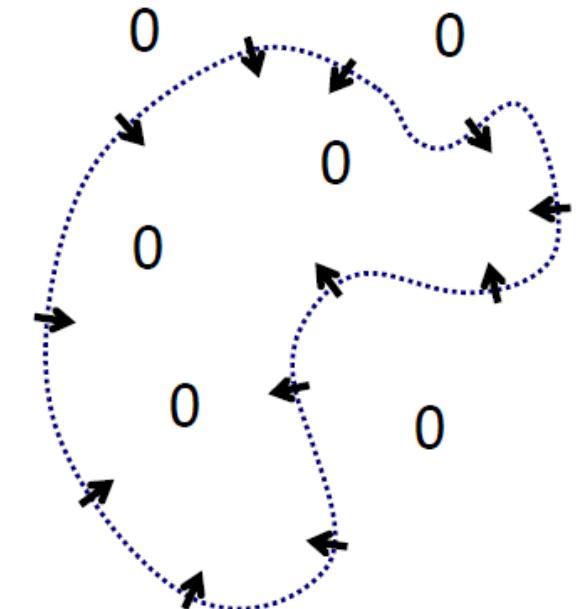
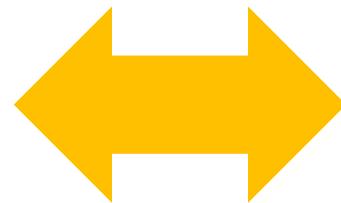
$$\chi_M$$

We don't know the indicator function ☹

Poisson method



Oriented points

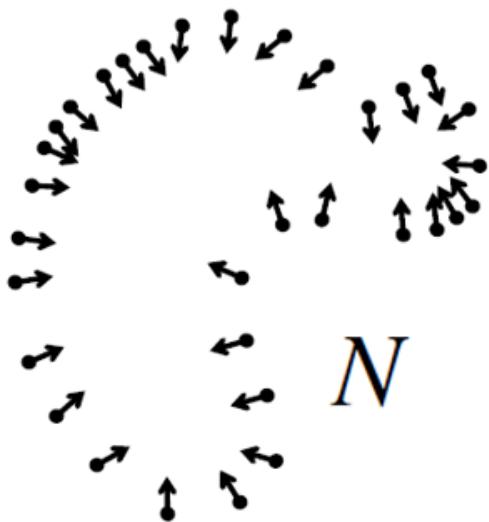


Indicator gradient

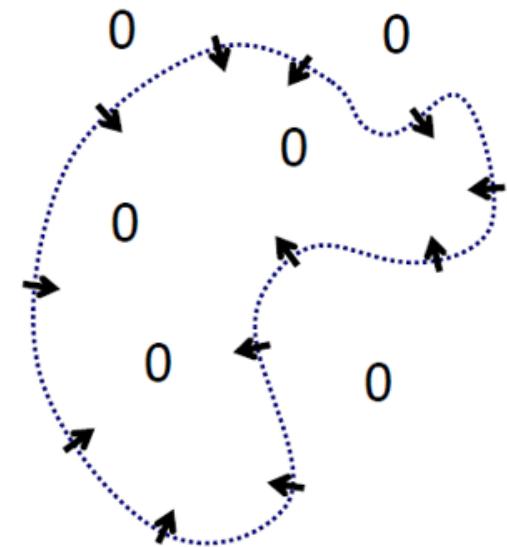
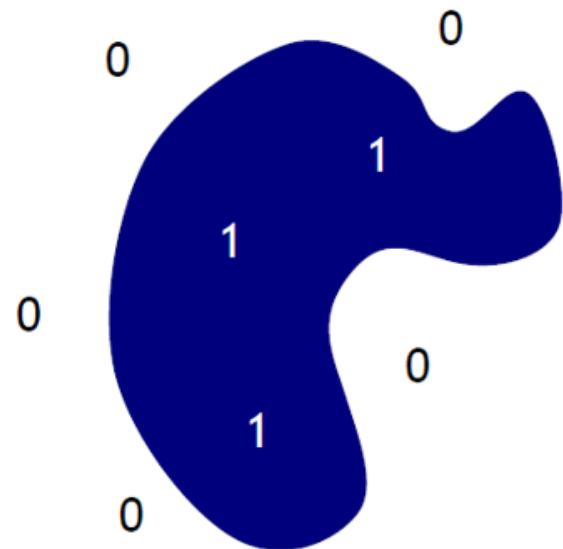
$$\nabla \chi_{\mathcal{M}}$$

But we can estimate its gradient! ☺

Poisson method



Oriented points



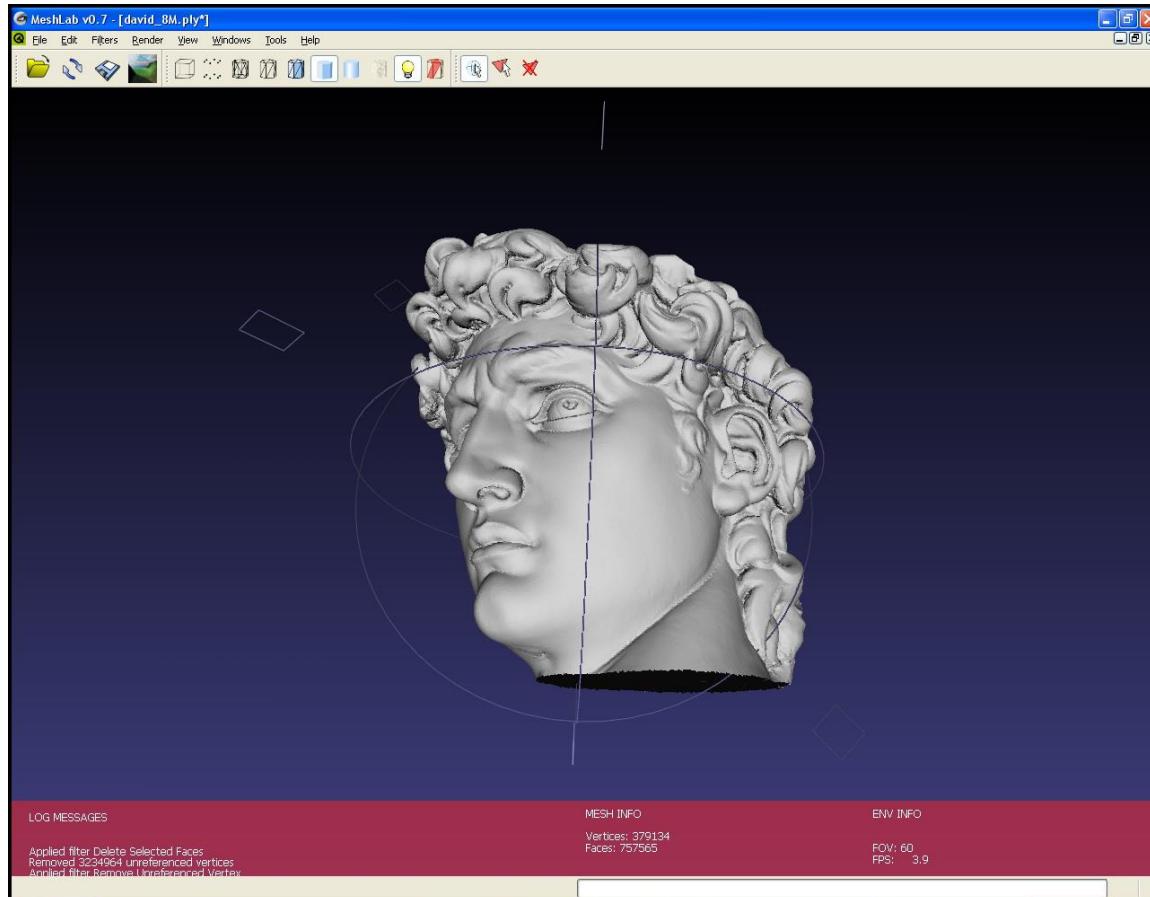
Reconstruct χ by solving
the Poisson equation

$$\Delta X_M = \nabla \cdot (\nabla X_M)$$
$$\nabla \cdot N$$

Poisson method



Meshlab



<https://www.meshlab.net/>

Homework1

1. Find public dataset of range images from different acquisition systems
 - E.g. <https://www.lidarusa.com/sample-data.html#>

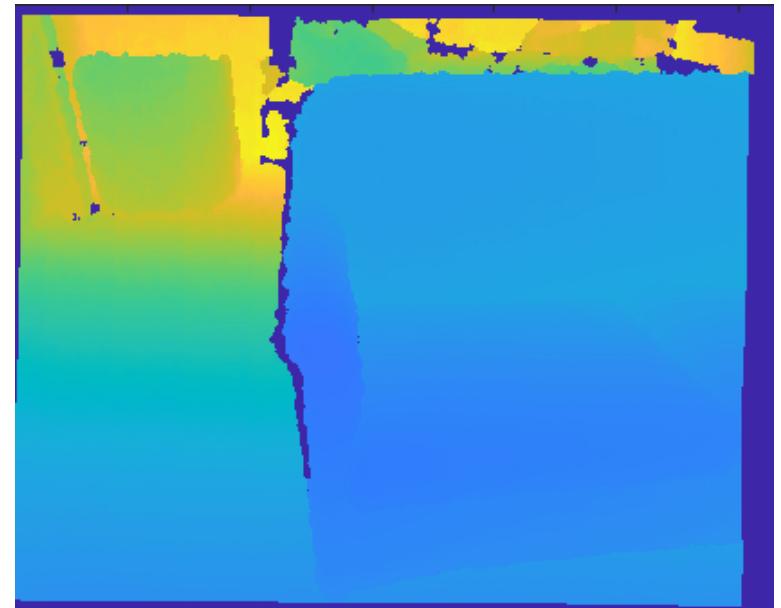


Homework2

2. Create a 3D cloud of points from a range image



Color image



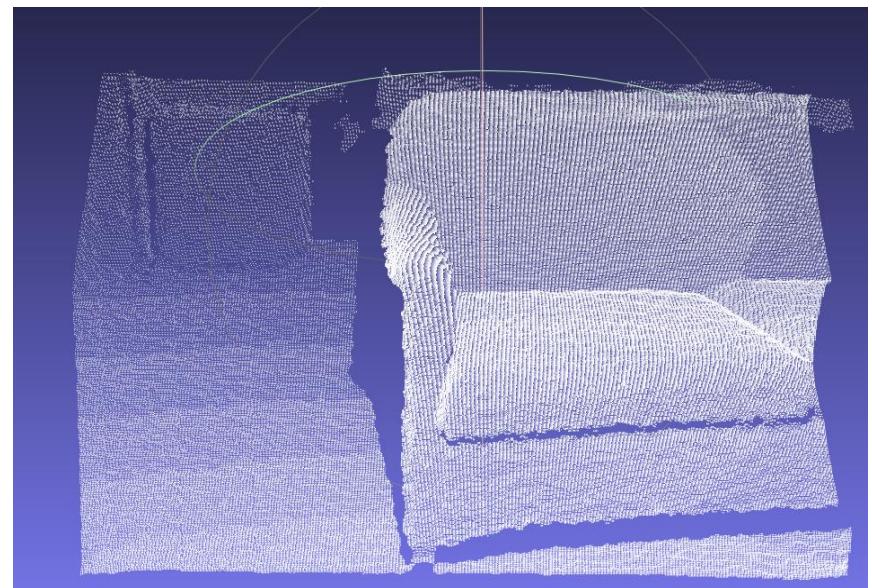
Range image

Homework2

2. Create a 3D cloud of points from a range image



Color image

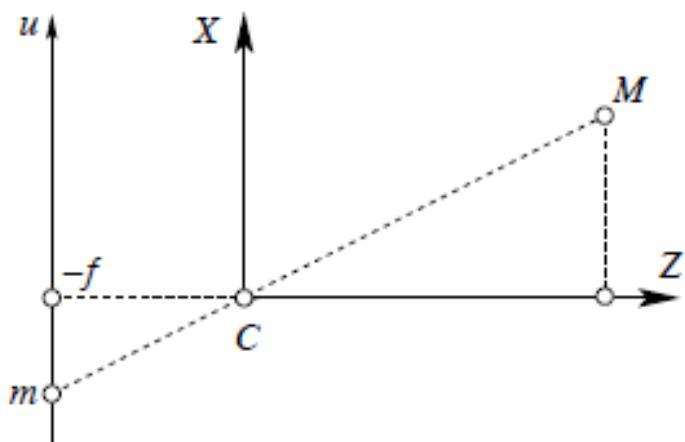


3D point cloud

Homework2

2. Create a 3D cloud of points from a range image

Suggestions: z is encoded in the range image, find the x and y coordinates from the projection equations



$$\begin{cases} u = k_u \frac{-f}{z} x + u_0 \\ v = k_v \frac{-f}{z} y + v_0 \end{cases},$$

The global reference system is on the camera, therefore only the main intrinsic parameters are required: (u_0, v_0) and $f_u = f k_u$, $f_v = f k_v$

Homework2

2. Create a 3 cloud of points from a range image

E.g.: take some samples from <http://redwood-data.org/3dscan/>



The RGB-D sequences were acquired with **PrimeSense Carmine** cameras

The focal length is 525 for both axes ($f_u=525$, $f_v=525$) and the principal point is ($u_0=319.5$, $v_0=239.5$).

Homework2

Export the cloud of point with this script and visualize it with Meshlab

```
function exportMeshToPly(vertices, faces, vertex_color, name)
if(max(max(vertex_color))<= 1.0)
    vertex_color = vertex_color.*256;
end
if(size(vertex_color,2) == 1)
    vertex_color = repmat(vertex_color,1,3);
end
vertex_color = uint8(vertex_color);
fidply = fopen([name '.ply'],'w');

fprintf(fidply, 'ply\n');
fprintf(fidply, 'format ascii 1.0\n');
fprintf(fidply, 'element vertex %d\n', size(vertices,1));
fprintf(fidply, 'property float x\n');
fprintf(fidply, 'property float y\n');
fprintf(fidply, 'property float z\n');
fprintf(fidply, 'property uchar red\n');
fprintf(fidply, 'property uchar green\n');
fprintf(fidply, 'property uchar blue\n');
fprintf(fidply, 'element face %d\n', size(faces,1));
fprintf(fidply, 'property list uchar int vertex_index\n');
fprintf(fidply, 'end_header\n');

for i=1:size(vertices,1)
    fprintf(fidply, '%f %f %f %d %d %d\n',vertices(i,1), vertices(i,2), vertices(i,3), vertex_color(i,1), vertex_color(i,2), vertex_color(i,3));
end

for i=1:size(faces,1)
    fprintf(fidply, '3 %d %d %d\n',faces(i,1)-1, faces(i,2)-1, faces(i,3)-1);
end

fclose(fidply);

end
```

Homework2

- Try other public available datasets like:

<https://rgbd-dataset.cs.washington.edu/index.html>

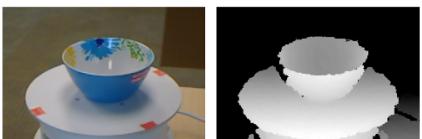
RGB-D Object Dataset

[Home](#) [People](#) [Dataset](#) [Results](#) [Software](#) [Demos](#)

News

- April 5, 2014 - The RGB-D Scenes Dataset v.2 is now available [here!](#)! It contains 14 new scenes reconstructed from RGB-D videos with furniture and tabletop objects, as well as Trimble 3D Warehouse objects use
- July 14, 2013 - Code for HMP features now available [here](#). It achieves state-of-the-art [results](#) on the RGB-D Object Dataset!
- December 13, 2012 - Software and data for detection-based object labeling in Kinect videos now available [here](#).
- October 3, 2012 - The dataset is now available for download directly from the website! No more sending emails necessary (questions and suggestions are, of course, still welcomed!).
- April 6, 2012 - RGB-D kernel descriptors are now available.
- March 22, 2012 - 3D reconstructions created by aligning video frames of all 8 scenes in the RGB-D Scenes Dataset are now available.
- June 20, 2011 - Pose annotations for all 300 objects in the RGB-D Object Dataset are now available.

Overview



The RGB-D Object Dataset is a large dataset of 300 common household objects. The objects are organized into 51 categories arranged using WordNet hypernym-hyponym relationships (similar to ImageNet). This dataset and aligned 640x480 RGB and depth images at 30 Hz. Each object was placed on a turntable and video sequences were captured for one whole rotation. For each object, there are 3 video sequences, each recorded with different angles with the horizon.

Unlike many existing datasets, such as Caltech 101 and ImageNet, objects in this dataset are organized into both *categories* and *instances*. In these datasets, the class *dog* contains images from many different dogs and the RGB-D Object Dataset the category *soda can* is divided into physically unique instances like *Pepsi Can* and *Mountain Dew Can*. The dataset also provides ground truth pose information for all 300 objects.

Here are some example objects that have been segmented from the background.

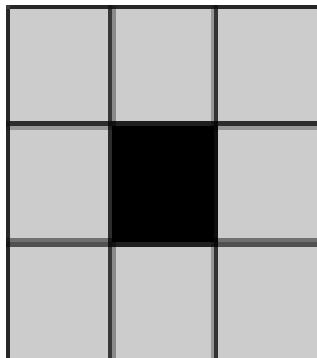


Use (and understand) the available code !

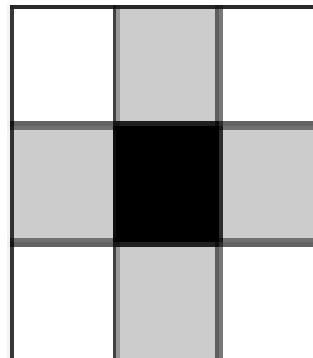
Homework 3

Mesh reconstruction from range image

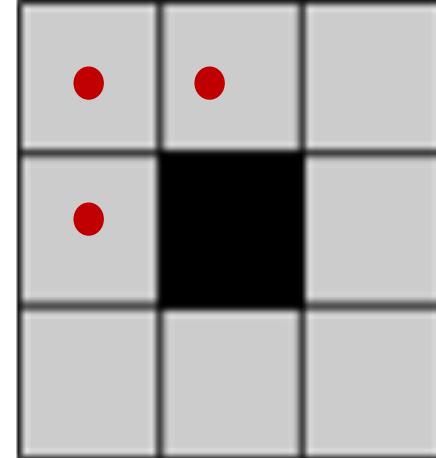
- **Idea:** points are on a regular grid where the connectivity can be inherited from the pixel neighbourhood



8-neigh



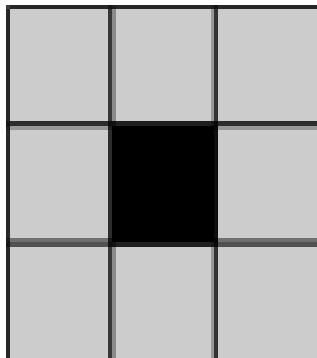
4-neigh



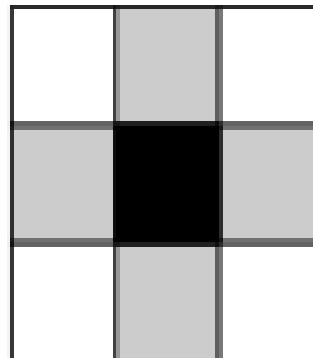
Homework 3

Mesh reconstruction from range image

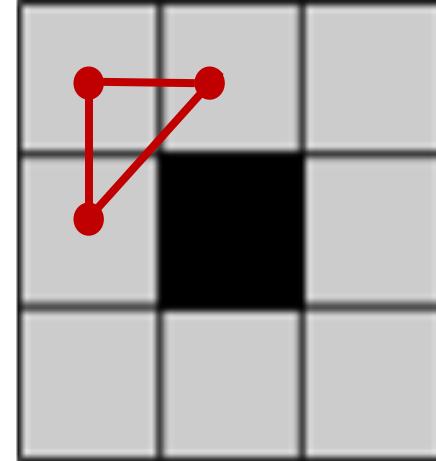
- **Idea:** points are on a regular grid where the connectivity can be inherited from the pixel neighbourhood



8-neigh



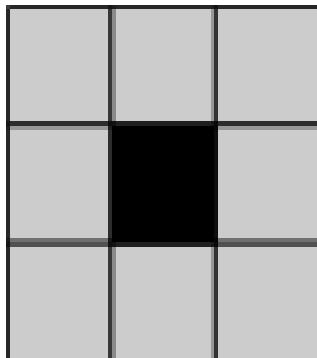
4-neigh



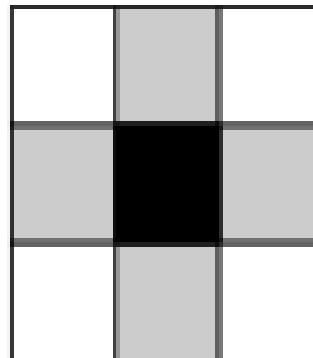
Homework 3

Mesh reconstruction from range image

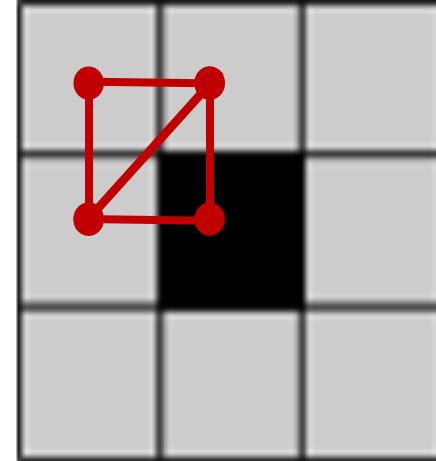
- **Idea:** points are on a regular grid where the connectivity can be inherited from the pixel neighbourhood



8-neigh



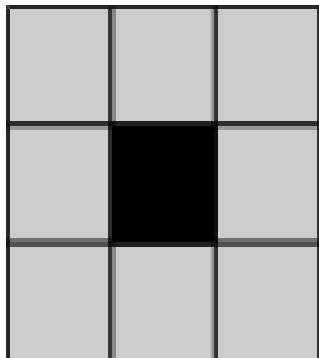
4-neigh



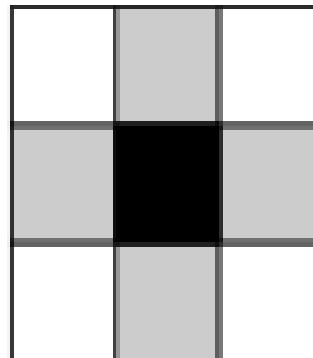
Homework 3

Mesh reconstruction from range image

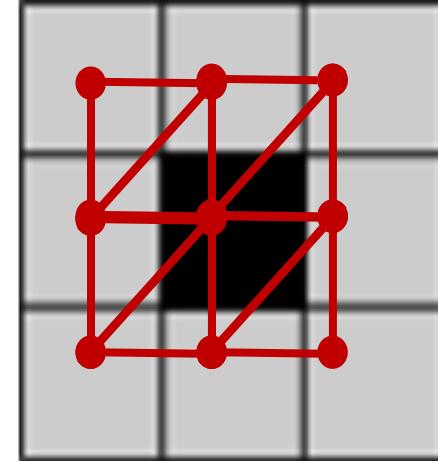
- **Idea:** points are on a regular grid where the connectivity can be inherited from the pixel neighbourhood



8-neigh



4-neigh



Homework 3

- Points are on a regular grid where the connectivity can be inherited from the pixel neighbourhood,

BUT...

- 1) Not all the pixels on the range image are the projection of a point on the 3D space!
 - a binary mask can be used to define valid points,
- 2) (optional) Nearby pixels should not correspond to nearby points on the 3D space!
 - a robust strategy can be used to remove long edges.