

# Registration

Umberto Castellani

Robotics Vision and Control

# 3D modelling from reality pipeline



# Homework1

## Play with Zephyr:

- **TUTORIAL 1:** <https://www.3dflow.net/technology/documents/3df-zephyr-tutorials/convert-photos-3d-models-3df-zephyr/>
- **VIDEO-TUTORIAL:** <https://www.3dflow.net/it/tutorial-per-3df-zephyr/>
- **VIDEO-TUTORIAL WITH DATA:** <https://www.3dflow.net/it/community-fotogrammetria/3df-zephyr-vetrina-di-ricostruzioni/>

 Cherub Statue  
by 3dflow



**Cherub Statue**

↓ Clicca qui per scaricare il progetto .zep

📷 Clicca qui per scaricare il dataset completo

Ricostruzione di una statuetta con 65 fotografie.

Storico dataset utilizzato nel **primo tutorial di 3DF Zephyr** nel quale viene mostrata la classica modalità di acquisizione per piccoli oggetti.

# Homework2

Create your 3D model of your physical object:

- This model will be used in our robotics, vision and control pipeline




The object should be grabbed by the robot



# Registration: overall aim

Align partially-  
overlapping views  
given initial guess  
for relative transform




# Overall aim




Two views registration

# Overall aim



**Multiple views registration**

# Overall aim



**Multiple views**



# Two-view registration

## Problem statement

Given a pair of views  $D$  and  $M$  representing two scans (partial 3D views) of the same object, **registration** is the problem of finding the parameters  $\mathbf{a}^*$  of the transformation function  $T(\mathbf{a}; D)$  which best aligns  $D$  to  $M$ .

The diagram illustrates the components of the registration equation  $\mathbf{a}^* = \arg \min_{\mathbf{a}} E(T(\mathbf{a}, D), M)$ . It features three labeled boxes with arrows pointing to parts of the equation: a red box labeled 'Moving view' with a red arrow pointing to  $D$  inside the transformation function  $T$ ; a blue box labeled 'Fixed view' with a blue arrow pointing to  $M$ ; and a black box labeled 'Error function' with a black arrow pointing to  $E$ . A black arrow also points from the label 'Transformation function' to the  $T$  symbol. The variable  $\mathbf{a}$  is the parameter being minimized over.

$$\mathbf{a}^* = \arg \min_{\mathbf{a}} E(T(\mathbf{a}, D), M)$$

Labels and arrows in the diagram:

- Moving view** (red box) points to  $D$  in  $T(\mathbf{a}, D)$ .
- Fixed view** (blue box) points to  $M$ .
- Error function** (black box) points to  $E$ .
- Transformation function** (black box) points to  $T$ .


# Two-view registration

## Problem statement

$D$  is the **data** view (moving)

$M$  is the **model** view (fixed)

When the transformation  $T$  is applied to the data view, data view is moved to the model view and the alignment is obtained



# Two-view registration

- **The transformation function.** The transformation function  $T$  usually implements a rigid transformation of the 3D space. It uses a translation vector  $\mathbf{t}$  and a rotation matrix  $\mathbf{R}$  whose values are encoded or parametrized in the parameter vector  $\mathbf{a}$ .

$$T(\mathbf{a}; \mathbf{D}) = \mathbf{R}\mathbf{D} + \mathbf{t}$$

Transformation function for rigid registration

**Note** that the transformation function may also handle deformations but this requires a more complex formulation...

# Two-view registration

- **The error function.** The error function  $E$  measures the registration error or dissimilarity between  $\mathbf{D}$  and  $\mathbf{M}$  after alignment. When the transformation function  $T$  is rigid,  $E$  is a measure of congruence between the two views.


$$\begin{aligned} E &= ||T(\mathbf{a}; \mathbf{D}) - \mathbf{M}||^2 \\ &= ||(\mathbf{R}\mathbf{D} + \mathbf{t}) - \mathbf{M}||^2 \end{aligned}$$

Error function for rigid registration (i.e.,  
extrinsic similarity)


**Note** that the error function may also handle deformations but this requires a more complex formulation...

# Two-view registration


- **The optimisation method.** This is the method or algorithm used to find the minimizer of error function. The gold standard is the **ICP algorithm** which was specifically designed for the problem at hand.




Starting views



Iter 1




Iter 2



Iter 3

...



Iter N

**Note:** General purpose optimisation methods such as **Levenberg-Marquardt** have also been used for this problem.

# Rigid alignment


- When a set of corresponding points between  $\mathbf{D}$  and  $\mathbf{M}$  is available the rigid alignment can be estimated solving an **Orthogonal Procrustes problem**
  - Data is represented by a set of points  $\mathbf{d}_1, \dots, \mathbf{d}_N$
  - Model is represented by a set of points  $\mathbf{m}_1, \dots, \mathbf{m}_N$Where each pair  $(\mathbf{d}_i, \mathbf{m}_i)$  represent the same surface point observed from two different views

## Computer vision course!

# Iterative closest point (ICP)

- In standard scenarios the pairs  $(\mathbf{d}_i, \mathbf{m}_i)$  of corresponding points are not available!

**Main idea:** estimate the corresponding points with **closest points**.



# Iterative closest point (ICP)

- Data  $\mathbf{D}$  is represented by a set of points  $\mathbf{d}_1, \dots, \mathbf{d}_{N_d}$
- Model  $\mathbf{M}$  is represented by a set of points  $\mathbf{m}_1, \dots, \mathbf{m}_{N_m}$ 
  - Fixing  $\mathbf{d}_i \in \mathbf{D}$  the closest point  $\mathbf{m}_j \in \mathbf{M}$  is computed such that:

$$j = \arg \min_{j \in \{1, \dots, N_m\}} \|(\mathbf{R}\mathbf{d}_i + \mathbf{t}) - \mathbf{m}_j\|^2$$



# Iterative closest point (ICP)

- The iterative closest point algorithm (ICP) for the registration of 3D points is based on the iteration between two main steps:
  - Estimate the pair of corresponding points ( $\mathbf{d}_i, \mathbf{m}_j$ ) by computing the closest points,
  - Solve the Orthogonal Procrustes problem to estimate  $\mathbf{R}$  and  $\mathbf{t}$



$$T(\mathbf{a}; \mathbf{D}) = \mathbf{R}\mathbf{D} + \mathbf{t} \quad \mathbf{a} = (\mathbf{R}, \mathbf{t}) \quad \text{Transformation function}$$

$$E_{ICP}(\mathbf{a}, D, M) = \sum_{i=1}^{N_d} \|(\mathbf{R}\mathbf{d}_i + \mathbf{t}) - \mathbf{m}_j\|^2 \quad \text{Error function}$$

Optimization method

# Iterative closest point (ICP)

- **Summary of ICP algorithm:**

1. For each data-point  $\mathbf{d}_i \in D$ , compute the closest point  $\mathbf{m}_j \in M$ ;
2. With the correspondences  $(\mathbf{d}_i, \mathbf{m}_j)$  from step 1, compute the new transformation parameters  $\mathbf{a} = (\mathbf{R}, \mathbf{t})$ ;
3. Apply the new transformation parameters  $\mathbf{a}$  from step 2 to the point cloud  $D$ ;
4. If the change in  $E_{ICP}(\mathbf{a}, D, M)$  between two successive iterations is lower than a threshold terminate, else goto step 1.

**This algorithm is guaranteed to converge monotonically to a local solution !**

# Iterative closest point (ICP)

- **Open issue** of standard ICP:
  - The two views must be close to each other. If not, ICP will probably get stuck in a local minimum.



a good initialization of  $\mathbf{a}$  is needed!

# Iterative closest point (ICP)

- **Open issue** of standard ICP:
  - The two views must fully overlap, or the data-view  $D$  must be a subset of the model-view  $M$ . Otherwise if a data point has no corresponding model point, this will create a spurious correspondence,



an outlier rejection strategy is needed!

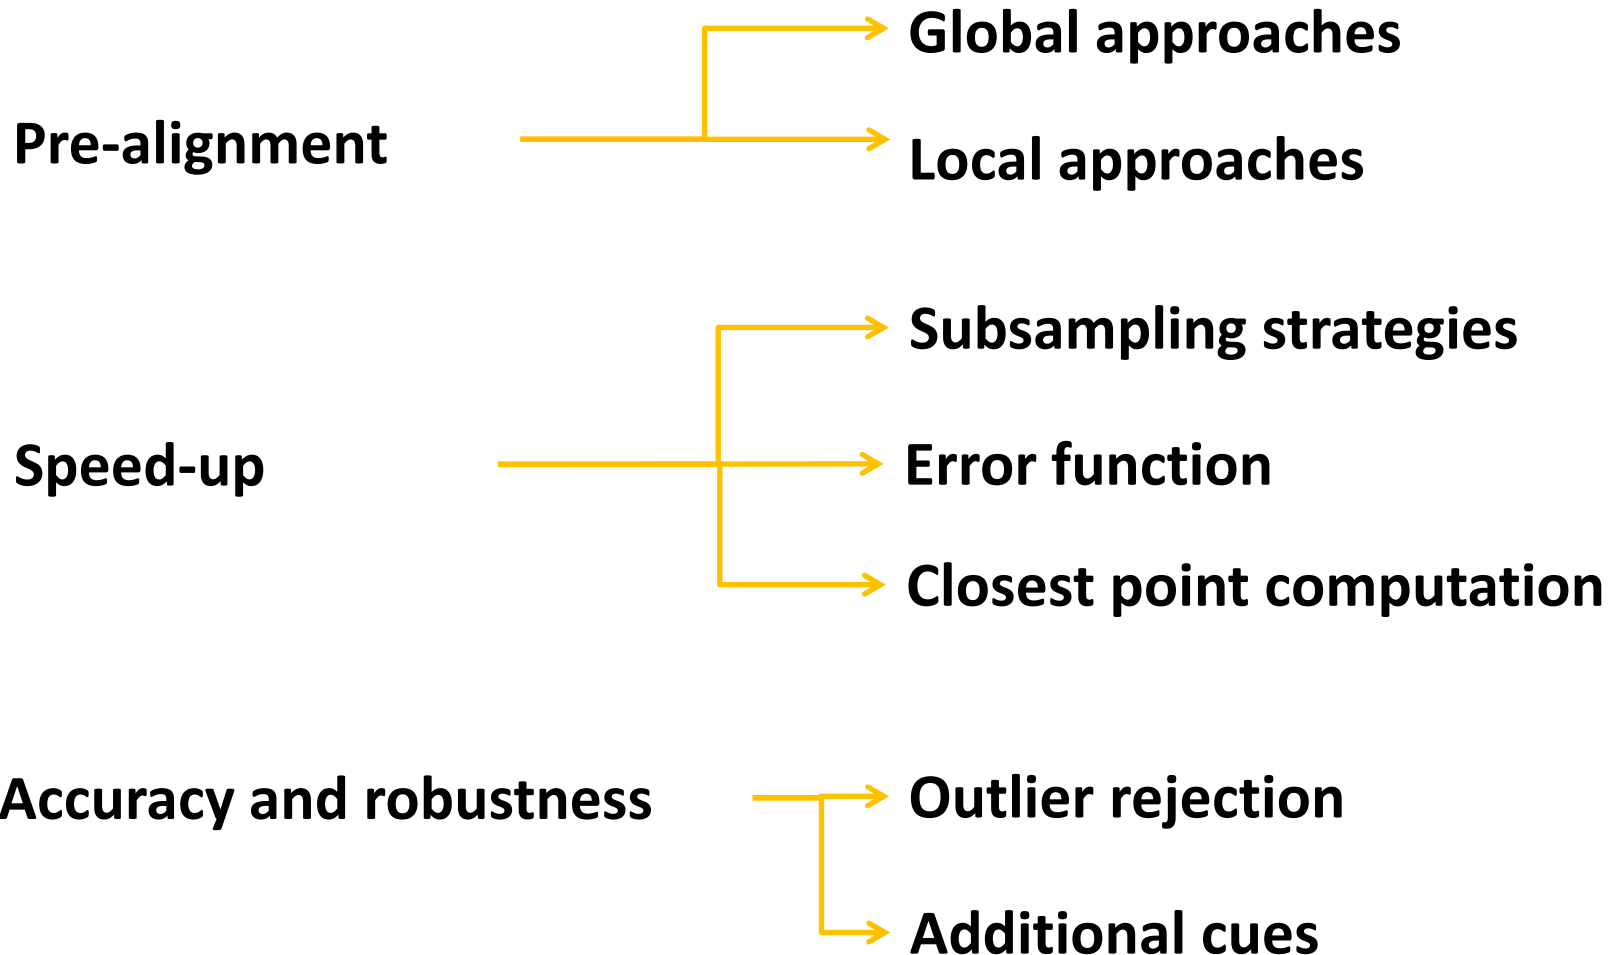
# Iterative closest point (ICP)

- **Open issue** of standard ICP:
  - The estimation of closest point is computational expensive,

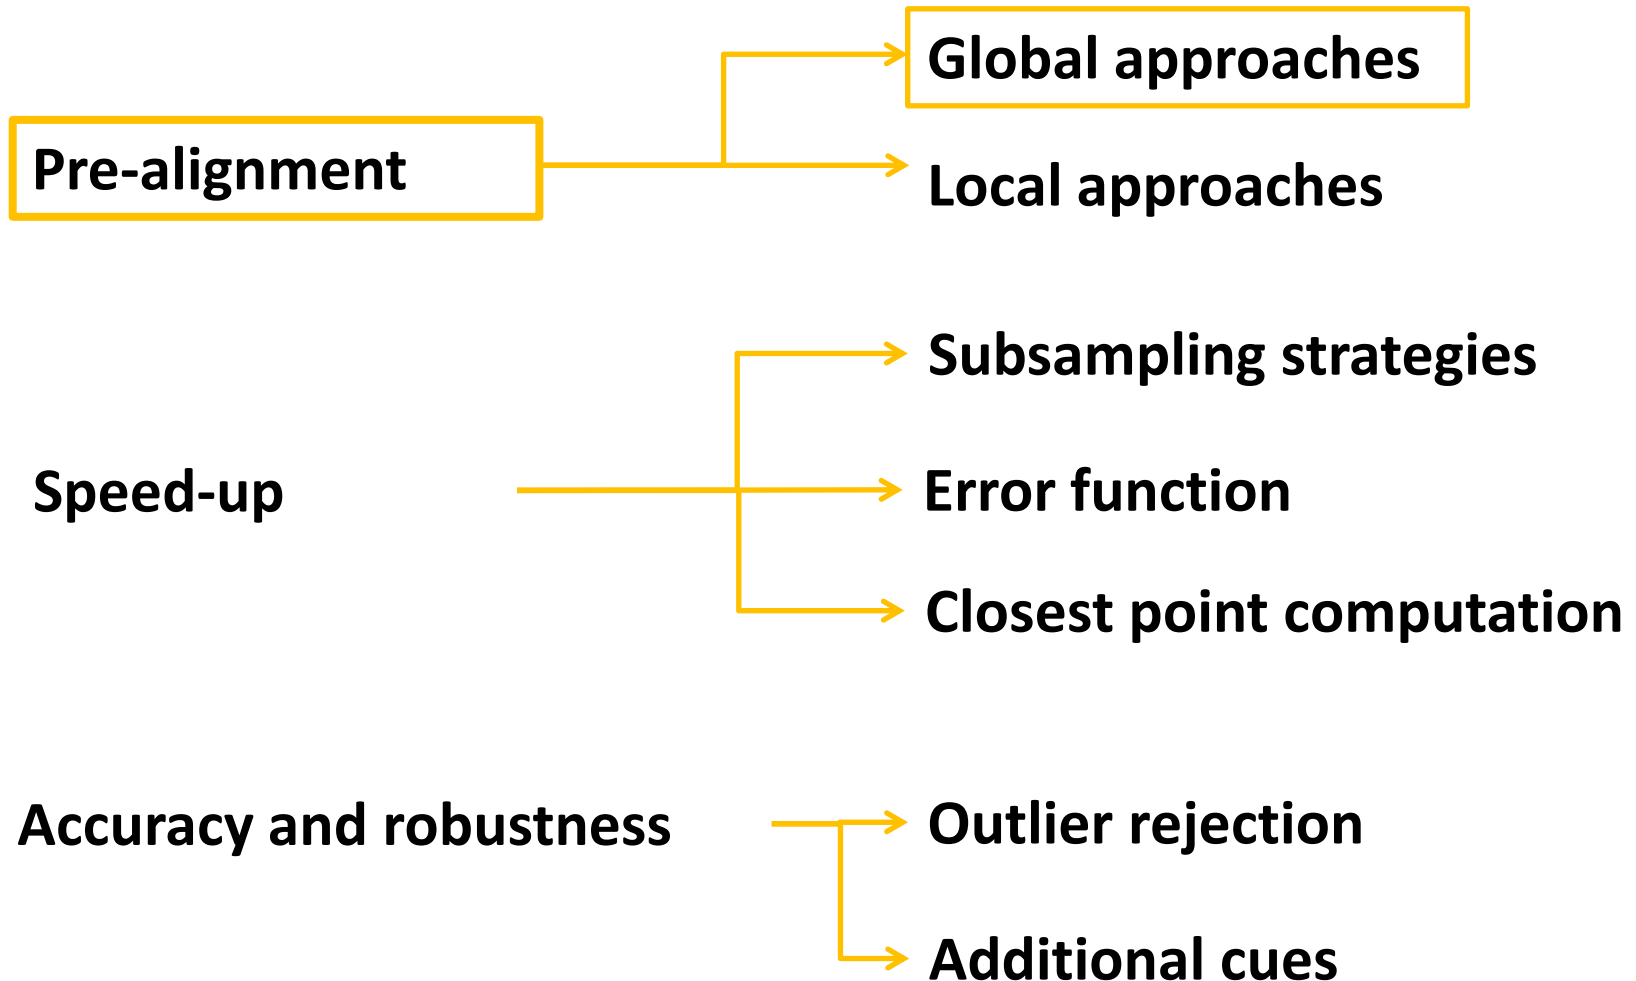


A strategy to improve the speed is needed!

# ICP variants



# ICP variants



# Pre-alignment

**Global pre-alignment** (i.e., all points are used)

$$x' = Rx + t$$

- How to remove **translation** and **rotation** ambiguity?
- Find some “**canonical**” placement of the shape  $X$  in  $\mathbb{R}^3$



# Pre-alignment

- **Extrinsic centroid** (a.k.a. **center of mass**, or **center of gravity**):

$$x_0 = \int_X x dx$$

- Set  $t = -x_0$  to resolve translation ambiguity.
- Three degrees of freedom remaining...

# Pre-alignment


- Find the direction  $d_1$  in which the surface has **maximum extent**.
- Maximize **variance** of projection of  $X$  onto  $d_1$

$$\begin{aligned}
 d_1 &= \arg \max_{d_1: \|d_1\|_2=1} \int_X (d^\top x)^2 dx && \text{(optimal } d \text{)} \\
 &= \arg \max_{d_1: \|d_1\|_2=1} d_1^\top \left( \int_X x x^\top dx \right) d_1 && (d^\top x)(x^\top d) \\
 &= \arg \max_{d_1: \|d_1\|_2=1} d_1^\top \Sigma_X d_1 && d^\top (x x^\top) d
 \end{aligned}$$


- $\Sigma_X$  is the **covariance matrix**
- **Second-order geometric moments** of  $X$ :  $\sigma_{ij} = \int_X x^i x^j dx$
- $d_1$  is the first **principal direction**

# Pre-alignment

- Project  $X$  on the plane orthogonal to  $d_1$ .
- Repeat the process to find second and third principal directions  $d_2, d_3$ .



# Pre-alignment



**Canonical basis**

- $d_1 \perp d_2 \perp d_3$  span a canonical orthogonal basis for  $X$  in  $\mathbb{R}^3$ .

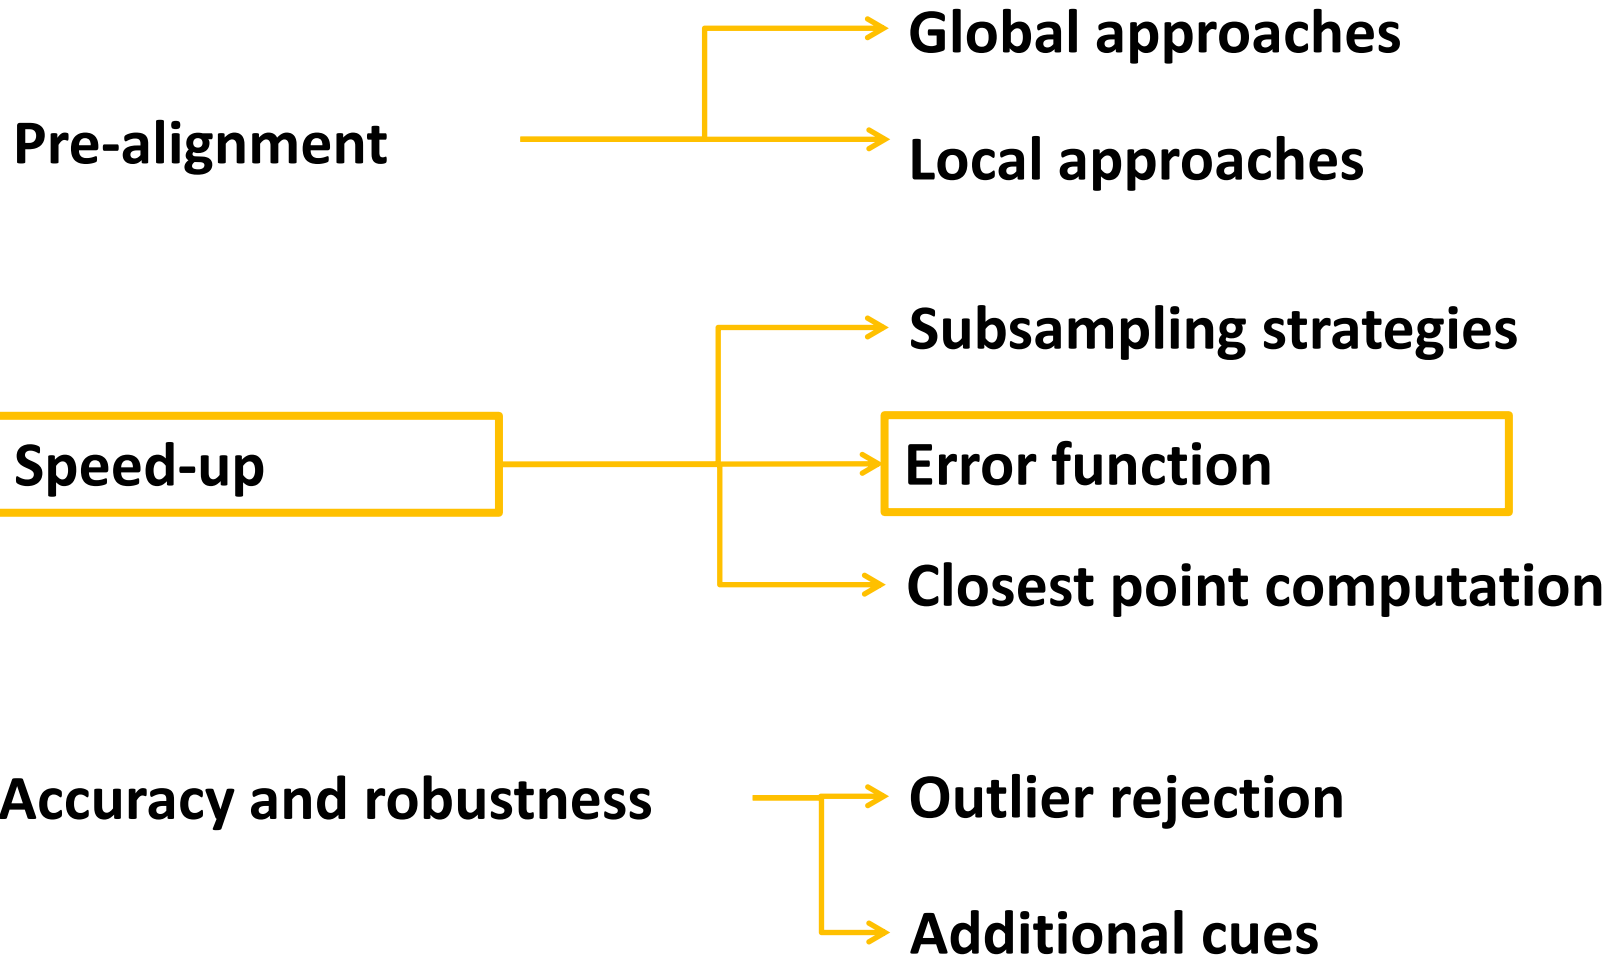
## How to get rid of the rotation ambiguity?

- Direction maximizing  $d_1^\top \Sigma_X d_1 = \text{largest eigenvector}$  of  $\Sigma_X$ .
- $d_2$  and  $d_3$  correspond to the second and third eigenvectors of  $\Sigma_X$ .
- $\Sigma_X$  admits **unitary diagonalization**  $\Sigma_X = U^\top \Lambda U$ .
- Setting  $R = U^\top$  aligns  $d_1, d_2, d_3$  with the **standard basis** axes  $e_1, e_2, e_3$ .
- **Principal component analysis** (PCA), a.k.a. **Karhunen-Loève transform** (KLT), or **Hotelling transform**.
- Bottom line: the transformation

$$(R, t) = (U^\top, -U^\top x_0)$$

brings the shape into a canonical configuration in  $\mathbb{R}^3$ .



# ICP variants



# Point-to-Plane Error Metric

Using point-to-plane distance instead of point-to-point lets flat regions slide along each other

[Chen & Medioni 91]



# Point-to-Plane Error Metric



- Error function:

$$E = \sum \left( (Rp_i + t - q_i) \cdot n_i \right)^2$$

where  $R$  is a rotation matrix,  $t$  is translation vector

- Linearize (i.e. assume that  $\sin \theta \approx \theta$ ,  $\cos \theta \approx 1$ ):

$$E \approx \sum \left( (p_i - q_i) \cdot n_i + r \cdot (p_i \times n_i) + t \cdot n_i \right)^2, \quad \text{where } r = \begin{pmatrix} r_x \\ r_y \\ r_z \end{pmatrix}$$

- Result: overconstrained linear system



# Point-to-Plane Error Metric

- Overconstrained linear system



$$\mathbf{A}x = b,$$
$$\mathbf{A} = \begin{pmatrix} \leftarrow & p_1 \times n_1 & \rightarrow & \leftarrow & n_1 & \rightarrow \\ \leftarrow & p_2 \times n_2 & \rightarrow & \leftarrow & n_2 & \rightarrow \\ \vdots & & & \vdots & & \end{pmatrix}, \quad x = \begin{pmatrix} r_x \\ r_y \\ r_z \\ t_x \\ t_y \\ t_z \end{pmatrix}, \quad b = \begin{pmatrix} -(p_1 - q_1) \cdot n_1 \\ -(p_2 - q_2) \cdot n_2 \\ \vdots \end{pmatrix}$$

- Solve

**See notes!**


$$x = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T b$$

# ICP variants




# Finding Corresponding Points

- Finding closest point is most expensive stage of the ICP algorithm
  - Brute force search –  $O(n)$
  - Spatial data structure (e.g., k-d tree) –  $O(\log n)$




# Projection to Find Correspondences

- Idea: use a simpler algorithm to find correspondences
- For range images, can simply project point [Blais 95]
  - Constant-time
  - Does not require precomputing a spatial data structure





# Projection-Based Matching

- Slightly worse performance per iteration
- Each iteration is one to two orders of magnitude faster than closest-point
- Result: can align two range images in a few milliseconds, vs. a few seconds



# ICP variants



**Accuracy and robustness**

**Outlier rejection**

**Additional cues**



# Outlier rejection

Given the residual errors  $\mathbf{e} = [e_1, \dots, e_{N_d}]$

**Robust statistics approach:**

Mean  Median


Variance  Median Absolute Deviation (MAD)

**X84 rule for automatic threshold estimation**

$$|e_i - \text{median}| < k \cdot MAD$$

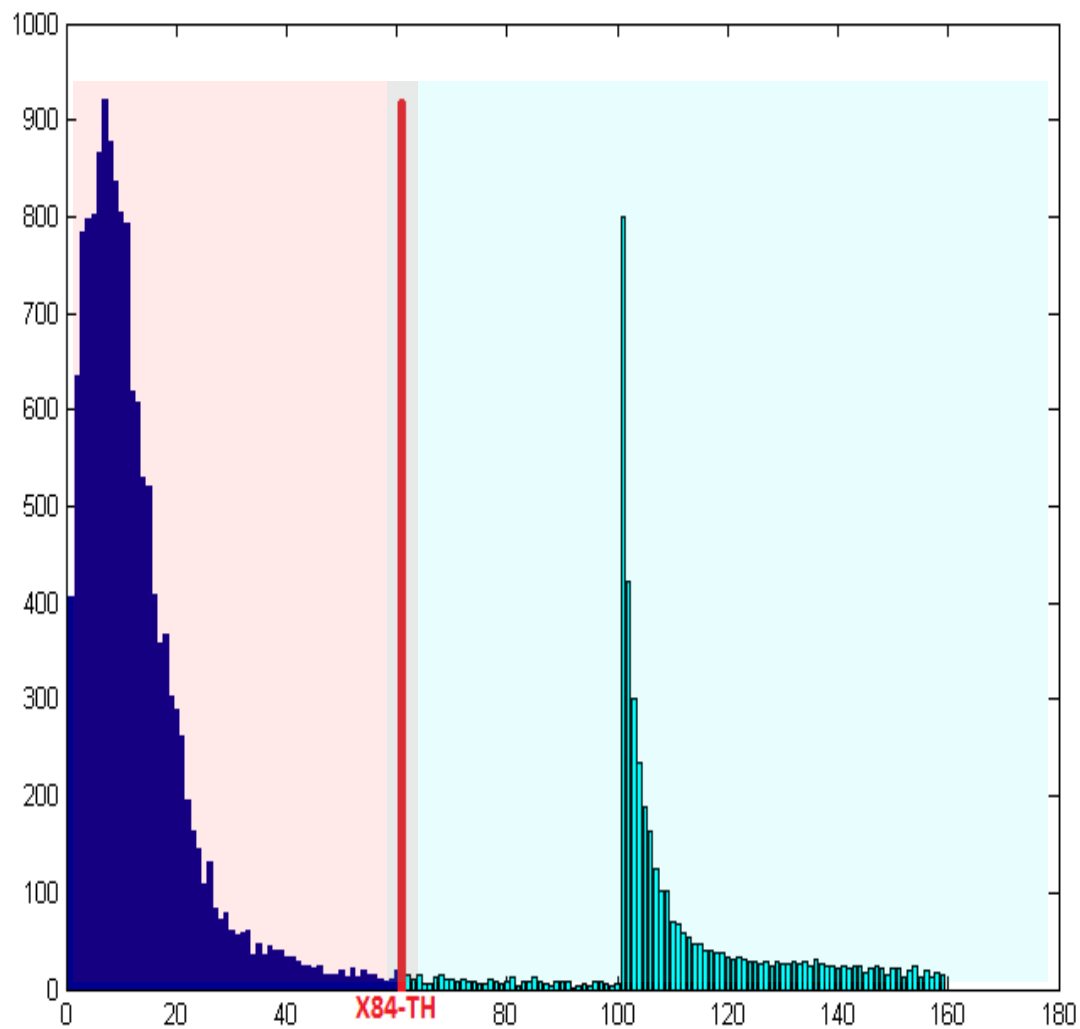
# See notes!

# Outlier rejection






# Outlier rejection

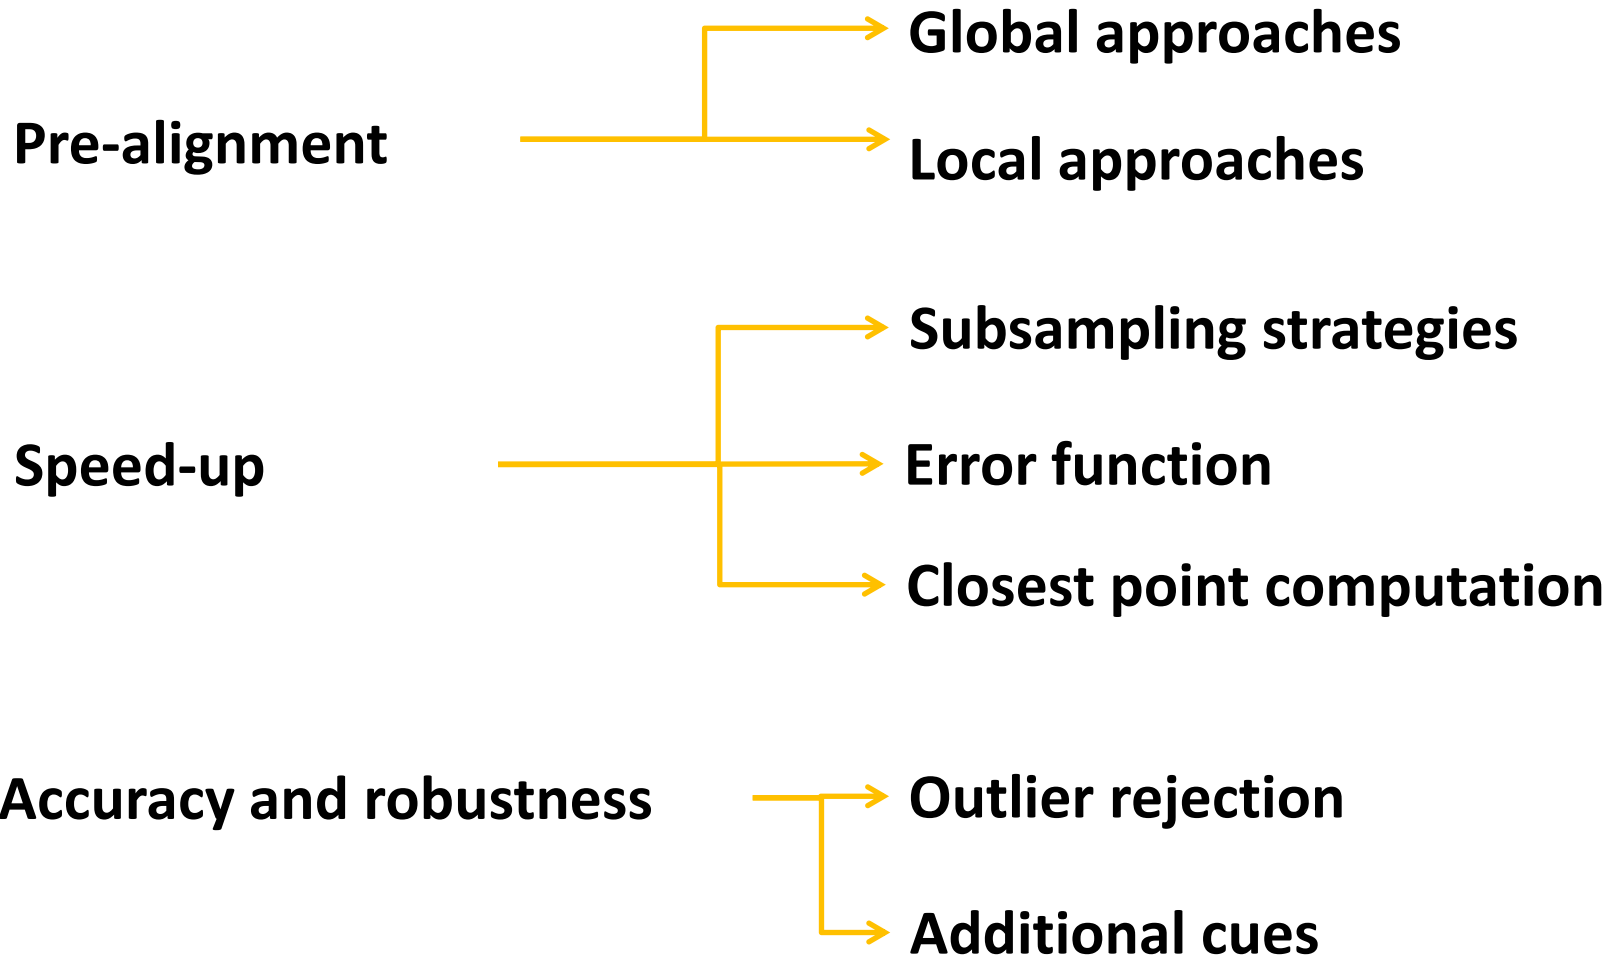


**Inliers**

**Outliers**




# ICP variants




# Global Registration Goal

- Given:  $n$  scans around an object
- Goal: align them all
- First attempt: ICP each scan to one other and then bring views to global ref system by concatenation



# Global Registration Goal

- Problem: small errors from pairwise views are amplified when sequential views are involved
- Advanced method for distributing accumulated error among all scans




# Bad ICP in Globalreg

One bad ICP can throw off the entire model!




Correct Globalreg



Globalreg Including Bad ICP

# Homework 3

Compute the canonical orientation of your object:



Zephyr brings the object to an arbitrary position

# Homework 4

- Global alignment by accumulation of rigid transforms